

Question - 1**SQL: Detecting Potential Payment Fraud in an Online Marketplace**

The company needs a report identifying users who have failed transactions using different payment methods. Failed transactions have "Failed" in the *status* field.

The result should have the following columns: *user_id* / *failed_transactions* / *distinct_payment_methods*.

- *user_id* – User attempting multiple failed transactions.
- *failed_transactions* – Total number of failed transactions.
- *distinct_payment_methods* – Total number of unique payment methods used.

Note:

- Only users who have made more than 5 failed transactions in the entire dataset should be included in the report.
- Row order does not matter.

▼ Schema

transactions

Name	Type	Constraint	Description
transaction_id	INT	PRIMARY KEY	Unique identifier for a transaction
user_id	INT		User attempting the payment
payment_method	VARCHAR(255)		Payment method used
amount	DECIMAL(10,2)		Transaction amount
transaction_date	DATE		Date of the transaction
status	VARCHAR(255)		Status of transaction

▼ Sample Data Tables

transactions

transaction_id	user_id	payment_method	amount	transaction_date	status
101	202	Credit Card	200.43	2025-02-16	Completed
102	203	Netbanking	3233.10	2025-03-11	Failed
103	203	Netbanking	1195.35	2025-02-24	Failed
104	203	Debit Card	376.11	2025-03-10	Failed
105	203	Netbanking	112.01	2025-04-04	Failed
106	203	Credit Card	111.1	2025-09-12	Failed
107	203	Debit Card	2344.5	2025-10-03	Failed

Sample Output

```
user_id    failed_transactions  distinct_payment_methods
203          6                      3
```

Explanation

The user with `user_id` 203 attempted 6 distinct transactions, which have the status as "Failed", thus it is a potential case of fraud.

Question - 2

SQL: Average Response Time

A customer support team wants to analyze response times for resolving tickets to identify performance metrics and improve service quality. The goal is to generate a report calculating the average response time for successfully resolved customer support tickets. Resolved tickets have a value in the `resolved_at` field.

The result should have the following columns: `average_response_time`.

- `average_response_time` - The average time between `created_at` and `resolved_at`, calculated in hours and set to two decimal places, including trailing zeros if necessary (e.g., 5.00).

▼ Schema

support_tickets

Name	Type	Constraint	Description
id	INT	PRIMARY KEY	Unique identifier for each support ticket
customer_id	INT		Reference to the customer who created the ticket
created_at	VARCHAR(19)		Date and time when the ticket was created
resolved_at	VARCHAR(19)		Date and time when the ticket was resolved

▼ Sample Data Tables

support_tickets

id	customer_id	created_at	resolved_at
1	1	2023-12-21 05:42:00	2024-01-01 05:42:00
2	2	2023-07-08 14:22:00	NULL
3	3	2023-05-22 08:54:00	2023-06-17 08:54:00

Sample Output

```
average_response_time
444.00
```

Explanation

The sample output shows the average response time for resolved tickets to be 444.00, excluding Ticket `id` 2 since it has not been resolved.

Question - 3

SQL: Highest-Spending Customers per City

A retail company wants to identify the highest-spending customer in each city to target them for personalized marketing campaigns and loyalty programs. The goal is to generate a report highlighting the top customer by total spending in each city location. Row order does not matter.

The result should have the following columns: *customer_id / name / city / total_spending*.

- *customer_id* - Unique identifier for the customer.
- *name* - Name of the customer.
- *city* - The city where the customer is located.
- *total_spending* - The total spending of the customer, calculated by summing all order amounts for each customer, and should be converted to an integer by rounding down using an appropriate function, e.g., 1.99 rounds to 1.

Note:

- Only include the highest-spending customer(s) from each city based on their total spending. That is, if the maximum highest amount spent in a city is 100, include all customers in that city that spent 100.

▼ Schema

customers

Name	Type	Constraint	Description
id	INT	PRIMARY KEY	Unique identifier for a customer
name	VARCHAR(255)		Name of the customer
city	VARCHAR(255)		City where the customer is located

orders

Name	Type	Constraint	Description
id	INT	PRIMARY KEY	Unique identifier for an order
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	Reference to the customer
amount	DECIMAL(10,2)		Total amount of the order

▼ Sample Data Tables

customers

id	name	city
1	Customer 1	Los Angeles
2	Customer 2	Chicago
3	Customer 3	Chicago

orders

id	customer_id	amount
1	1	150.75
2	2	230.50
3	3	345.25

Sample Output

<i>customer_id</i>	<i>name</i>	<i>city</i>	<i>total_spending</i>
1	Customer 1	Los Angeles	150
3	Customer 3	Chicago	345

Explanation

The sample output shows the highest-spending customers in each city with their customer ID, name, city, and total spending. For "Los Angeles", "Customer 1" has the highest spending of 150. In "Chicago", "Customer 2" has a total spending of 230, but "Chicago" top spender is "Customer 3", with a total of 345.

Question - 4

SQL: E-commerce Product Request Report

An e-commerce platform maintains a database to track its products and customer requests for these products. The task is to generate a report that lists each available product's name and the total number of requests received for it.

The result should have the following columns: *product_name / total_requests*.

- *product_name* - the name of the available product.
- *total_requests* - the total number of requests received for the product.

The result should be sorted in descending order based on *total_requests*, and in case of a tie, by *product_name* in ascending order.

Note:

- Only products that are currently available should be included in the report.
- The *is_available* field in the products table indicates whether a product is available (1 for available, 0 for not available).

▼ Schema

products

<i>name</i>	<i>type</i>	<i>constraints</i>	<i>description</i>
<i>id</i>	INT	NOT NULL PRIMARY KEY	The identifier of the product
<i>name</i>	VARCHAR(255)		The name of the product
<i>category</i>	VARCHAR(255)		The category of the product
<i>is_available</i>	SMALLINT		The flag indicating if the product is available

requests

<i>name</i>	<i>type</i>	<i>constraints</i>	<i>description</i>
<i>product_id</i>	INT	FOREIGN KEY REFERENCES products(id)	The reference to the product
<i>client_email</i>	VARCHAR(255)		The email address of the client

▼ Sample Data Tables

products

id	name	category	is_available
1	PromoPro	beauty products	1
2	AdVantage	outdoor gear	1
3	MarketMagnet	sports equipment	1
5	AdBlitz	beauty products	0

requests

product_id	client_email
1	salgate1@fc2.com
1	lwycliff6@list-manage.com
1	ekimbleyf@scientificamerican.com
2	bgooro@spotify.com
2	vsamwayest@bbb.org
3	apappin0@yellowbook.com
3	ringreyb@businessinsider.com
3	mrysonm@istockphoto.com
5	ayushin1c@opera.com
5	bcoulston1q@hubpages.com

Sample Output

```
MarketMagnet 3
PromoPro    3
AdVantage   2
```

Explanation

The sample output lists the available products with their total request counts. 'PromoPro' and 'MarketMagnet' both have 3 requests, but 'MarketMagnet' appears first due to alphabetical ordering. 'AdVantage' has 2 requests and is listed after the others.

Question - 5

SQL: Active Campaign Engagement Report

A marketing company maintains a database to track its advertising campaigns and engagements. The task is to generate a report that includes the name of each active campaign, the total number of engagements, and the sum of views and clicks for those engagements.

The result should have the following columns: *campaign_name / total_engagements / total_views_and_clicks*.

- *campaign_name* - the name of the active campaign
- *total_engagements* - the number of engagements
- *total_views_and_clicks* - the combined number of views and clicks

The result should be sorted in ascending order by *campaign_name*.

Note:

- Only active campaigns should be included in the report.
- The *is_active* field in the campaigns table indicates whether a campaign is active (1 for active, 0 for inactive).

▼ Schema

campaigns

name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the campaign
name	VARCHAR(255)		The name of the campaign
is_active	SMALLINT		The flag indicates if the campaign is active

engagements

name	type	constraints	description
campaign_id	INT	FOREIGN KEY REFERENCES campaigns(id)	The reference to the campaign
views	INT		The total number of views of the engagement
clicks	INT		The total number of clicks of the engagement

▼ Sample Data Tables

campaigns

id	name	is_active
1	SummerSavings	1
2	FallFrenzy	1
3	WinterWonderland	0

engagements

campaign_id	views	clicks
1	100	10
1	150	20
2	200	30
2	250	40
3	300	50
1	120	15
2	180	25
3	220	35
1	130	18
2	210	28

Sample Output

```
FallFrenzy      3  963
SummerSavings  4  443
```

Explanation

The output includes only the active campaigns. For each active campaign, the total number of engagements is calculated by counting the number of records in the engagements table for that campaign. The number of views and clicks is calculated by summing (views + clicks) for each engagement related

to the campaign. The campaigns are then sorted by their names in ascending order.

Question - 6

SQL: Tax Report Summary

An online tax reporting application needs a summary report of account activity. It should list each account's email and the sum of reported amounts during 2023.

The result should have the following columns: *email / total_report_amount*.

- *email* - the email address of the account.
- *total_report_amount* - the sum of reported amounts submitted in 2023, rounded to two decimal places.

The results should be sorted in ascending order by *email*.

Note:

- Only reports submitted in the year 2023 should be included in the report.
- Ensure all decimal values are formatted to include trailing zeros if necessary (e.g., 5.00).

▼ Schema

accounts

name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the account
email	VARCHAR(255)		The email address of the account

reports

name	type	constraints	description
account_id	INT	FOREIGN KEY REFERENCES accounts(id)	The reference to the account
dt	VARCHAR(19)		The date and time of report
amount	DECIMAL(6, 2)		The reported amount

▼ Sample Data Tables

accounts

id	email
1	hratke0@disqus.com
2	lcaiger1@si.edu
3	gburkett2@vinaora.com

reports

account_id	dt	amount
1	2023-05-27 01:46:19	830.45
2	2023-01-15 09:23:21	2518.18
3	2023-05-08 01:44:41	4637.39
1	2023-06-30 15:02:03	3953.69

2	2023-12-05 04:39:31	3357.99
3	2023-02-03 09:41:00	1907.38
1	2022-12-30 04:05:57	1217.29
2	2024-01-24 14:18:07	2441.66
3	2024-01-05 23:19:31	3055.2
1	2023-05-26 01:54:24	2077.36

Sample Output

```
gburkett2@vinaora.com 6544.77
hratke0@disqus.com 6861.50
lcaiger1@si.edu 5876.17
```

Explanation

The sample output shows the total amount of reports for each account email in 2023. The amounts are summed and rounded to two decimal places. The output is sorted by the account email in ascending order.

Question - 7

SQL: Antivirus Device Scan Report

An antivirus software company maintains a database to track devices and the files scanned on each device. The task is to generate a report that lists each device's MAC address along with the total number of files scanned and the total number of infected files for that device.

The result should have the following columns: *mac_address* / *total_files_scanned* / *total_infected_files*.

- *mac_address* - the MAC address of the device.
- *total_files_scanned* - the total number of files scanned on the device.
- *total_infected_files* - the total number of infected files on the device.

The result should be sorted in ascending order by *mac_address*.

Note:

- The *is_infected* field in the scanned files table indicates whether a file is infected (1 for infected, 0 for not infected).

▼ Schema

devices

name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the device
mac_address	VARCHAR(255)		The MAC address of the device

scanned_files

name	type	constraints	description
device_id	INT	FOREIGN KEY REFERENCES devices(id)	The reference to the device
filename	VARCHAR(255)		The name of the file
is_infected	SMALLINT		The flag indicating if the file is infected

▼ Sample Data Tables

devices

id	mac_address
1	66-0F-84-41-B8-8E
2	A6-1A-2F-3A-7B-83
3	76-CD-24-48-F0-DD

scanned_files

device_id	filename	is_infected
1	File1.mp3	0
1	File2.xls	1
2	File3.doc	0
2	File4.ppt	1
2	File5.mp3	1
3	File6.xls	0
3	File7.doc	1
3	File8.ppt	0
3	File9.mp3	1
3	File10.xls	0

Sample Output

```
66-0F-84-41-B8-8E 2 1
A6-1A-2F-3A-7B-83 3 2
76-CD-24-48-F0-DD 5 2
```

Explanation

The report shows the MAC address of each device along with the total number of files scanned and the total number of infected files. For example, the device with MAC address '66-0F-84-41-B8-8E' has 2 files scanned, out of which 1 is infected.

Question - 8

SQL: Cryptocurrency Transactions Report

In the cryptocurrency market, a database engineer is tasked with generating a report for all cryptocurrency coins and their associated transactions. The report should include the name of each coin, the total amount of transactions, and the total number of transactions for each coin in the year 2023.

The result should have the following columns: *coin_name / total_transaction_amount / total_transactions*.

- *coin_name* - the name of the cryptocurrency coin.
- *total_transaction_amount* - the sum of transaction amounts in 2023, rounded to two decimal places.
- *total_transactions* - the total number of transactions in 2023

The result should be sorted in ascending order by *coin_name*.

Note:

- Only transactions that occurred in 2023 should be included in the report.
- Ensure decimal values are formatted to include trailing zeros if necessary, e.g., 5.00.

▼ Schema

coins			
name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the cryptocurrency coin
name	VARCHAR(255)		The name of the cryptocurrency coin

transactions

name	type	constraints	description
coin_id	INT	FOREIGN KEY REFERENCES coins(id)	The reference to the cryptocurrency coin
dt	VARCHAR(19)		The date and time of the transaction
amount	DECIMAL(5, 2)		The amount of the transaction

▼ Sample Data Tables

coins	
id	name
1	BitCash
2	Etherium
3	Litecoin

transactions

coin_id	dt	amount
1	2023-07-03 12:16:53	34.32
1	2023-12-08 12:14:58	47.59
2	2022-12-16 20:42:10	45.54
2	2023-11-05 09:27:11	53.3
3	2023-12-05 06:45:23	71.51
3	2023-01-19 01:43:25	97.18
3	2024-01-24 13:34:00	86.68
1	2023-05-07 05:30:06	25.6
2	2023-03-08 08:07:20	40.11
3	2023-08-13 10:44:54	87.54

Sample Output

```
BitCash 107.51 3
Etherium 93.41 2
Litecoin 256.23 3
```

Explanation

The sample output shows the total transaction amount and count for each coin in the year 2023. For 'BitCash', there are three transactions totaling 107.51. 'Etherium' has two transactions totaling 93.41. 'Litecoin' has three transactions totaling 256.23. All amounts are rounded to two decimal places.

Question - 9

SQL: Customer Domain Ownership Report

A domain hosting company maintains a database to manage its customers and the domains they own. The task is to generate a report that lists each customer's email address along with the total number of domains they own.

The result should have the following columns: *email / total_domains*.

- *email* - the email address of the customer.
- *total_domains* - the total number of domains owned by the customer.

The result should be sorted in ascending order by *email*.

▼ Schema

customers

name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the customer
email	VARCHAR(255)		The email address of the customer

domains

name	type	constraints	description
customer_id	INT	FOREIGN KEY REFERENCES customers(id)	The reference to the customer
name	VARCHAR(255)		The name of the domain

▼ Sample Data Tables

customers

id	email
1	ebayldon0@washingtonpost.com
2	agammade1@comcast.net
3	goloshkin2@reference.com
4	cantonescu3@earthlink.net
5	fparzis4@ow.ly
6	cpetroulis5@shutterfly.com
7	tbeels6@bbb.org
8	zmacturlough7@4shared.com
9	eshury8@skype.com
10	jfehnerns9@github.io

domains

customer_id	name
1	bfilipa.net

1	gsparshtoli.net
1	jhughsr.org
2	scopas8.net
2	cglison1u.org
3	tginiz.com
3	arubinowitsch2l.net
3	clockyear2m.org
4	sfinnigand.com
4	vborrelt.net

Sample Output

```
agammade1@comcast.net 2
cantonescu3@earthlink.net 2
ebayldon0@washingtonpost.com 3
goloshkin2@reference.com 3
```

Explanation

The output lists the email addresses of customers along with the total number of domains they own. For instance, the customer with the email 'ebayldon0@washingtonpost.com' owns 3 domains, while 'agammade1@comcast.net' owns 2 domains. The result is sorted by email addresses in ascending order.

Question - 10

SQL: E-commerce Wishlist Report

Generate a report from an e-commerce database that lists the product names and prices, along with the total number of times each is on a wishlist.

The result should have the following columns: *product_name / price / total_wishlist_count*.

- *product_name* - the name of the product.
- *price* - its price
- *total_wishlist_count* - the total number of times it appears in wishlists

The result should be sorted in ascending order by *product_name*.

Note:

- Only include products that are currently in stock.
- Ensure all decimal values are formatted to include trailing zeros if necessary, e.g., 5.00.

▼ Schema

products

name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the product
name	VARCHAR(255)		The name of the product
price	DECIMAL(6, 2)		The price of the product
in_stock	SMALLINT		1 indicates 'in stock', 0 indicates 'out of stock'

wishlists

name	type	constraints	description
product_id	INT	FOREIGN KEY REFERENCES products(id)	The reference to the product
customer_email	VARCHAR(255)		The email address of the customer

▼ Sample Data Tables

products

id	name	price	in_stock
1	TechGadget Pro X	324.24	1
2	LuxuryHome Decor Set	884.9	1
3	FitnessTracker Elite	698.59	0

wishlists

product_id	customer_email
1	user1@example.com
1	user2@example.com
2	user3@example.com
2	user4@example.com
2	user5@example.com
3	user6@example.com
1	user7@example.com
2	user8@example.com
1	user9@example.com
3	user10@example.com

Sample Output

```
LuxuryHome Decor Set 884.9 3
TechGadget Pro X      324.24 4
```

Explanation

The report includes only products that are in stock. 'TechGadget Pro X' is in stock and appears 4 times in wishlists, while 'LuxuryHome Decor Set' is also in stock and appears 3 times. 'FitnessTracker Elite' is not included in the report as it is not in stock.

Question - 11

SQL: Email Campaign Report

An email campaign tracking platform maintains data on various campaigns and their email statistics. The task is to generate a report that lists each campaign's name along with the total number of emails sent, emails opened, and emails not opened.

The result should have the following columns: *campaign_name / total_emails_sent / total_emails_opened / total_emails_not_opened*.

- *campaign_name* - the name of the email campaign.
- *total_emails_sent* - the total number of emails sent in the campaign.
- *total_emails_opened* - the total number of emails opened in the campaign.

- *total_emails_not_opened* - the total number of emails not opened.

The result should be sorted in ascending order by *campaign_name*.

Note:

- The number of emails not opened is calculated as the difference between the total emails sent and the emails opened.

▼ Schema

campaigns

name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the email campaign
name	VARCHAR(255)		The name of the email campaign

email_stats

name	type	constraints	description
campaign_id	INT	FOREIGN KEY REFERENCES campaigns(id)	The reference to the email campaign
emails_sent	INT		The number of emails sent in the email campaign
emails_opened	INT		The number of emails opened in the email campaign

▼ Sample Data Tables

campaigns

id	name
1	SummerSale2021
2	FallPromo
3	WinterWonderland

email_stats

campaign_id	emails_sent	emails_opened
1	1000	800
2	1500	1200
3	2000	1800
1	500	300
2	700	500
3	800	600
1	300	200
2	400	300
3	600	500
3	400	300

Sample Output

FallPromo	2600	2000	600
SummerSale2021	1800	1300	500
WinterWonderland	3800	3200	600

Explanation

The report shows each campaign's name along with the total emails sent, opened, and not opened. For example, the 'FallPromo' campaign had a total of 2600 emails sent, 2000 emails opened, and 600 emails not opened, calculated as the difference between sent and opened.

Question - 12

SQL: Auction Lot Offers Report

In an e-commerce auction platform, the task is to generate a report that provides insights into the bidding activities on various lots. The report should list each lot's name, the highest offer made for that lot, and the total number of offers received.

The result should have the following columns: *lot_name / highest_offer / total_offers*.

- *lot_name* - the name of the lot.
- *highest_offer* - the highest offer made for the lot.
- *total_offers* - the total number of offers received for the lot.

The result should be sorted in ascending order by *lot_name*.

Note:

- Ensure all decimal values are formatted to include trailing zeros if necessary (e.g., 5.00).

▼ Schema

lots			
name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the lot
name	VARCHAR(255)		The name of the lot

offers			
name	type	constraints	description
lot_id	INT	FOREIGN KEY REFERENCES lots(id)	The identifier of the lot for which the offer is made
amount	DECIMAL(6, 2)		The amount of the offer

▼ Sample Data Tables

lots	
id	name
1	Acacia parramattensis Tindale
2	Poa arctica R. Br. ssp. aperta (Scribn. & Merr.) Soreng
3	Calophyllum inophyllum L.

offers	
lot_id	amount
1	260.91
1	802.83

1	986.78
2	814.57
2	999.06
2	414.67
3	200.41
3	593.07
3	701.88
3	972.87

Sample Output

```
Acacia parramattensis Tindale           986.78 3
Calophyllum inophyllum L.                972.87 4
Poa arctica R. Br. ssp. aperta (Scribn. & Merr.) Soreng 999.06 3
```

Explanation

The sample output shows the highest offer and the total number of offers for each lot. For example, the lot 'Acacia parramattensis Tindale' received a maximum offer of 986.78 and a total of 3 offers. The decimal values are rounded to two decimal places.

Question - 13

SQL: Online Banking Transactions Report

A financial services company needs a report of account transactions for September 2022.

For each account, the report should have the following columns: *IBAN* / *min_transaction* / *max_transaction* / *avg_transaction* / *total_transactions*.

- *IBAN* - the account's International Bank Account Number
- *min_transaction* - the minimum transaction amount
- *max_transaction* - the maximum transaction amount
- *avg_transaction* - the average transaction amount
- *total_transactions* - the total number of transactions

The results should be sorted in ascending order by *IBAN* to facilitate easy reference.

Note:

- Only transactions from September 2022 should be included in the report.
- Ensure all decimal values are formatted to include trailing zeros if necessary, e.g., 5.00.

▼ Schema

accounts			
name	type	constraints	description
id	INT	NOT NULL PRIMARY KEY	The identifier of the account
iban	VARCHAR(255)		The IBAN of the account

transactions

name	type	constraints	description
account_id	INT	FOREIGN KEY REFERENCES accounts(id)	The identifier of the account related to the transaction

dt	CHAR(19)		The date and time of the transaction
amount	DECIMAL(5, 2)		The amount of the transaction

▼ Sample Data Tables

accounts

id	iban
1	BG40 RFFX 4898 53DD CZD6 KQ
2	PT42 5267 0592 8451 8001 2180 3
3	FR96 8758 8909 81LR DJ71 ERKN D56

transactions

account_id	dt	amount
1	2022-09-02 06:33:39	33.31
1	2022-09-20 08:14:39	31.77
1	2022-09-25 06:41:45	72.84
2	2022-09-04 22:28:12	35.26
2	2022-09-17 07:57:29	33.27
2	2022-09-27 22:30:36	70.78
3	2022-09-16 21:54:12	75.04
3	2022-09-19 18:27:39	71.19
3	2022-09-28 01:38:56	14.34
3	2022-08-30 01:35:31	69.19

Sample Output

```
BG40 RFFX 4898 53DD CZD6 KQ      31.77 72.84 45.97 3
PT42 5267 0592 8451 8001 2180 3  33.27 70.78 46.44 3
FR96 8758 8909 81LR DJ71 ERKN D56 14.34 75.04 53.52 3
```

Explanation

The sample output shows the IBAN for each account along with the minimum, maximum, and average transaction amounts for September 2022, rounded to two decimal places. It also includes the total number of transactions for each account during that month.

Question - 14

SQL: Top Wishlist Products Summary

In the competitive e-commerce landscape, understanding customer preferences and product popularity is crucial for inventory management and marketing strategies. A development team is creating a report that identifies the top 3 products in customers' wishlists. This report will highlight in-stock products, allowing immediate action to capitalize on customer interest.

The result should have the following columns: *name / price / total_wishes*.

- *name* - the name of the product
- *price* - the price of the product
- *total_wishes* - the total number of times the product appears in customers' wishlists

The results should be sorted in descending order by *total_wishes* to easily identify the most desired products, then in ascending order by *name*.

The result should be limited to the top 3 products.

Note:

- Only products that are in stock should be included.

▼ Schema

products

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the product
name	VARCHAR(255)		The name of the product
price	DECIMAL(6,2)		The price of the product
in_stock	SMALLINT		The flag indicating if the product is in stock

wishlists

Name	Type	Constraints	Description
product_id	INT	FOREIGN KEY(product_id => products.id)	The reference to the product
customer_email	VARCHAR(255)		The email address of the customer

▼ Sample Data Tables

products

id	name	price	in_stock
1	TechGadget Pro X	274.80	1
2	LuxuryHome Decor Set	262.84	1
3	FitnessTracker Elite	637.92	0
4	GourmetCookware Set	535.34	1
5	Fashionista Wardrobe Collection	525.44	1

wishlists

product_id	customer_email
1	crabbec@redcross.org
1	efindlow2@tinypic.com
1	jmachoste5@issuu.com
1	nselle@simplemachines.org
2	aonn1@ebay.co.uk
2	bbolton0@google.cn
2	ebockett3@storify.com

2	fzunguyg@symantec.com
2	slowried@cbsnews.com
3	jgately7@goo.ne.jp
3	ospearettj@bandcamp.com
3	rpanonsb@paypal.com
3	ydevauxh@toplist.cz
3	zbabbage9@imageshack.us
4	dpauleya@cnbc.com
4	jletterick4@dailymotion.com
4	khunnisett6@princeton.edu
4	rkernelf@uiuc.edu
5	blodin8@wikimedia.org
5	lyusupovi@nps.gov

Sample Output

```
+-----+-----+-----+
| name           | price | total_wishes |
+-----+-----+-----+
| LuxuryHome Decor Set|262.84|5
| GourmetCookware Set |535.34|4
| TechGadget Pro X   |274.80|4
+-----+-----+-----+
```

Question - 15

SQL: E-commerce Customer Purchases Report

In e-commerce, understanding customer purchasing behavior is essential for tailoring marketing strategies and improving customer service. A development team is creating a query for a report that details customer activity in March 2024. This report will provide insights into the total number and sum of purchases made by each customer, facilitating targeted engagement and promotional efforts.

The result should have the following columns: *email / total_purchases / total_purchase_amount*.

- *email* - the email address of the customer
- *total_purchases* - the total number of purchases made by the customer
- *total_purchase_amount* - the total sum of the purchases made by the customer, with two decimal places, including trailing zeros if necessary, e.g., 500.00

The results should be sorted in ascending order by *email*.

Note:

- Only purchases made in March 2024 should be included.

▼ Schema

customers			
Name	Type	Constraints	Description

id	INT	PRIMARY KEY	The identifier of the customer
email	VARCHAR(255)		The email address of the customer

purchases

Name	Type	Constraints	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
dt	VARCHAR(19)		The date and time of the purchase
amount	DECIMAL(6,2)		The amount of the purchase

▼ Sample Data Tables

customers

id	email
1	floggie0@newsvine.com
2	sgillbe1@ca.gov
3	jgohn2@elegantthemes.com

purchases

customer_id	dt	amount
2	2024-02-21 02:56:12	228.58
2	2024-02-23 09:32:47	972.41
1	2024-03-14 15:50:13	109.16
1	2024-03-17 00:31:44	11.49
1	2024-03-17 04:15:42	692.64
2	2024-03-01 04:35:09	589.74
2	2024-03-13 14:42:23	508.75
2	2024-03-17 07:57:36	933.91
2	2024-03-19 08:24:38	488.26
2	2024-03-31 23:30:54	55.07
3	2024-03-03 11:34:30	816.67
3	2024-03-08 23:46:07	672.93
3	2024-03-15 18:09:56	260.66
3	2024-03-20 15:18:11	321.07
3	2024-03-20 17:40:35	29.06
3	2024-03-20 23:41:39	314.85
3	2024-03-25 11:41:07	67.12
1	2024-04-05 03:05:10	417.78

2	2024-04-09 08:16:17	697.53
3	2024-04-02 07:56:48	156.27

Sample Output

```
+-----+-----+-----+
|email           |total_purchases|total_purchase_amount|
+-----+-----+-----+
|floggie0@newsvine.com| 3          |813.29
|ljgohn2@elegantthemes.com| 7          |2482.36
|sgillbe1@ca.gov      | 5          |2575.73
+-----+-----+-----+
```

Question - 16

SQL: Report on Applicants Pending Consular Service

In managing consular services, addressing delays caused by system issues is critical. A development team is creating a query for the system's dashboard to generate a report that identifies applicants who have not received service due to delays, emphasizing the extent of these delays. This report is essential for prioritizing these applicants and efficiently managing the backlog.

The result should have the following columns: *email / scheduled_appointment / days_of_delay*.

- *email* - the email address of the applicant
- *scheduled_appointment* - the date when the appointment was scheduled
- *days_of_delay* - the number of full days between the scheduled appointment date and the current date, April 10, 2024

The result should be sorted in ascending order by *scheduled_appointment*, then in ascending order by *email*.

Note:

- Only applicants who have not received service before the current date should be included in the result.
- The current date is fixed as April 10, 2024.

▼ Schema

applicants			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the applicant
email	VARCHAR(255)		The email address of the applicant

appointments			
Name	Type	Constraints	Description
applicant_id	INT	FOREIGN KEY(applicant_id => applicants.id)	The reference to the applicant
dt	VARCHAR(19)		The date of the scheduled appointment
is_received	BOOLEAN		The flag indicating if the service was received

▼ Sample Data Tables

applicants	
id	email

1	nkienzle0@spiegel.de
2	alaste1@bbc.co.uk
3	jjochanany2@ow.ly
4	bsenn3@salon.com
5	bwhittall4@nhs.uk

appointments

applicant_id	dt	is_received
1	2024-04-27	0
2	2024-04-01	0
3	2024-04-15	0
4	2024-03-27	0
5	2024-03-26	1

Sample Output

```
+-----+-----+-----+
| email           | scheduled_appointment | days_of_delay |
+-----+-----+-----+
| bsenn3@salon.com | 2024-03-27            |      14       |
| alaste1@bbc.co.uk | 2024-04-01            |      9        |
+-----+-----+-----+
```

Question - 17

SQL: Weekend Appointments for Consular Services Rescheduling

In consular services, efficiently managing appointment schedules is key to effective service delivery and resource utilization. A development team is creating a query for the system's dashboard to list all appointments scheduled on weekends (Saturday or Sunday). This initiative aims to identify appointments that need to be rescheduled to a working day.

The result should have the following columns: *email / scheduled_appointment*.

- *email* - the email address of the user who made the appointment
- *scheduled_appointment* - the name of the day when the appointment is scheduled

The results should be sorted in ascending order by *email*.

Note:

- Only appointments that are scheduled on weekends (Saturday or Sunday) should be included.

▼ Schema

applicants

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the applicant
email	VARCHAR(255)		The email address of the applicant

Name	Type	Constraints	Description
applicant_id	INT	FOREIGN KEY(applicant_id => applicants.id)	The reference to the applicant
dt	VARCHAR(19)		The date and time of the scheduled appointment

▼ Sample Data Tables

applicants

id	email
1	rastlatt0@instagram.com
2	gcarmody1@stanford.edu
3	mgreenset2@state.tx.us

appointments

applicant_id	dt
1	2024-05-26 01:36:43
2	2024-05-27 16:30:28
3	2024-05-18 19:28:52

Sample Output

```
+-----+-----+
| email           | scheduled_appointment |
+-----+-----+
|mgreenset2@state.tx.us | Saturday          |
|rastlatt0@instagram.com| Sunday           |
+-----+-----+
```

Question - 18

SQL: Active Domains Registration by Country with Totals

In managing a domain hosting service, understanding the distribution of active domain registrations across different countries is crucial for strategic planning and market analysis. A development team is creating a query for the system's dashboard to generate a report that lists countries with active domain registrations and the total number of domains per country.

The results should have the following columns: *country_name / total_domains*.

- *country_name* - the name of the country from where the domain was registered
- *total_domains* - the total number of active domains registered from that country

The results should be sorted in ascending order by *country_name*.

Note:

- Only active domains should be included in the result.

▼ Schema

countries

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the country
name	VARCHAR(255)		The name of the country

domains

Name	Type	Constraints	Description
country_id	INT	FOREIGN KEY(country_id => countries.id)	The reference to the country
name	VARCHAR(255)		The name of the domain
is_active	BOOLEAN		The status of the domain

▼ Sample Data Tables

countries

id	name
1	Azerbaijan
2	Colombia
3	China

domains

country_id	name	is_active
1	angelfire.com	1
1	free.fr	1
1	google.cn	1
1	nationalgeographic.com	1
1	ovh.net	1
1	surveymonkey.com	1
1	twitpic.com	1
2	ameblo.jp	1
2	berkeley.edu	1
2	multiply.com	1
2	redcross.org	1
2	sourceforge.net	1
3	hc360.com	1
3	liveinternet.ru	1
3	squidoo.com	1
3	technorati.com	1
3	webnode.com	1

3	yahoo.co.jp	1
1	1und1.de	0
1	qq.com	0

Sample Output

```
+-----+-----+
|country_name|total_domains|
+-----+-----+
|Azerbaijan | 7
|China       | 6
|Colombia   | 5
+-----+-----+
```

Question - 19

SQL: Domain Renewal Overview

In managing a domain hosting panel, tracking domain renewals is crucial for service continuity and customer satisfaction. A development team is creating a query to generate a report of all domains, including the current date, next renewal date, and days until renewal. This report helps domain owners plan their renewals, ensuring timely actions to maintain their online presence.

The result should have the following columns: *name* / *today_date* / *next_renewal_date* / *days_until_renewal*.

- *name* - the name of the domain
- *today_date* - the current date, April 10, 2024, in the format YYYY-MM-DD
- *next_renewal_date* - the date when the domain is scheduled for its next renewal
- *days_until_renewal* - the number of full days from the current date until the domain's next renewal date

The results should be sorted in ascending order by *days_until_renewal*, then in ascending order by *name*.

Note:

- The current date is fixed as April 10, 2024.

▼ Schema

domains

Name	Type	Constraints	Description
name	VARCHAR(255)		The name of the domain
next_renewal_date	VARCHAR(19)		The date when the domain is scheduled for next renewal

▼ Sample Data Tables

domains

name	next_renewal_date
wired.com	2024-06-14 00:10:12
blogger.com	2024-07-18 05:54:57
com.com	2024-07-21 02:57:25

Sample Output

```
+-----+-----+-----+
| name | today_date | next_renewal_date | days_until_renewal |
+-----+-----+-----+
| wired.com | 2024-04-10 | 2024-06-14 | 65 |
| blogger.com | 2024-04-10 | 2024-07-18 | 99 |
| com.com | 2024-04-10 | 2024-07-21 | 102 |
+-----+-----+-----+
```

Question - 20

SQL: User Transaction Details

In payment systems, analyzing transaction data provides insights into user behavior and system performance. A development team is creating a query for the system's dashboard to generate a report detailing user transactions for March 2024. This report will show the total number, minimum, maximum, and total amount of transactions for each user. These insights help tailor user engagement strategies and improve the user experience.

The results should have the following columns: *email / total_transactions / min_amount / max_amount / total_amount*.

- *email* - the email address of the user
- *total_transactions* - the total number of transactions made by the user during the specified period
- *min_amount* - the minimum transaction amount by the user during the specified period, with two decimal places, including trailing zeros if necessary, e.g., 500.00
- *max_amount* - the maximum transaction amount by the user during the specified period, with two decimal places, including trailing zeros if necessary, e.g., 500.00
- *total_amount* - the total amount of all transactions made by the user during the specified period, with two decimal places, including trailing zeros if necessary, e.g., 500.00

The results should be sorted in ascending order by *email*.

Note:

- Only transactions that occurred in March 2024 should be included.

▼ Schema

users			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the user
email	VARCHAR(255)		The email address of the user

transactions

Name	Type	Constraints	Description
user_id	INT	FOREIGN KEY(user_id => users.id)	The reference to the user
dt	VARCHAR(19)		The date and time of transaction
amount	DECIMAL(5,2)		The amount of transaction

▼ Sample Data Tables

users	
id	email
1	bblaszcynski0@devhub.com

2	dwokey1@chronoengine.com
3	fлерway2@wikipedia.org

transactions

user_id	dt	amount
1	2024-02-23 19:30:03	942.50
1	2024-03-07 09:01:15	855.22
1	2024-04-01 04:18:41	253.35
1	2024-04-07 02:40:58	886.88
2	2024-02-25 05:11:39	957.77
2	2024-03-06 03:00:40	413.39
2	2024-03-07 14:41:03	906.16
2	2024-03-10 00:58:13	116.59
2	2024-03-13 23:38:29	550.31
2	2024-03-22 03:07:46	196.23
2	2024-03-24 00:23:14	399.76
2	2024-03-25 12:28:18	398.07
2	2024-03-27 09:11:15	212.33
2	2024-04-09 06:33:26	97.85
3	2024-03-01 17:24:48	323.11
3	2024-03-05 10:16:06	673.23
3	2024-03-08 14:19:46	236.74
3	2024-03-23 15:37:47	234.87
3	2024-04-05 20:55:45	989.35
3	2024-04-07 05:26:35	369.20

Sample Output

```
+-----+-----+-----+-----+-----+
|email           |total_transactions|min_amount|max_amount|total_amount|
+-----+-----+-----+-----+-----+
|bbłaszczyński@devhub.com|1           |855.22    |855.22    |855.22    |
|dwokey1@chronoengine.com|8           |116.59    |906.16    |3192.84   |
|fлерway2@wikipedia.org|4           |234.87    |673.23    |1467.95   |
+-----+-----+-----+-----+-----+
```

Question - 21

SQL: Total Transactions and Sum for Each User

In the digital age, payment systems are crucial for online transactions. For businesses operating these systems, understanding user activity helps optimize services and enhance engagement. A development team is creating a report that details the total number and sum of transactions for each user. User transaction patterns will help identify highly active users and those who need additional engagement.

The result should have the following columns: *email* / *total_transactions* / *total_amount*.

- *email* - the email address of the user
- *total_transactions* - the total number of transactions made by the user
- *total_amount* - the total sum of the transactions made by the user, with two decimal places, including trailing zeros if necessary, e.g., 500.00

The results should be sorted in ascending order by *email*.

Note:

- Only transactions that occurred in the year 2023 should be included.

▼ Schema

users			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the user
email	VARCHAR(255)		The email address of the user

transactions

Name	Type	Constraints	Description
user_id	INT	FOREIGN KEY(user_id => users.id)	The reference to the user
dt	VARCHAR(19)		The date and time of transaction
amount	DECIMAL(5,2)		The amount of transaction

▼ Sample Data Tables

users	
id	email
1	lvasilevich0@google.co.uk
2	hscholey1@sina.com.cn
3	mmcjury2@hibu.com

transactions

user_id	dt	amount
3	2022-12-05 00:16:56	162.11
1	2023-05-20 03:20:58	81.58
1	2023-06-08 19:24:02	52.46
1	2023-06-27 21:16:07	447.59
1	2023-07-20 08:19:32	136.68
1	2023-12-11 17:08:05	852.55

1	2023-12-15 04:45:54	77.11
1	2023-12-22 00:46:34	670.71
1	2023-12-29 12:43:23	948.46
2	2023-01-04 00:51:46	793.50
2	2023-04-07 16:29:14	762.52
2	2023-06-17 17:42:50	527.18
2	2023-10-10 11:16:51	733.47
2	2023-10-18 23:32:00	920.14
3	2023-03-27 18:31:41	408.13
3	2023-04-08 09:57:55	817.88
3	2023-05-18 09:47:14	916.98
3	2023-09-14 14:00:54	53.30
3	2023-09-30 01:34:01	589.37
3	2024-01-27 15:13:58	666.37

Sample Output

```
+-----+-----+-----+
|email           |total_transactions|total_amount|
+-----+-----+-----+
|hscholey1@sina.com.cn|5          |3736.81      |
|lvasilevich0@google.co.uk|8          |3267.14      |
|mcmcjury2@hibu.com|5          |2785.66      |
+-----+-----+-----+
```

Question - 22

SQL: Top Cryptocurrencies by Average Transaction Amount

In the rapidly changing cryptocurrency market, understanding average transaction amounts provides valuable insights into market behavior and investor sentiment. A development team is creating a report for a financial dashboard that highlights the top 3 cryptocurrencies based on the average transaction amount for all transactions in 2023. Cryptocurrencies with higher average transaction values indicate higher investor confidence or larger institutional involvement.

The results should include the following columns: *name*, *avg_transaction_amount*.

- *name* - the name of the cryptocurrency
- *avg_transaction_amount* - the average amount of transactions for the cryptocurrency during the specified period, rounded to two decimal places, including trailing zeros if necessary, e.g., 5.00

The result should be sorted in ascending order by *avg_transaction_amount*.

The results should be limited to the top 3 coins.

Note:

- Only transactions that occurred in 2023 should be included.

▼ Schema

coins			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of cryptocurrency
name	VARCHAR(255)		The name of the cryptocurrency

transactions

Name	Type	Constraints	Description
coin_id	INT	FOREIGN KEY(coin_id => coins.id)	The reference to the cryptocurrency
dt	VARCHAR(19)		The date and time of transaction
amount	DECIMAL(5,2)		The amount of transaction

▼ Sample Data Tables

coins	
id	name
1	BitCash
2	Etherium
3	Litecoin
4	Ripple
5	Dogecoin

transactions

coin_id	dt	amount
1	2022-12-09 19:40:17	60.91
1	2023-01-02 09:35:37	76.35
1	2023-03-21 09:34:39	23.11
1	2023-08-11 03:43:27	80.20
1	2023-10-21 19:42:46	29.59
2	2023-07-08 19:47:20	69.49
2	2023-09-22 14:23:40	23.13
3	2023-01-08 10:22:10	72.45
3	2023-01-28 00:54:51	98.72
3	2023-02-24 00:13:32	70.36
3	2023-05-16 15:13:19	93.59
4	2023-05-24 13:43:44	9.34
4	2023-07-25 14:59:09	78.52

5	2023-01-20 15:49:38	81.66
5	2023-08-21 17:19:45	94.89
5	2023-10-25 00:44:42	64.40
5	2023-11-30 02:38:47	86.84
5	2023-12-31 03:26:39	58.99
2	2024-01-21 10:25:26	29.36
5	2024-01-08 03:09:00	95.25

Sample Output

```
+-----+-----+
| name      | avg_transaction_amount |
+-----+-----+
|Ripple    | 43.93                |
| Ethereum | 46.31                |
| BitCash   | 52.31                |
+-----+-----+
```

Question - 23

SQL: Cryptocurrency Transactions Summary Report

In the fast-changing world of cryptocurrency, tracking transaction metrics is crucial for investors and financial analysts to understand market dynamics. A development team is creating a report that summarizes transactions in March 2024 for each cryptocurrency, providing users with critical insights into transaction activity to help them make informed investment decisions.

The results should include the following columns: *name*, *total_transactions*, *min_amount*, *max_amount*, *avg_amount*.

- *name* - the name of the cryptocurrency
- *total_transactions* - the total number of transactions for the cryptocurrency during the specified period
- *min_amount* - the minimum transaction amount for the cryptocurrency during the specified period, with two decimal places, including trailing zeros if necessary, e.g., 500.00
- *max_amount* - the maximum transaction amount for the cryptocurrency during the specified period, with two decimal places, including trailing zeros if necessary, e.g., 500.00
- *avg_amount* - the average transaction amount for the cryptocurrency during the specified period, rounded to two decimal places, including trailing zeros if necessary, e.g., 5.00

The results should be sorted in descending order by *total_transactions* to easily identify the most actively traded coins, then in ascending order by *name*.

Note:

- Only transactions that occurred in March 2024 should be included.

▼ Schema

coins

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of cryptocurrency coin
name	VARCHAR(255)		The name of the cryptocurrency coin

transactions

Name	Type	Constraints	Description

coin_id	INT	FOREIGN KEY(coin_id => coins.id)	The reference to the cryptocurrency coin
dt	VARCHAR(19)		The date and time of transaction
amount	DECIMAL(5,2)		The amount of transaction

▼ Sample Data Tables

coins

id	name
1	BitCash
2	Etherium
3	Litecoin

transactions

coin_id	dt	amount
1	2024-02-24 14:56:31	69.38
2	2024-02-24 17:23:54	46.79
1	2024-03-07 15:07:57	73.45
1	2024-03-13 00:47:18	2.10
1	2024-03-13 06:21:06	5.68
1	2024-03-14 15:06:59	25.32
1	2024-03-20 02:10:37	72.85
2	2024-03-09 12:06:47	67.79
2	2024-03-16 06:17:14	82.87
2	2024-03-24 11:11:23	5.96
2	2024-03-26 14:36:34	21.80
3	2024-03-20 08:28:56	5.07
3	2024-03-23 06:45:13	11.85
3	2024-03-27 02:40:23	34.25
3	2024-03-28 07:37:50	81.12
3	2024-03-29 07:34:32	19.06
1	2024-04-02 08:49:17	13.04
2	2024-04-04 17:29:13	4.74
2	2024-04-04 20:17:36	94.42
2	2024-04-07 01:47:46	64.76

Sample Output

name	total_transactions	min_amount	max_amount	avg_amount
BitCash	5	12.10	73.45	35.88
Litecoin	5	15.07	81.12	30.27
Etherium	4	15.96	82.87	44.61

Question - 24

SQL: Antivirus Suspicious File Extensions Report

In cybersecurity, knowing the types of file extensions of suspicious files scanned by antivirus software is essential to understanding digital threats. A development team is creating a report for an antivirus dashboard that shows the top 5 suspicious file extensions for March 2024. It aims to give clients insights into common threats so they can protect their systems effectively.

The results should include the following columns: *extension*, *total_suspicious_files*.

- *extension* - the extension of the suspicious file
- *total_suspicious_files* - the number of times files with this extension were flagged as suspicious

The results should be sorted in descending order by *total_suspicious_files* and then in ascending order by *extension*.

The results should be limited to the top 5 extensions.

Note:

- Only files that were scanned and flagged as suspicious in March 2024 should be included.

▼ Schema

suspicious_files

Name	Type	Constraints	Description
filename	VARCHAR(255)		The name of the file
extension	VARCHAR(255)		The extension of the file
scan_dt	VARCHAR(19)		The date and time when the file was scanned
is_suspicious	BOOLEAN		The flag indicating whether the file was flagged as suspicious

▼ Sample Data Tables

suspicious_files

filename	extension	scan_dt	is_suspicious
Mauris.pdf	.pdf	2024-04-05 23:55:27	1
Augue.xls	.xls	2024-02-28 18:11:28	1
Sapien.avi	.avi	2024-03-30 12:24:10	1
Pulvinar.doc	.doc	2024-03-08 22:00:41	1
TemporConvallisNulla.gif	.gif	2024-03-29 21:32:41	1
InFaucibus.mp3	.mp3	2024-03-20 14:18:32	1

EleifendPedeLibero.ppt	.ppt	2024-03-05 04:47:56	1
VestibulumAntelpsum.ppt	.ppt	2024-03-05 17:34:34	1
IntegerPede.ppt	.ppt	2024-03-12 17:11:28	1
VenenatisNon.tiff	.tiff	2024-03-20 18:04:47	1
IaculisDiam.xls	.xls	2024-03-01 05:18:03	1
QuisqueArcuLibero.xls	.xls	2024-03-09 09:00:32	1
MaurisSit.png	.png	2024-04-03 23:20:03	0
SitAmetSem.mp3	.mp3	2024-02-23 22:06:43	0
Nisi.mp3	.mp3	2024-02-29 09:40:45	0
Magna.tiff	.tiff	2024-02-27 00:25:16	0
EratVestibulum.gif	.gif	2024-03-30 04:19:52	0
Neque.jpeg	.jpeg	2024-03-07 07:11:26	0
VolutpatQuam.ppt	.ppt	2024-03-23 04:33:43	0
NonQuam.xls	.xls	2024-03-10 19:12:29	0

Sample Output

```
+-----+-----+
|extension|total_suspicious_files|
+-----+-----+
|.ppt    | 3
|.xls    | 2
|.avi    | 1
|.doc    | 1
|.gif    | 1
+-----+-----+
```

Question - 25

SQL: Antivirus Scanned Devices Report

In cybersecurity, monitoring the number of devices scanned by antivirus software is essential for digital security. A development team is creating a feature for an antivirus dashboard to provide clients with a report detailing the total number of devices scanned. This report will focus on devices actively scanned in March 2024, giving clients insights into their antivirus activity.

The results should include the following columns: *email*, *total_scanned_devices*.

- *email* - the email address of the client
- *total_scanned_devices* - the total number of devices that have been scanned for each client during the specified period

The results should be sorted in ascending order by *email*.

Note:

- Only devices that were scanned in March 2024 should be included.

▼ Schema

clients

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the client
email	VARCHAR(255)		The email address of the client

devices

Name	Type	Constraints	Description
client_id	INT	FOREIGN KEY(client_id => clients.id)	The reference to the client
mac_address	VARCHAR(255)		The MAC address of the device
is_scanned	BOOLEAN		The flag indicating whether the device was scanned
scheduled_scan_dt	VARCHAR(19)		The date and time the device is scheduled to be scanned

▼ Sample Data Tables

clients

id	email
1	cbracegirdle0@irs.gov
2	gwickardt1@msu.edu
3	mpaulon2@edublogs.org

devices

client_id	mac_address	is_scanned	scheduled_scan_dt
1	87-EF-C7-BD-DF-A2	1	2024-02-28 08:13:55
1	37-FE-45-2B-9D-2A	1	2024-03-07 00:00:05
1	13-82-F2-48-88-FD	1	2024-03-11 06:53:47
1	0D-56-2A-B2-33-EF	1	2024-03-30 19:41:31
1	17-C3-E3-2E-37-7E	1	2024-04-06 13:14:06
1	93-25-74-C5-07-32	0	2024-02-21 16:23:31
2	64-E1-5B-12-AC-F9	0	2024-02-22 10:40:23
2	0F-66-56-E2-B0-3A	0	2024-02-23 11:03:58
2	40-F4-40-12-C8-A5	0	2024-02-24 18:01:50
2	B0-2B-99-84-68-7C	0	2024-02-25 21:31:06
2	07-1F-BD-16-AC-23	0	2024-02-29 21:45:22
2	CA-79-F4-B4-9E-69	0	2024-03-29 00:05:10
3	71-EB-63-A2-3C-AF	1	2024-03-03 07:50:20
3	0B-40-DF-14-53-0F	1	2024-03-21 11:10:52
3	0A-77-ED-ED-50-28	1	2024-04-03 04:11:25

3	A4-79-0C-6D-B8-4C	1	2024-04-04 13:50:43
3	44-A5-56-27-C8-70	0	2024-03-28 21:15:26
3	93-64-42-51-62-6F	0	2024-03-31 20:26:01
3	87-5E-B3-51-38-2D	0	2024-04-05 10:51:00
3	DE-F2-F6-AD-76-4A	0	2024-04-08 23:15:03

Sample Output

```
+-----+-----+
|email      |total_scanned_devices|
+-----+-----+
|cbracegirdle0@irs.gov| 3
|mpaulon2@edublogs.org| 2
+-----+-----+
```

Question - 26

SQL: Resource Usage Report for Online Hosting Panel

In web hosting, monitoring resource utilization is essential for providers and customers. A development team is adding a new report that shows customers whose websites use more than 50% of any resource (CPU, memory, or disk). It will provide a detailed overview of their average resource consumption to enable more effective hosting management.

The results should include the following columns: *email*, *average_cpu_usage*, *average_memory_usage*, *average_disk_usage*.

- *email* - the email address of the customer
- *average_cpu_usage* - the average CPU usage across all sites for each customer, rounded to two decimal places, including trailing zeros if necessary, e.g., 500.00
- *average_memory_usage* - the average memory usage across all sites for each customer, rounded to two decimal places, including trailing zeros if necessary, e.g., 500.00
- *average_disk_usage* - the average disk usage across all sites for each customer, rounded to two decimal places, including trailing zeros if necessary, e.g., 500.00

The results should be sorted in ascending order by *email*.

Note:

- Only customers with at least one average value of a metric (CPU, memory, or disk usage) greater than 50% should be included.

▼ Schema

customers			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the customer
email	VARCHAR(255)		The email address of the customer

site_metrics

Name	Type	Constraints	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
cpu_usage	DECIMAL(5,2)		The CPU usage percentage
memory_usage	DECIMAL(5,2)		The memory usage percentage

disk_usage	DECIMAL(5,2)		The disk usage percentage
------------	--------------	--	---------------------------

▼ Sample Data Tables

customers

id	email
1	lrathke0@usa.gov
2	epearsall1@fema.gov
3	sivasechko2@cisco.com

site_metrics

customer_id	cpu_usage	memory_usage	disk_usage
1	31.53	80.84	1.51
1	12.54	26.47	47.74
1	12.34	46.24	34.43
1	26.64	84.98	17.56
2	80.45	50.05	10.63
2	40.14	86.67	15.98
2	30.14	34.38	17.67
2	1.11	83.44	2.95
3	30.60	18.60	28.02
3	41.64	33.64	5.20
3	31.88	7.37	91.14
3	43.20	9.56	40.40
3	2.33	34.29	18.65
3	11.50	32.89	71.39
3	39.57	4.49	48.05
3	25.06	23.77	33.00
3	32.81	1.59	25.85
3	48.38	79.21	8.31
3	11.62	26.75	71.71
3	54.43	6.48	4.86

Sample Output

```
+-----+-----+-----+-----+
| email | average_cpu_usage | average_memory_usage | average_disk_usage |
+-----+-----+-----+-----+
| epearsall1@fema.gov | 37.96 | 63.64 | 11.81 |
```

Question - 27

SQL: Dashboard Report for Online Hosting Customers Panel

In the dynamic world of web hosting, providing customers with a clear overview of their assets is key to enhancing user experience. A development team must create a feature for an online hosting customers panel that generates a dashboard report of all customers and the number of active websites they own.

The results should include the following columns: *email*, *total_active_sites*.

- *email* - the email address of the customer
- *total_active_sites* - the total number of active websites for each customer

The results should be sorted in ascending order by *email*.

Note:

- Only active websites should be included.

▼ Schema

customers			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the customer
email	VARCHAR(255)		The email address of the customer

sites			
Name	Type	Constraints	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
url	VARCHAR(255)		The URL of the website
is_active	BOOLEAN		The status of the website

▼ Sample Data Tables

customers	
id	email
1	dchristofol0@slashdot.org
2	mbillanie1@japanpost.jp
3	hmainz2@utexas.edu

sites		
customer_id	url	is_active
1	https://trellian.com	1
1	https://www.google.de	1

1	https://merriam-webster.com	1
1	https://wordpress.com	1
1	https://nsw.gov.au	1
1	https://www.barnesandnoble.com	1
1	https://www.yahoo.com	1
2	https://cloudflare.com	0
2	https://www.is.gd	1
2	https://www.unesco.org	1
3	https://www.sina.com.cn	0
3	https://xinhuanet.com	1
3	https://cyberchimps.com	1
3	https://ask.com	1
3	https://businessinsider.com	1
3	https://www.dailymail.co.uk	1
3	https://www.guardian.co.uk	1
3	https://www.microsoft.com	1
3	https://www.gizmodo.com	1
3	https://www.163.com	1

Sample Output

```
+-----+-----+
|email           |total_active_sites|
+-----+-----+
|dcristofol0@slashdot.org|7   |
|hmainz2@utexas.edu    |9   |
|mbillanie1@japanpost.jp|2   |
+-----+-----+
```

Question - 28

SQL: Average Income Report in Online Tax Application

A development team is enhancing an online tax application to provide users with actionable financial insights. A new feature under development aims to generate a report showcasing the accounts with the highest average income based on income received within a specific timeframe. The immediate focus is on income recorded in the first quarter of 2024.

The result should include the following columns: *iban*, *average_income*, *total_income*.

- *iban* - the IBAN of the account
- *average_income* - the average income for each account during the specified period, rounded to two decimal places, including trailing zeros if necessary, e.g., 500.00
- *total_income* - the total amount of income for each account during the specified period, with two decimal places, including trailing zeros if necessary, e.g., 500.00

The results should first be sorted in descending order by *average_income*, then in ascending order by *iban*.

The results should identify the top 3 accounts by their average income during this period.

Note:

- Only income recorded in the first quarter of 2024 (from January 1, 2024, to March 31, 2024) should be included.

▼ Schema

accounts			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the account
iban	VARCHAR(255)		The IBAN of the account

income

Name	Type	Constraints	Description
account_id	INT	FOREIGN KEY(account_id => accounts.id)	The reference to the account
dt	VARCHAR(19)		The date and time of income
amount	DECIMAL(6,2)		The income amount

▼ Sample Data Tables

accounts

id	iban
1	SK39 8924 2092 2997 1101 4161
2	PL28 9141 8610 8442 2367 7521 0000
3	CH93 8418 0F7G KQK4 NEHF Q
4	GT41 TBM8 DPFH MTNS BVW5 D4CX VIRR
5	IT27 Y015 0159 036T W7E5 I6ZD EQZ

income

account_id	dt	amount
1	2024-01-17 16:43:20	4061.53
1	2024-02-28 05:30:15	4488.11
1	2024-04-07 05:41:27	4001.91
2	2023-12-21 07:38:45	4313.69
2	2024-01-08 04:48:45	3640.82
2	2024-01-20 17:31:20	3385.15
3	2024-01-06 23:18:30	2347.15
3	2024-03-08 12:53:20	3814.86

3	2024-04-01 21:18:16	2764.27
4	2024-01-02 23:52:06	3526.08
4	2024-02-04 12:32:28	2221.91
4	2024-02-11 19:44:53	4197.07
4	2024-03-06 06:28:34	1357.44
4	2024-03-16 16:13:49	1854.52
5	2023-12-31 22:08:57	2819.54
5	2024-01-14 18:03:47	2641.20
5	2024-01-23 07:50:22	3692.56
5	2024-02-28 23:43:28	1999.09
5	2024-03-20 10:29:44	1670.18
5	2024-03-27 11:12:04	1193.15

Sample Output

```
+-----+-----+-----+
|iban           |average_income|total_income |
+-----+-----+-----+
|SK39 8924 2092 2997 1101 4161    |4274.82      |8549.64     |
|PL28 9141 8610 8442 2367 7521 0000|3512.99      |7025.97     |
|CH93 8418 0F7G KQK4 NEHF Q       |3081.01      |6162.01     |
+-----+-----+-----+
```

Question - 29

SQL: Tax Calculation for Online Tax Application

In the world of finance, tax calculation is a fundamental operation. A development team is working on an online tax application designed to simplify tax calculations for users. One of the key features they aim to implement is a tax calculation mechanism. This mechanism will calculate a fixed 20% tax on the total income for each account in the system.

The results should have the following columns: *iban* / *total_income* / *tax_rate* / *calculated_tax*.

- *iban* - the IBAN of the account
- *total_income* - the total income for the account, with two decimal places, including trailing zeros if necessary, e.g., 500.00
- *tax_rate* - the tax rate applied, which is a fixed text 20%
- *calculated_tax* - the tax calculated, rounded to two decimal places, including trailing zeros if necessary, e.g., 500.00

The results should be sorted in ascending order by *iban*.

Note:

- Only income in 2023 should be included.

▼ Schema

accounts			
Name	Type	Constraints	Description

id	INT	PRIMARY KEY	The identifier of the account
iban	VARCHAR(255)		The IBAN of the account

income

Name	Type	Constraints	Description
account_id	INT	FOREIGN KEY(account_id => accounts.id)	The reference to the account
dt	VARCHAR(19)		The date and time of income
amount	DECIMAL(6,2)		The income amount

▼ Sample Data Tables

accounts

id	iban
1	FR55 4477 6154 73ND TN3F HMOU T36
2	DK46 1272 1831 2573 01
3	RS53 5237 5794 6016 5411 43

income

account_id	dt	amount
1	2022-12-31 10:03:42	2779.19
1	2023-02-04 08:50:14	1777.68
1	2023-02-13 04:22:07	1954.81
1	2023-03-04 14:46:04	1547.79
1	2023-05-23 15:42:13	1208.49
1	2023-05-24 23:24:07	1521.72
1	2023-07-28 11:01:46	1792.75
1	2023-12-07 14:19:09	2374.25
1	2024-01-27 05:55:36	2803.39
2	2022-12-03 18:04:34	1826.65
2	2023-02-17 00:59:57	3074.11
2	2023-03-01 08:17:15	1007.30
2	2023-08-19 09:16:41	4515.04
2	2024-01-08 04:14:22	3321.78
2	2024-01-10 15:16:28	2033.87
3	2023-05-09 07:28:27	3158.66
3	2023-05-22 04:39:34	3851.20
3	2023-07-21 19:51:14	4152.29

3	2023-10-05 05:42:49	4722.20
3	2023-11-11 02:42:59	1592.16

Sample Output

```
+-----+-----+-----+
|iban           |total_income|tax_rate|calculated_tax|
+-----+-----+-----+
|DK46 1272 1831 2573 01      |8596.45    |20%     |1719.29    |
|FR55 4477 6154 73ND TN3F HMOU T36|12177.49   |20%     |2435.50    |
|RS53 5237 5794 6016 5411 43      |17476.51    |20%     |3495.30    |
+-----+-----+-----+
```

Question - 30

SQL: Monthly Budget Report for Online Budgeting Application

A development team is working on an online budgeting application and needs to implement a feature that generates a monthly report for each customer. The report should provide insights into the customers' total expenses and total income, allowing them to better manage their finances.

The results should have the following columns: *email / total_expenses / total_income*.

- *email* - the email address of the customer
- *total_expenses* - the sum of all expenses recorded for the customer, showing two decimal places, for example, 500.00
- *total_income* - the sum of income recorded for the customer, showing two decimal places, for example, 500.00

The results should be sorted in ascending order by *email*.

Note:

- Only expenses and incomes that were recorded in March 2024 should be included in the results.

▼ Schema

customers			
Name	Type	Constraint	Description
id	INT	PRIMARY KEY	The identifier of the customer
email	VARCHAR(255)		The email address of the customer

expenses			
Name	Type	Constraint	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
dt	VARCHAR(19)		The date and time of expense
amount	DECIMAL(6,2)		The expense amount

income			
Name	Type	Constraint	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
dt	VARCHAR(19)		The date and time of income
amount	DECIMAL(6,2)		The income amount

▼ Sample Data Tables

customers

id	email
1	otoohey0@elpais.com
2	egrebbin1@state.gov
3	arides2@sohu.com

expenses

customer_id	dt	amount
1	2024-02-21 22:12:12	90.41
1	2024-02-27 06:48:37	792.88
1	2024-03-10 05:19:43	442.01
1	2024-03-11 19:48:25	327.35
1	2024-03-24 22:03:06	639.62
1	2024-03-29 00:37:46	150.12
1	2024-04-02 03:36:50	257.67
2	2024-02-21 06:11:26	400.22
2	2024-03-11 15:34:19	298.41
2	2024-03-25 04:36:27	376.87
2	2024-03-29 19:05:51	530.07
2	2024-03-30 07:07:28	287.84
2	2024-04-02 15:44:22	868.03
3	2024-03-01 16:02:47	33.30
3	2024-03-06 11:53:42	838.51
3	2024-03-20 23:34:48	968.08
3	2024-03-21 21:18:08	35.36
3	2024-03-30 06:51:13	956.12
3	2024-03-31 10:11:56	896.32
3	2024-03-31 22:36:57	740.94

income

customer_id	dt	amount
1	2024-02-20 21:00:55	366.66
1	2024-03-11 03:25:04	769.38

1	2024-03-15 00:49:53	84.10
1	2024-03-21 18:32:51	839.48
1	2024-03-29 15:34:13	333.97
1	2024-04-01 00:34:24	253.13
1	2024-04-02 11:13:49	263.56
2	2024-02-20 15:03:26	822.75
2	2024-02-26 14:57:39	277.23
2	2024-03-19 09:24:47	24.08
2	2024-03-20 15:54:24	988.34
2	2024-04-02 08:28:38	990.54
3	2024-02-21 10:23:33	430.82
3	2024-02-29 08:25:32	482.85
3	2024-03-01 05:10:42	962.60
3	2024-03-04 08:27:34	30.21
3	2024-03-19 12:12:01	80.00
3	2024-03-21 00:32:10	674.76
3	2024-03-23 14:14:32	863.79
3	2024-04-09 13:37:07	51.42

Sample Output

```
+-----+-----+-----+
|email      |total_expenses|total_income|
+-----+-----+-----+
|arides2@sohu.com|4468.63      |2611.36      |
|legrebbin1@state.gov|1493.19      |1012.42      |
|lotoohhey0@elpais.com|1559.10      |2026.93      |
+-----+-----+-----+
```

Question - 31

SQL: Balance Report for Online Budgeting Application

In personal finance, knowing one's income versus expenses is crucial. A development team is creating an online budgeting application to provide insights into financial health. One key feature is a balance report, showing the difference between each customer's total income and total expenses.

The results should have the following columns: *email* / *balance*.

- *email* - the email address of the customer
- *balance* - the difference between the total income and total expenses for the customer, showing two decimal places, for example, 500.00

The results should be sorted in ascending order by *email*.

Note:

- Only customers that have a negative balance should be included.

▼ Schema

customers

Name	Type	Constraint	Description
id	INT	PRIMARY KEY	The identifier of the customer
email	VARCHAR(255)		The email address of the customer

expenses

Name	Type	Constraint	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
amount	DECIMAL(6,2)		The expense amount

income

Name	Type	Constraint	Description
customer_id	INT	FOREIGN KEY(customer_id => customers.id)	The reference to the customer
amount	DECIMAL(6,2)		The income amount

▼ Sample Data Tables

customers

id	email
1	dtollmache0@typepad.com
2	eclutterbuck1@baidu.com
3	mdensun2@ustream.tv

expenses

customer_id	amount
1	136.18
1	323.28
1	383.37
1	505.41
1	841.21
2	5.23
2	408.33
2	489.45
2	545.40
2	591.43
2	706.13

2	716.82
2	761.75
2	796.30
3	152.26
3	211.30
3	447.57
3	685.03
3	966.89
3	967.30

income

customer_id	amount
1	39.44
1	49.49
1	292.19
1	419.36
1	529.26
1	695.43
1	763.72
1	797.92
1	833.34
2	139.42
2	422.18
2	506.59
2	566.00
2	697.92
2	938.51
3	304.66
3	345.03
3	371.86
3	371.88
3	552.08

Sample Output

```
+-----+-----+
|email |balance |
+-----+-----+
|eclutterbuck1@baidu.com|-1750.22|
|mdensun2@ustream.tv |-1484.84|
+-----+-----+
```

Question - 32

SQL: Monthly Sales Report

The sales department of a company wants to generate a monthly sales report for the first quarter of the year to track the performance of each product. The goal is to identify the top-selling products based on the total sales amount.

The result should have the following columns: *name / month / total_sales*.

- *name* - the name of the product.
- *month* - the month name of the sale (e.g., "January", "February", "March").
- *total_sales* - the total sales amount for the product in a specific month, rounded to two decimal places, including trailing zeros if necessary (e.g., 5.00).

The result should be sorted first by *month* in ascending calendar order, then by *total_sales* in descending order within each month.

Note:

- Consider only transactions that occurred in the first quarter of 2024.

▼ Schema

products

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the product
name	VARCHAR(255)		The name of the product

sales

Name	Type	Constraints	Description
product_id	INT	FOREIGN KEY(product_id => products.id)	The reference to the product
dt	VARCHAR(19)		The date and time of sale
amount	DECIMAL(7,2)		The amount of the sale

▼ Sample Data Tables

products

id	name
1	Luxury Gold Watch
2	Smartphone Holder Stand
3	Stainless Steel Water Bottle

sales

product_id	dt	amount
1	2024-01-13 17:12:22	7008.16
1	2024-01-03 03:15:27	6191.64
1	2024-01-22 18:29:09	4527.86
1	2024-01-26 19:38:53	7828.36
1	2024-02-17 09:27:13	5273.16
1	2024-02-11 09:51:24	3364.73
1	2024-02-22 23:53:15	8584.33
2	2024-01-28 11:33:58	3710.06
2	2024-01-25 14:47:25	5221.02
2	2024-01-21 07:58:53	2525.72
2	2024-03-15 14:16:18	8158.08
2	2024-03-12 17:02:01	6760.77
3	2024-01-13 19:27:51	1942.79
3	2024-02-15 08:04:40	9186.38
3	2024-03-06 08:02:37	5821.97
3	2024-03-03 15:39:18	8676.24
2	2024-04-08 09:53:01	6327.20
1	2023-12-26 05:48:22	8360.43
1	2023-12-15 14:52:51	9101.30
3	2023-12-21 15:57:50	3857.98

Sample Output

```
+-----+-----+-----+
|name          |month   |total_sales|
+-----+-----+-----+
|Luxury Gold Watch |January |25556.02 |
|Smartphone Holder Stand |January |11456.80 |
|Stainless Steel Water Bottle|January |1942.79 |
|Luxury Gold Watch |February|17222.22 |
|Stainless Steel Water Bottle|February|9186.38 |
|Smartphone Holder Stand |March   |14918.85 |
|Stainless Steel Water Bottle|March   |14498.21 |
+-----+-----+-----+
```

Question - 33

SQL: IT Project Resource Analysis

An IT company is looking to analyze the allocation of resources across various projects. The goal is to understand the distribution of employees among projects, focusing on projects that require more resources due to their complexity or scale. The analysis should provide insights into the number of employees assigned to each project, the average experience level of the team, and identify projects that are potentially understaffed or overstaffed.

The result should have the following columns: *project_name* / *employee_count* / *avg_experience_years* / *is_understaffed*.

- *project_name* - the name of the project.
- *employee_count* - the total number of employees assigned to the project.
- *avg_experience_years* - the average years of experience per employee on the project, rounded up to the nearest integer (the ceiling)
- *is_understaffed* - the derived column that shows Yes if the employee count is less than 5, otherwise No .

The result should be sorted in descending order by *employee_count*, and then in ascending order by *project_name*.

Note:

- The same employee can be assigned to two or more projects.
- The same employee can have two or more assignments to the same project, so the number of employees on a project should be counted as the number of assignments on that project.
- Only include projects with more than 2 years of average experience per employee.

▼ Schema

projects

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the project
name	VARCHAR(255)		The name of the project

employees

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the employee
ein	VARCHAR(255)		The employee identification number
experience_years	INT		Years of experience

projects_employees

Name	Type	Constraints	Description
project_id	INT	FOREIGN KEY(project_id => projects.id)	The reference to the project
employee_id	INT	FOREIGN KEY(employee_id => employees.id)	The reference to the employee

▼ Sample Data Tables

projects

id	name
1	Project X
2	Sunshine Project
3	Blue Sky Initiative

employees

id	ein	experience_years
1	62-0524667	4

2	62-1435366	1
3	29-3144922	1
4	80-9606443	1
5	63-6630813	1

projects_employees

project_id	employee_id
1	1
1	1
1	2
1	3
1	5
2	1
2	1
2	2
2	5
3	1
3	1
3	2
3	3
3	3
3	4
3	4
3	5
3	5
3	5
3	5

Sample Output

```
+-----+-----+-----+-----+
|project_name |employee_count|avg_experience_years|is_understaffed|
+-----+-----+-----+-----+
|Project X    |5            |3                 |No              |
|Sunshine Project|4            |3                 |Yes             |
+-----+-----+-----+-----+
```

Question - 34

SQL: Ethereum Market Dashboard Analysis

A cryptocurrency trading platform is interested in analyzing the Ethereum transactions on its platform. Specifically, the platform wants to understand the buying and selling activities of Ethereum for each wallet address.

The result should have the following columns: *wallet / total_transactions / total_bought / total_sold*.

- *wallet* - the unique identifier of a wallet on the platform.
- *total_transactions* - the total number of transactions for the specific wallet address.
- *total_bought* - the total amount of Ethereum bought by the wallet, rounded to two decimal places, including trailing zeros if necessary (e.g., 5.00).
- *total_sold* - the total amount of Ethereum sold by the wallet, rounded to two decimal places, including trailing zeros if necessary (e.g., 5.00).

The result should be sorted in ascending order by *wallet*.

Note:

- Positive transaction values represent purchases, negative values represent sales.
- Only include transactions that occurred in February 2024.

▼ Schema

transactions

Name	Type	Constraints	Description
dt	VARCHAR(19)		The date and time when the transaction occurred
wallet	VARCHAR(255)		The wallet address involved in the transaction
amount	DECIMAL(4,2)		The amount of Ethereum transacted

▼ Sample Data Tables

transactions

dt	wallet	amount
2024-01-31 13:42:19	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	-5.78
2024-01-24 06:07:14	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	-2.79
2024-01-25 06:39:19	0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F	-2.73
2024-01-29 04:37:45	0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F	4.68
2024-02-28 06:20:04	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	-7.36
2024-02-12 07:45:28	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	-3.71
2024-02-25 10:49:54	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	-3.53
2024-02-03 19:43:00	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	4.01
2024-02-14 08:55:30	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	8.20
2024-02-16 04:31:26	0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F	-8.96
2024-02-06 23:45:31	0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F	-7.88
2024-02-11 01:00:35	0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F	-7.66
2024-02-25 09:39:01	0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F	-7.45
2024-02-14 04:04:15	0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F	4.17

2024-02-15 11:47:23	0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F	7.56
2024-02-24 14:58:54	0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F	-1.45
2024-02-18 21:17:24	0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F	1.05
2024-02-19 11:12:32	0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F	3.67
2024-03-09 16:52:14	0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c	1.07
2024-03-07 10:29:46	0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F	7.26

Sample Output

```
+-----+-----+-----+-----+
|wallet           |total_transactions|total_bought|total_sold|
+-----+-----+-----+-----+
|0x1a2b3c4d5e6f7a8b9c0d1e2f3a4b5c | 5          | 12.21     | 14.60    |
|0x3a4FbC5Df2E1bBfDdE5c4fA7bF8dE7aC1F| 6          | 11.73     | 31.95    |
|0x9B8aDc2eFf4cC3DdEe5f6a7B8dE9aC1F | 3          | 4.72      | 1.45     |
+-----+-----+-----+-----+
```

Question - 35

SQL: Employee Leave Tracker

A human resources department needs to track the number of leave days taken by employees in a specific year. The aim is to identify employees who have exceeded their annual leave allowance.

The result should have the following columns: *email / leave_days_taken / leave_status*.

- *email* - the email address of the employee.
- *leave_days_taken* - the total number of leave days taken by the employee.
- *leave_status* - a derived column that shows `Within Limit` if the total number of leave days taken by the employee is less than or equal to 20, otherwise `Exceeded`.

The result should be sorted in ascending order by *email*.

Note:

- Only consider leave records that created in 2023.

▼ Schema

employees			
Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the employee
email	VARCHAR(255)		The email address of the employee

leave_records

Name	Type	Constraints	Description
employee_id	INT	FOREIGN KEY(employee_id => employees.id)	The reference to the employee
leave_dt	VARCHAR(19)		The date and time when the leave record was created
days_taken	INT		The number of leave days taken

▼ Sample Data Tables

employees

id	email
1	jquarly0@macromedia.com
2	cchastand1@stanford.edu
3	lpuckrin2@creativecommons.org

leave_records

employee_id	leave_dt	days_taken
1	2022-11-10 11:52:14	4
1	2022-09-07 23:22:46	1
2	2022-11-11 01:47:50	7
3	2022-11-06 23:12:27	7
3	2022-11-17 07:43:18	7
1	2023-05-19 04:40:25	2
1	2023-12-25 16:29:51	7
1	2023-03-12 18:54:29	1
1	2023-08-23 12:33:56	6
2	2023-04-20 04:19:10	5
2	2023-04-28 00:41:50	7
3	2023-06-11 18:49:25	2
3	2023-12-23 15:53:10	7
3	2023-03-13 13:46:16	2
3	2023-10-08 11:57:43	2
3	2023-04-12 07:49:02	4
3	2023-01-17 06:05:35	6
1	2024-02-05 16:01:59	1
1	2024-01-05 22:15:30	7
2	2024-02-21 00:50:11	4

Sample Output

```
+-----+-----+-----+
| email          | leave_days_taken | leave_status |
+-----+-----+-----+
| cchastand1@stanford.edu | 12           | Within Limit |
| jquarly0@macromedia.com   | 16           | Within Limit |
| lpuckrin2@creativecommons.org | 23           | Exceeded     |
+-----+-----+-----+
```

Question - 36

SQL: Email Platform Engagement Stats

An email marketing platform wants to analyze user engagement by generating a report that shows the total number of emails sent, the total number of emails opened, and the open rate for each campaign.

The result should have the following columns: *name / total_emails_sent / total_emails_opened / open_rate*.

- *name* - the name of the email campaign.
- *total_emails_sent* - the total number of emails sent in the campaign.
- *total_emails_opened* - the total number of emails that were opened in the campaign.
- *open_rate* - percentage of emails opened out of the total emails sent, calculated as `(emails_opened / emails_sent) * 100`, rounded to two decimal places, including trailing zeros if necessary (e.g., 5.00).

The result should be sorted in descending order by *open_rate*, and then alphabetically by *name*.

Note:

- Only include email campaigns with an open rate greater than 50%.

▼ Schema

campaigns

Name	Type	Constraints	Description
id	INT	PRIMARY KEY	The identifier of the email campaign
name	VARCHAR(255)		The name of the email campaign

email_stats

Name	Type	Constraints	Description
campaign_id	INT	FOREIGN KEY(campaign_id => campaigns.id)	The reference to the email campaign
emails_sent	INT		The number of emails sent in the email campaign
emails_opened	INT		The number of emails opened in the email campaign

▼ Sample Data Tables

campaigns

id	name
1	SummerSale2021
2	FallPromo
3	WinterWonderland

email_stats

campaign_id	emails_sent	emails_opened
1	1749	775
1	641	423

1	976	598
1	756	121
1	975	107
1	752	367
1	1068	809
1	1046	589
1	1212	939
1	567	214
2	1084	283
2	992	478
2	1505	604
3	899	315
3	742	554
3	1744	917
3	1163	423
3	1501	948
3	736	451
3	537	434

Sample Output

```
+-----+-----+-----+-----+
| name           | total_emails_sent | total_emails_opened | open_rate |
+-----+-----+-----+-----+
| WinterWonderland | 7322             | 4042              | 55.20     |
| SummerSale2021   | 9742             | 4942              | 50.73     |
+-----+-----+-----+-----+
```

Question - 37

SQL: Bond Maturity Analysis

A team of financial analysts wants a comprehensive assessment of the maturity dates associated with mortgages within various mortgage-backed bonds held by their company. Draft a query that delivers insights for each bond.

The insights required are a count of the mortgage maturity dates per bond, the earliest maturity date, the latest maturity date, and the average days to maturity.

The result should have the following columns: *name / maturity_dates / earliest_maturity / latest_maturity / avg_days_to_maturity*.

- *name* - name of the bond
- *maturity_dates* - total number of maturity dates for a specific bond
- *earliest_maturity* - earliest maturity date for a specific bond
- *latest_maturity* - latest maturity date for a specific bond

- *avg_days_to_maturity* - average days to maturity for a specific bond, rounded up to the nearest whole number, e.g., 4.01 rounds up to 5.

The result should be sorted in ascending order by *name*.

Note:

- Only bonds with an average days to maturity greater than 365 days should be included in the result.
- Today is September 13, 2023.
- Use actual days rather than 30-day months.

▼ Schema

bonds			
name	type	constraint	description
id	INT	PRIMARY KEY	Bond ID
name	VARCHAR(255)		Name of the bond

maturities			
name	type	constraint	description
bond_id	INT	FOREIGN KEY(bond_id => bonds.id)	Reference to the bond
maturity	DATE		Maturity date of the bond

▼ Sample Data Tables

bonds	
id	name
1	Alpha Mortgage Bond
2	Beta Mortgage Bond
3	Gamma Mortgage Bond

maturities	
bond_id	maturity
1	2024-01-26
1	2024-02-22
1	2024-03-26
1	2024-05-13
1	2024-07-06

1	2024-08-23
1	2024-09-06
1	2024-11-30
1	2024-12-30
1	2025-04-30
1	2025-05-03
2	2024-07-25
2	2024-12-07
3	2023-12-16
3	2024-01-25
3	2024-01-26
3	2024-05-04
3	2024-10-02
3	2024-12-14
3	2025-01-15

▼ Expected Output

name	maturity_dates	earliest_maturity	latest_maturity	avg_days_to_maturity
Beta Mortgage Bond	2	2024-07-25	2024-12-07	384

Question - 38

SQL: Bond Interest Rate Analysis

A finance team is keen to analyze the coupon rates associated with their variable-rate mortgage-backed bonds. The insights required for each bond are a count of the interest rates seen, the lowest interest rate, the highest interest rate, and the average interest rate.

The result should have the following columns: *name / interest_rates / lowest_rate / highest_rate / avg_rate*.

- *name* - name of the bond
- *interest_rates* - total number of interest rates for a specific bond
- *lowest_rate* - lowest interest rate for a specific bond with one place after the decimal
- *highest_rate* - highest interest rate for a specific bond with one place after the decimal
- *avg_rate* - average interest rate for a specific bond, rounded to two decimal places

The results should be sorted in ascending order by *name*.

Note:

- Only bonds with an average interest rate greater than 3% should be included in the result.
- Retain trailing zeros after the decimal as required, e.g., 5.0 or 3.10.

▼ Schema

bonds			
name	type	constraint	description
id	INT	PRIMARY KEY	Bond ID
name	VARCHAR(255)		Name of the bond

interest_rates			
name	type	constraint	description
bond_id	INT	FOREIGN KEY(bond_id => bonds.id)	Reference to the bond
rate	DECIMAL(2,1)		Interest rate of the bond (in percent)

▼ Sample Data Tables

bonds	
id	name
1	Alpha Mortgage Bond
2	Beta Mortgage Bond
3	Gamma Mortgage Bond

interest_rates	
bond_id	rate
1	1.4
1	1.8
1	2.0
1	2.4
1	3.4
1	4.6
1	4.7
1	4.9

2	2.0
2	2.1
2	3.0
2	3.2
2	4.0
3	1.2
3	1.3
3	1.4
3	2.1
3	2.5
3	3.5
3	4.0

▼ Expected Output

name	interest_rates	lowest_rate	highest_rate	avg_rate
Alpha Mortgage Bond	8	1.4	4.9	3.15

Question - 39

SQL: Bond Cash Flow Analysis for Bondholders

A company's finance department wishes to understand the total cash flow expected from bond investments for each bondholder. Create a query to calculate this based on coupon payments of the bonds each bondholder possesses.

The result should have the following columns: *name / total_cash_flow*.

- *name* - name of the bondholder
- *total_cash_flow* - total expected cash flow from all bonds the bondholder possesses with 2 places after the decimal, e.g., 10 shows as 10.00.

The result should be sorted in descending order by *total_cash_flow*.

Note:

- Only bondholders with a total expected cash flow greater than 10,000.00 should be included in the result.
- The *total_cash_flow* = *annual_coupon* x *coupons_remaining*.

▼ Schema

bondholders

name	type	constraint	description
id	INT	PRIMARY KEY	Bondholder ID
name	VARCHAR(255)		Name of the bondholder

bonds			
name	type	constraint	description
id	INT	PRIMARY KEY	Bond ID
name	VARCHAR(255)		Name of the bond
annual_coupon	DECIMAL(5,2)		Annual coupon payment
coupons_remaining	INT		Number of remaining coupon payments

bondholders_bonds			
name	type	constraint	description
bondholder_id	INT	FOREIGN KEY(bondholder_id => bondholders.id)	Reference to the bondholder
bond_id	INT	FOREIGN KEY(bond_id => bonds.id)	Reference to the bond

▼ Sample Data Tables

bondholders	
id	name
1	Alex Smith
2	Taylor Johnson
3	Jordan Davis

bonds			
id	name	annual_coupon	coupons_remaining
1	Golden Bonds	150.00	4
2	Silver Lining	200.00	2
3	Diamond Trust	100.00	4
4	Emerald Wealth	350.00	5
5	Ruby Returns	150.00	8
6	Sapphire Security	450.00	5

7	Amber Assurance	100.00	8
8	Topaz Treasury	100.00	2
9	Opal Opportunities	150.00	5
10	Pearl Prosperity	450.00	5
11	Platinum Promise	450.00	9
12	Jade Investments	350.00	1
13	Garnet Growth	150.00	4
14	Onyx Returns	350.00	2
15	Quartz Capital	100.00	2
16	Citrine Securities	250.00	2
17	Aquamarine Assets	250.00	2
18	Peridot Portfolio	300.00	8
19	Tourmaline Trust	100.00	6
20	Moonstone Money	150.00	9

bondholders_bonds	
bondholder_id	bond_id
1	1
1	2
1	6
1	8
1	9
1	13
1	14
1	16
1	17
2	4
2	5
2	7
2	11
2	15

2	18
3	3
3	10
3	12
3	19
3	20

▼ Expected Output

name	total_cash_flow
Taylor Johnson	10400.00

Question - 40

SQL: Sum of the Cash Flows Analysis

A financial analysis firm is building a tool to analyze the sum of the cash flows from bond investments. Create a query to extract insights on the cash flows per investor.

The required statistics are a list of all investors and their total, minimum, maximum, and average cash flows from investments.

The result should have the following columns: *email* / *investments* / *min_cash_flow* / *max_cash_flow* / *avg_cash_flow*.

- *email* - investor email
- *investments* - total number of investments by a specific investor
- *min_cash_flow* - minimum cash flow from investments for a specific investor
- *max_cash_flow* - maximum cash flow from investments for a specific investor
- *avg_cash_flow* - average cash flow from investments for a specific investor, rounded to two decimal places

The result should be sorted in ascending order by *email*.

Note:

- Only investors who have a total cash flow greater than 1,000,000 should be included in the results.

▼ Schema

investors			
name	type	constraint	description
id	INT	PRIMARY KEY	Investor ID
email	VARCHAR(255)	UNIQUE	Email address

cash_flows

name	type	constraint	description
investor_id	INT	FOREIGN KEY(investor_id => investors.id)	Reference to the investor
cash_flow	DECIMAL(8,2)		Cash flow amount

▼ Sample Data Tables

investors	
id	email
1	ematson0@ebay.co.uk
2	lsalvadore1@msn.com
3	aclowser2@patch.com

cash_flows	
investor_id	cash_flow
1	184040.12
1	179280.08
1	179374.42
1	79302.21
1	87466.20
1	194588.36
1	153563.92
1	56377.92
2	59039.14
2	167247.23
2	59311.03
2	183883.00
2	118851.21
3	58868.62
3	96909.26
3	103735.73

3	171261.97
3	86463.11
3	56931.73
3	194699.58

▼ Expected Output

email	investments	min_cash_flow	max_cash_flow	avg_cash_flow
ematson0@ebay.co.uk	8	56377.92	194588.36	139249.15

Question - 41

SQL: Expected Cash Flow Analysis

A financial analytics platform is developing a feature to help investors understand the expected cash flows from their bond investments. Create a query that provides insights into each investor's expected cash flows.

The output should list all investors with their total number of investments, the sum of expected cash flows, and the range of expected cash flows (the difference between the highest and the lowest expected cash flow).

The result should have the following columns: *email* / *investment_count* / *total_expected_flow* / *range_expected_flow*.

- *email* - investor email
- *investment_count* - total number of investments held by a specific investor
- *total_expected_flow* - sum of expected cash flows for a specific investor
- *range_expected_flow* - the range of expected cash flows for a specific investor (maximum - minimum)

Sort the result in ascending order by *email*.

Note:

- Only investors who have a sum of expected cash flows greater than 100,000 should be included in the result.

▼ Schema

investors			
name	type	constraint	description
id	INT	PRIMARY KEY	Investor ID
email	VARCHAR(255)	UNIQUE	Email address

cash_flows			
name	type	constraint	description

investor_id	INT	FOREIGN KEY(investor_id => investors.id)	Reference to the investor
expected_flow	DECIMAL(8,2)		Expected cash flow amount

▼ Sample Data Tables

investors	
id	email
1	tdowner0@timesonline.co.uk
2	cgarza1@opera.com
3	nbarwise2@si.edu

cash_flows	
investor_id	expected_flow
1	24923.83
1	30212.10
1	87126.50
1	56018.65
1	93357.47
1	55073.54
1	27095.07
2	44165.12
2	43658.84
2	35835.34
2	12660.46
2	58676.60
2	95929.25
2	47161.23
2	80283.91
2	54427.20
2	93223.98
3	19741.35

3	12243.25
3	50470.06

▼ Expected Output

email	investment_count	total_expected_flow	range_expected_flow
cgarza1@opera.com	10	566021.93	83268.79
tdowner0@timesonline.co.uk	7	373807.16	68433.64

Question - 42

SQL: Online Store Coupon Codes Report

An online store needs a report on the usage of its coupon codes.

The result should have the following columns: *coupon_code / description / total_uses / min_discount / max_discount / avg_discount*.

- *coupon_code* - the unique code of the coupon
- *description* - a description of the coupon code
- *total_uses* - the total number of times the coupon code has been redeemed
- *min_discount* - the smallest discount amount given using this coupon code
- *max_discount* - the largest discount amount given using this coupon code
- *avg_discount* - the average discount amount given using this coupon code, rounded to two places after the decimal

The result should be sorted in ascending order by *coupon_code*.

Note:

- Only active coupons should be included in the report.

▼ Schema

coupons			
name	type	constraint	description
id	INT	PRIMARY KEY	Coupon Code ID
coupon_code	VARCHAR(255)	UNIQUE	Unique code of the coupon
description	VARCHAR(255)		Description of the coupon code
is_enabled	SMALLINT		Coupon enable flag

coupon_uses			
name	type	constraint	description
coupon_id	INT	FOREIGN KEY(coupon_id => coupons.id)	Reference to the coupon

amount	DECIMAL(4,2)		Discount amount given using the coupon
--------	--------------	--	--

▼ Sample Data Tables

coupons			
id	coupon_code	description	is_enabled
1	COUPON123	nisi nam ultrices libero non	0
2	SAVE20	ac est lacinia	1
3	DISCOUNT50	quis odio consequat	1

coupon_uses	
coupon_id	amount
1	36.68
1	3.56
1	2.10
1	39.58
2	39.81
2	24.07
2	28.42
2	31.03
2	3.24
2	36.33
3	8.89
3	30.44
3	36.94
3	42.65
3	33.61
3	41.92
3	1.78
3	20.26
3	27.92

3	0.23
---	------

▼ Expected Output

coupon_code	description	total_uses	min_discount	max_discount	avg_discount
DISCOUNT50	quis odio consequat	10	0.23	42.65	24.46
SAVE20	ac est lacinia	6	3.24	39.81	27.15

Question - 43

SQL: Freelancer Platform Yearly Income Report

A freelancer platform wants to generate a yearly report detailing the income they have earned from their top professions.

The result should have the following columns: *title / total_projects / total_income / total_freelancers / average_income_per_freelancer*.

- *title* - the profession title (e.g., Web Development, Graphic Design)
- *total_projects* - the total number of projects in each profession
- *total_income* - the total income the platform earned from each profession
- *total_freelancers* - the total number of freelancers in each profession
- *average_income_per_freelancer* - the average income earned from each project in the profession, rounded to two decimal places

The result should be sorted in descending order by *total_income*.

Note:

- Only completed projects should be included in the report.

▼ Schema

professions			
name	type	constraint	description
id	INT	PRIMARY KEY	Profession ID
title	VARCHAR(255)	UNIQUE	Title of profession (e.g., Web Development)

freelancers			
name	type	constraint	description
id	INT	PRIMARY KEY	Freelancer ID
profession_id	INT	FOREIGN KEY(profession_id => professions.id)	Reference to the profession
email	VARCHAR(255)	UNIQUE	Email of the freelancer

projects

name	type	constraint	description
id	INT	PRIMARY KEY	Project ID
freelancer_id	INT	FOREIGN KEY(freelancer_id => freelancers.id)	Reference to the freelancer
status	ENUM('Completed','Ongoing','Cancelled')		Status of the project
income	DECIMAL(6,2)		Income earned by the platform for the project

▼ Sample Data Tables

professions	
id	title
1	Artificial Intelligence Engineer
3	Game Developer
2	Network Administrator

freelancers		
id	profession_id	email
1	1	lfernez0@microsoft.com
3	2	mbrydone2@delicious.com
4	2	jhamp3@4shared.com
5	3	cparfett4@twitter.com

projects			
id	freelancer_id	status	income
5	1	Completed	8562.13
11	1	Completed	6727.56
10	3	Completed	3753.46
20	3	Completed	6659.39
6	4	Completed	8459.28
13	4	Completed	5899.31
16	4	Completed	2709.63

4	5	Completed	5029.44
7	5	Completed	1763.94
9	5	Completed	6988.36
8	3	Cancelled	8699.67
1	5	Cancelled	5403.21
19	3	Ongoing	72.51
3	4	Ongoing	8561.14
15	4	Ongoing	9235.78
17	4	Ongoing	4307.76

▼ Expected Output

title	total_projects	total_income	total_freelancers	average_income_per_freelancer
Network Administrator	5	27481.07	2	5496.21
Artificial Intelligence Engineer	2	15289.69	1	7644.85
Game Developer	3	13781.74	1	4593.91

Question - 44

SQL: Ecommerce Warehouse Stock Report

An ecommerce website admin panel needs to generate a report to display product stock status in their warehouse. The report should display a list of all product categories and the total number of products available in each category. It should be limited to products with more than 10 units available.

The result should have the following columns: *category / title / total_stock*.

- *category* - the name of the product category (e.g., Electronics, Clothing)
- *title* - the title of the product
- *total_stock* - total number of items available in stock for that product

The result should be sorted in ascending order by *category*, then in ascending order by *title*, and finally in descending order by *total_stock*.

Note:

- Only products with a total stock of more than 10 items should be included in the report.

▼ Schema

categories			
name	type	constraint	description
id	INT	PRIMARY KEY	Category ID

title	VARCHAR(255)	UNIQUE	Title of the product category
-------	--------------	--------	-------------------------------

products			
name	type	constraint	description
id	INT	PRIMARY KEY	Product ID
category_id	INT	FOREIGN KEY(category_id => categories.id)	Category ID reference
title	VARCHAR(255)		Title of the product
sku	VARCHAR(255)	UNIQUE	SKU of the product
stock_number	INT		Number of items available in stock

▼ Sample Data Tables

categories	
id	title
1	Electronics
2	Clothing
3	Home & Kitchen

products				
id	category_id	title	sku	stock_number
11	1	Elegant Gadget	EG-11	4
3	1	Luxury Gizmo	LG-3	10
19	1	Sleek Widget	SW-19	8
8	1	Sleek Widget	SW-8	8
14	2	Elegant Gadget	EG-14	2
16	2	Elegant Gadget	EG-16	6
10	2	Elegant Gadget	EG-10	10
7	2	Luxury Gizmo	LG-7	3
2	2	Luxury Gizmo	LG-2	8
18	2	Luxury Gizmo	LG-18	9
1	2	Sleek Widget	SW-1	3
6	2	Sleek Widget	SW-6	7

20	3	Elegant Gadget	EG-20	10
9	3	Luxury Gizmo	LG-9	4
12	3	Luxury Gizmo	LG-12	5
13	3	Luxury Gizmo	LG-13	5
5	3	Luxury Gizmo	LG-5	9
4	3	Sleek Widget	SW-4	8
15	3	Sleek Widget	SW-15	9
17	3	Sleek Widget	SW-17	9

▼ Expected Output

category	title	total_stock
Clothing	Elegant Gadget	18
Clothing	Luxury Gizmo	20
Electronics	Sleek Widget	16
Home & Kitchen	Luxury Gizmo	23
Home & Kitchen	Sleek Widget	26

Question - 45

SQL: Antivirus Database Quarantine Report

A prominent cyber-security company aims to generate a report detailing quarantined URLs. This report should list URL domains, their associated threat types, the total number of times they have been quarantined, and the total number of users affected by each threat from a domain.

The result should have the following columns: *domain_name / threat_type / total_occurrences / total_users_affected*.

- *domain_name* - the website's domain (e.g., example.com)
- *threat_type* - the type of threat identified (e.g., Malware, Phishing)
- *total_occurrences* - the number of times this domain was quarantined for the specific threat
- *total_users_affected* - the number of users affected by each threat from the domain

The result should be sorted in descending order by *total_users_affected*, then ascending by *domain_name*.

Note:

- Only URLs with a status of "Quarantined" should be included in the report.

▼ Schema

threat_types

name	type	constraint	description
id	INT	PRIMARY KEY	Threat type ID
threat_type	VARCHAR(255)		Type of threat (e.g., Malware, Phishing)

quarantine_urls			
name	type	constraint	description
id	INT	PRIMARY KEY	URL ID
threat_id	INT	FOREIGN KEY(threat_id => threat_types.id)	Reference to the type of threat
domain_name	VARCHAR(255)		Domain name of the quarantined URL
status	ENUM('Quarantined','Safe','Deleted')		URL status in the system
users_affected	INT		Number of users affected by the quarantined URL

▼ Sample Data Tables

threat_types	
id	threat_type
1	Phishing
2	Rootkit
3	Malware

quarantine_urls				
id	threat_id	domain_name	status	users_affected
17	1	amazon.com	Quarantined	862
16	1	google.com	Quarantined	63
9	1	amazon.com	Quarantined	41
18	2	amazon.com	Quarantined	149
12	2	yahoo.com	Quarantined	967
4	3	amazon.com	Quarantined	377
10	3	yahoo.com	Quarantined	721
11	1	yahoo.com	Deleted	551
20	1	amazon.com	Safe	407
19	1	amazon.com	Deleted	665

15	1	Facebook.com	Safe	52
2	1	google.com	Safe	309
1	2	twitter.com	Safe	562
13	2	Facebook.com	Safe	208
14	2	google.com	Deleted	731
8	2	twitter.com	Safe	924
7	2	twitter.com	Safe	982
6	2	google.com	Deleted	864
3	2	facebook.com	Safe	136
5	3	yahoo.com	Safe	949

▼ Expected Output

domain_name	threat_type	total_occurrences	total_users_affected
yahoo.com	Rootkit	1	967
amazon.com	Phishing	2	903
yahoo.com	Malware	1	721
amazon.com	Malware	1	377
amazon.com	Rootkit	1	149
google.com	Phishing	1	63

Question - 46

SQL: Online Streaming Service Traffic Report

Create a query for an online streaming service. It should return a list of clients, their number of streams, and the total amount of traffic.

The result should have the following columns: *mac_address / streams / total_traffic*.

- *mac_address* - client MAC address
- *streams* - total number of streams for a specific client
- *total_traffic* - total traffic amount of all streams for a specific client

The result should be sorted in descending order by *total_traffic*.

Note:

- Only streams of "720p" quality or higher should be included in the result.

▼ Schema

clients			
name	type	constraint	description
id	INT	PRIMARY KEY	Client ID
mac_address	VARCHAR(255)		MAC address

streams			
name	type	constraint	description
client_id	INT	FOREIGN KEY (client_id => clients.id)	Client ID
title	VARCHAR(255)		Title
quality	ENUM('240p','360p','480p','720p','1080p','1440p','2160p')		Quality
traffic	INT		Traffic

▼ Sample Data Tables

clients	
id	mac_address
1	2F-80-8E-F2-0E-4C
2	A1-F7-D4-48-B9-E6
3	9F-72-DB-7C-73-FC

streams			
client_id	title	quality	traffic
1	Monte Carlo	360p	71928308
1	Separation, The (Sparation, La)	480p	35221785
1	Felidae	480p	54617023
1	Dirty Dancing	1440p	56419563
1	Ragtime	1440p	12404457
1	Oscar	1440p	49717246
1	Barb Wire	2160p	83761463
1	Jason and the Argonauts	2160p	27364051
2	Carry on Cruising	240p	33226462

2	Best of the Best	240p	62793858
2	Ecstasy (xtasis)	240p	73079415
2	Go Go Tales	480p	48836837
2	Nights and Weekends	1440p	32708277
3	Coneheads	480p	92308213
3	Silences of the Palace, The (Saimt el Qusur)	480p	52917945
3	Good Pick	720p	71890218
3	Wuthering Heights	720p	19813053
3	Big Kahuna, The	1080p	28786846
3	Work of Director Michel Gondry, The	2160p	18789351
3	My Best Friends	2160p	44347338

▼ Expected Output

mac_address	streams	total_traffic
2F-80-8E-F2-0E-4C	5	229666780
9F-72-DB-7C-73-FC	5	183626806
A1-F7-D4-48-B9-E6	1	32708277

Question - 47

SQL: Cloud Hosting Instances Performance Statistics

A cloud hosting provider dashboard that is in development requires a query that returns a list of networks where at least one instance within them has reached its maximum CPU usage threshold.

The result should have the following columns: *cidr / instances / avg_cpu_usage / avg_memory_usage / avg_network_usage*.

- *cidr* - network CIDR
- *instances* - total number of instances for a specific network
- *avg_cpu_usage* - average CPU usage of instances for a specific network
 - Record format is `##%`, where the placeholders are `##` in the order they appear:
 1. Average CPU usage across all instances of a specific network, rounded up to the nearest integer
- *avg_memory_usage* - average memory usage of instances for a specific network
 - Record format is `##%`, where the placeholders are `##` in the order they appear:
 1. Average memory usage across all instances of a specific network, rounded up to the nearest integer
- *avg_network_usage* - average network usage of instances for a specific network
 - Record format is `##%`, where the placeholders are `##` in the order they appear:
 1. Average network usage across all instances of a specific network, rounded up to the nearest integer

The result should be sorted in ascending order by *cidr*.

Note:

- Only include networks where there is at least one instance with *cpu_usage* of 80% or greater.

▼ Schema

networks			
name	type	constraint	description
id	INT	PRIMARY KEY	Network ID
cidr	VARCHAR(255)		IP Address v4 CIDR

instances			
name	type	constraint	description
network_id	INT	FOREIGN KEY (network_id => networks.id)	Network ID
cpu_usage	VARCHAR(255)		CPU usage
memory_usage	VARCHAR(255)		Memory usage
network_usage	VARCHAR(255)		Network usage

▼ Sample Data Tables

networks	
id	cidr
1	24.77.36.156/9
2	74.213.138.70/7
3	167.244.163.58/29

instances			
network_id	cpu_usage	memory_usage	network_usage
1	20%	74%	74%
3	26%	9%	99%
3	2%	21%	97%
1	51%	19%	89%
2	2%	27%	79%
3	92%	35%	41%

2	27%	5%	44%
3	67%	47%	79%
1	14%	28%	43%
3	47%	0%	53%
1	38%	3%	46%
2	71%	51%	6%
3	77%	74%	53%
3	31%	48%	80%
2	31%	42%	24%
1	77%	65%	46%
2	51%	94%	41%
3	8%	3%	57%
1	1%	56%	62%
2	15%	66%	65%

▼ Expected Output

cidr	instances	avg_cpu_usage	avg_memory_usage	avg_network_usage
167.244.163.58/29	8	44%	30%	70%

Question - 48

SQL: AI Video Processing Service Usage Time Calculation

A video processing service in development uses AI to improve the quality of the input video. The service uses a process separation mechanism to improve performance and minimize overall execution time.

Create a query that returns a list of all tasks and the total execution time for all processes of those tasks.

The result should have the following columns: *hash / usage_time*.

- *hash* - task hash
- *usage_time* - total execution time in seconds of all processes for a specific task

The result should be sorted in descending order by *usage_time*.

▼ Schema

tasks			
name	type	constraint	description
id	INT	PRIMARY KEY	Task ID
hash	VARCHAR(255)		Hash

processes			
name	type	constraint	description
task_id	INT	FOREIGN KEY (task_id => tasks.id)	Task ID
start_dt	VARCHAR(19)		Start timestamp
end_dt	VARCHAR(19)		End timestamp

▼ Sample Data Tables

tasks	
id	hash
1	208f95e0fcff792f617ade3cebf33ad9
2	0f44a9ffead2f18a7f25425c1260fc74
3	dbcf54e94395c32e01ec09a5db731912

processes		
task_id	start_dt	end_dt
1	2023-04-20 02:01:16	2023-04-20 02:11:35
1	2023-04-09 15:11:10	2023-04-09 15:26:43
1	2023-04-07 23:41:49	2023-04-08 00:34:10
2	2023-04-07 23:05:47	2023-04-08 00:00:05
2	2023-04-19 18:39:33	2023-04-19 18:54:57
2	2023-04-28 13:17:11	2023-04-28 13:24:37
2	2023-04-16 00:13:06	2023-04-16 01:02:39
2	2023-04-16 15:02:26	2023-04-16 15:58:14
2	2023-04-27 02:23:07	2023-04-27 02:59:13
2	2023-04-10 23:33:47	2023-04-11 00:09:35
2	2023-04-16 17:29:51	2023-04-16 18:10:22

2	2023-04-23 12:16:01	2023-04-23 12:48:07
3	2023-04-01 02:25:12	2023-04-01 02:49:26
3	2023-04-04 03:02:43	2023-04-04 03:42:03
3	2023-04-10 22:42:26	2023-04-10 23:14:42
3	2023-04-09 17:46:12	2023-04-09 18:10:19
3	2023-04-25 15:09:36	2023-04-25 15:19:54
3	2023-04-19 14:39:52	2023-04-19 15:21:23
3	2023-04-12 04:22:29	2023-04-12 04:25:10
3	2023-04-25 07:40:26	2023-04-25 08:01:30

▼ Expected Output

hash	usage_time
0f44a9ffead2f18a7f25425c1260fc74	19620
dbcf54e94395c32e01ec09a5db731912	11731
208f95e0fcff792f617ade3cebf33ad9	4693

Question - 49

SQL: Benchmarking Tool Report

A benchmarking tool that is under development needs a query that returns all devices and information about the achieved performance class.

The result should have the following columns: *device*.

- *device* - device benchmark result:
 - Record format is `Device ## has class: ##`, where the placeholders are `##` in the order they appear:
 1. Device ID
 2. The benchmark performance evaluation class, which can be calculated as:
 - "A", when `score ≥ 80`
 - "B", when `60 ≤ score < 80`
 - "C", when `40 ≤ score < 60`
 - "D", when `20 ≤ score < 40`
 - "F", when `score < 20`

The result should be sorted in ascending order by device ID.

▼ Schema

devices

name	type	constraint	description
id	INT	PRIMARY KEY	Device ID
score	INT		Score

▼ Sample Data Tables

devices	
id	score
1	20
2	50
3	50
4	68
5	95

▼ Expected Output

device
Device 1 has class: D
Device 2 has class: C
Device 3 has class: C
Device 4 has class: B
Device 5 has class: A

Question - 50

SQL: Smart Home Application Customer Report

As part of a smart home application, create a query that, based on data from meter readings, calculates the total electricity consumption, the name of the most expensive tariff consumed, and the total cost for invoicing.

The result should have the following columns: *username* / *email* / *highest_tariff* / *consumption* / *total_cost*.

- *username* - account username
- *email* - account email address
- *highest_tariff* - name of the most expensive tariff that is consumed
- *consumption* - total consumption amount
- *total_cost* - total cost of all consumed electricity, rounded to two decimal places, e.g., 8.404 shows 8.40.

The results should be sorted in ascending order by *username*.

Note:

- The total cost of all electricity consumed is the sum of the products of all meter readings and their respective consumed tariffs.

▼ Schema

accounts			
name	type	constraint	description
id	INT	PRIMARY KEY	Account ID
username	VARCHAR(255)		Account username
email	VARCHAR(255)		Account email address

tariffs			
name	type	constraint	description
id	INT	PRIMARY KEY	Tariff ID
name	ENUM('A','B','C','D','E')		Tariff name
cost	DECIMAL(4,3)		Tariff cost

readings			
name	type	constraint	description
account_id	INT	FOREIGN KEY (account_id => accounts.id)	Account ID
tariff_id	INT	FOREIGN KEY (tariff_id => tariffs.id)	Tariff ID
amount	SMALLINT		Readings amount

▼ Sample Data Tables

accounts		
id	username	email
1	hshillabeare0	rcalkin0@sourceforge.net
2	sdandy1	agaule1@businessweek.com
3	sgreiswood2	toppy2@lulu.com

tariffs		
id	name	cost

1	A	0.010
2	B	0.020
3	C	0.050
4	D	0.075
5	E	0.100

readings		
account_id	tariff_id	amount
1	2	54
1	3	19
1	3	37
1	3	89
1	3	119
2	1	12
2	1	44
2	1	81
2	2	60
2	2	164
2	2	199
2	3	79
2	5	186
3	1	31
3	1	59
3	1	77
3	1	95
3	1	110
3	1	125
3	2	31

▼ Expected Output

username	email	highest_tariff	consumption	total_cost
hshillabeare0	rcalkin0@sourceforge.net	C	318	14.28
sdandy1	agaule1@businessweek.com	E	825	32.38
sgreiswood2	toppy2@lulu.com	B	528	5.59

Question - 51

SQL: MMORPG Game Inventory Overload Notification

An MMORPG game is under development. For the profile and inventory mechanics, it needs a query that calculates a list of game accounts whose inventory is overloaded with game items.

The result should have the following columns: *username* / *email* / *items* / *total_weight*.

- *username* - account username
- *email* - account email address
- *items* - total number of items in inventory
- *total_weight* - total weight of items in inventory

The result should be sorted in descending order by *total_weight*, then in ascending order by *username*.

Note:

- Each item in the inventory has its own weight.
- Only accounts where the total weight of all items in the inventory exceeds the overload threshold should be included in the result.
- The overload threshold is 20.

▼ Schema

accounts			
name	type	constraint	description
id	INT	PRIMARY KEY	Account ID
username	VARCHAR(255)		Account username
email	VARCHAR(255)		Account email address

items			
name	type	constraint	description
id	INT	PRIMARY KEY	Item ID
type	ENUM("sword','shield','armor")		Item type
name	VARCHAR(255)		Item name
weight	SMALLINT		Item weight

accounts_items			
name	type	constraint	description
account_id	INT	FOREIGN KEY (account_id => accounts.id)	Account ID
item_id	INT	FOREIGN KEY (item_id => items.id)	Item ID

▼ Sample Data Tables

accounts		
id	username	email
1	esoane0	alefwich0@nytimes.com
2	jrafter1	bmcniff1@census.gov
3	rcawston2	fnickoll2@flickr.com

items			
id	type	name	weight
1	shield	Shield of Asteraceae	3
2	sword	Sword of Cyperaceae	3
3	shield	Shield of Apiaceae	3
4	sword	Sword of Onagraceae	3
5	sword	Sword of Campanulaceae	3

accounts_items	
account_id	item_id
1	2
1	3
1	3
1	4
1	4
1	5
1	5
1	5

2	1
2	1
2	2
2	2
2	2
2	2
2	2
2	3
2	3
2	5
3	3
3	4

▼ Expected Output

username	email	items	total_weight
jrafter1	bmcniff1@census.gov	10	30
esoane0	alefwich0@nytimes.com	8	24

Question - 52

SQL: Outdoor Banner Digital Marketplace Placement Report

In developing a digital marketplace for outdoor banners, a query is needed to return information about the banners in each of the cities.

The result should have the following columns: *city / banners / min_area / avg_area / max_area / total_area*.

- *city* - city name
- *banners* - total number of banners for a specific city
- *min_area* - minimum available banner area for a specific city
- *avg_area* - average banner area for a specific city, rounded up to the nearest integer, e.g., $\text{ceiling}(1.1) = 2$
- *max_area* - maximum available banner area for a specific city
- *total_area* - total available banner area for a specific city

The result should be sorted in ascending order by *city*.

Note:

- Banners are rectangular.

▼ Schema

cities			
name	type	constraint	description
id	INT	PRIMARY KEY	City ID
name	VARCHAR(255)		City name

banners			
name	type	constraint	description
city_id	INT	FOREIGN KEY (city_id => cities.id)	City ID
width	SMALLINT		Banner width
height	SMALLINT		Banner height

▼ Sample Data Tables

cities	
id	name
1	Kayu Agung
2	Yangkou
3	Marseille

banners		
city_id	width	height
3	6	20
1	20	14
1	6	17
1	15	6
2	16	8
2	6	7
3	6	9
1	20	16
3	19	14
2	9	17

2	8	12
1	12	16
3	15	14
3	11	7
3	6	14
2	12	7
3	7	20
1	13	6
3	10	13
2	19	15

▼ Expected Output

city	banners	min_area	avg_area	max_area	total_area
Kayu Agung	6	78	177	320	1062
Marseille	8	54	136	266	1081
Yangkou	6	42	132	285	788

Question - 53

SQL: Auction Web Service Lot Statistics

An auction web service needs a query that returns a list of all available lots with offers left by buyers.

The result should have the following columns: *name* / *min_offer* / *avg_offer* / *max_offer* / *offers*.

- *name* - lot name
- *offers* - total number of offers for a specific lot
- *min_offer* - minimum offer for a specific lot
- *avg_offer* - average offer for a specific lot, rounded to two decimal places
- *max_offer* - maximum offer for a specific lot

The result should be sorted in descending order by *offers*.

Notes:

- Some lots may not have offers.
- Values for *min_offer*, *avg_offer*, and *max_offer* should always show two places after the decimal, e.g., 10 is shown as 10.00.

▼ Schema

lots			
name	type	constraint	description
id	INT	PRIMARY KEY	Lot ID
name	VARCHAR(255)		Lot name

offers			
name	type	constraint	description
lot_id	INT	FOREIGN KEY (lot_id => lots.id)	Lot ID
amount	DECIMAL(6,2)		Offer amount

▼ Sample Data Tables

lots	
id	name
1	Merremia quinquefolia (L.) Hallier f.
2	Plantago maritima L.
3	Hohenbergia antillana Mez
4	Penstemon eriantherus Pursh var. argillosus M.E. Jones

offers	
lot_id	amount
1	510.51
2	703.80
2	181.80
1	38.06
2	368.78
3	91.40
2	413.80
3	157.99
3	885.82
2	863.99
1	307.61

2	120.39
1	771.96
2	801.42
3	871.59
1	541.61
3	477.62
2	303.29
2	612.83
3	464.98

▼ Expected Output

name	offers	min_offer	avg_offer	max_offer
Plantago maritima L.	9	120.39	485.57	863.99
Hohenbergia antillana Mez	6	91.40	491.57	885.82
Merremia quinquefolia (L.) Hallier F.	5	38.06	433.95	771.96
Penstemon eriantherus Pursh var. argillosus M.E. Jones	0	NULL	NULL	NULL

Question - 54

SQL: Tax Calculator Web Service Simple Report

Create a query for a tax calculator web service. It should return a list of all accounts, the total taxable income, the account's individual tax rate, and the amount of tax payable.

The result should have the following columns: *full_name* / *iban* / *income* / *rate* / *tax*.

- *full_name* - full profile name:
 - Record format is ## ##, where the placeholders are ## in the order they appear:
 1. Last name of the profile
 2. First name of the profile
- *iban* - account IBAN
- *income* - total taxable income
- *rate* - individual tax rate of the account
- *tax* - payable amount of tax, rounded to two decimal places

The result should be sorted in ascending order by *full_name*.

Note:

- The payable amount of tax is calculated as the total taxable income for the periods Q1'21-Q4'21 multiplied by the individual tax rate.
- The individual tax rate is 10% and applies to all accounts.

▼ Schema

accounts			
name	type	constraint	description
id	INT	PRIMARY KEY	Account ID
first_name	VARCHAR(255)		Account first name
last_name	VARCHAR(255)		Account last name
iban	VARCHAR(255)		Account IBAN

declarations			
name	type	constraint	description
account_id	INT	FOREIGN KEY (account_id => accounts.id)	Account ID
quarter	ENUM('Q1','Q2','Q3','Q4')		Declaration quarter
income	DECIMAL(7,2)		Declaration taxable income

▼ Sample Data Tables

accounts			
id	first_name	last_name	iban
1	Alex	Cantua	IL29 9590 1551 0560 0553 712
2	Chris	Lashmore	AZ54 CNUI 01DR XEXZ ASKY QM4W F8JI
3	Taylor	Blum	HR20 2041 7741 5014 9873 9
4	Robin	Neachell	NL87 PPCD 0429 1849 92
5	Drew	Barbier	FR72 7843 3990 42WM QC8P GVNV F78

declarations		
account_id	quarter	income
1	Q1	49235.67
1	Q2	46653.11
1	Q3	63739.99
1	Q4	43222.54
2	Q1	69743.50

2	Q2	29641.01
2	Q3	97725.49
2	Q4	91481.98
3	Q1	68402.43
3	Q2	12660.12
3	Q3	59601.65
3	Q4	54701.74
4	Q1	55220.27
4	Q2	87752.41
4	Q3	44447.06
4	Q4	45876.26
5	Q1	42511.74
5	Q2	22022.78
5	Q3	88396.81
5	Q4	67252.54

▼ Expected Output

full_name	iban	income	rate	tax
Barbier Drew	FR72 7843 3990 42WM QC8P GNVF78	220183.87	10%	22018.39
Blum Taylor	HR20 2041 7741 5014 9873 9	195365.94	10%	19536.59
Cantua Alex	IL29 9590 1551 0560 0553 712	202851.31	10%	20285.13
Lashmore Chris	AZ54 CNUI 01DR XEXZ ASKY QM4WF8JI	288591.98	10%	28859.20
Neachell Robin	NL87 PPCD 0429 1849 92	233296.00	10%	23329.60

Question - 55

SQL: Social Network Relationship Statistics

A social network under development needs a query that returns all profiles and statistics about their relationships.

The result should have the following columns: *full_name / email / total_relations / approved_relations / pending_relations*.

- *full_name* - full profile name:
 - Record format is ## ##, where the placeholders are ## in the order they appear:
 1. Last name of the profile
 2. First name of the profile

- *email* - email address
- *total_relations* - total number of relations for a specific profile
- *approved_relations* - total number of "approved" relations for a specific profile
- *pending_relations* - total number of "pending" relations for a specific profile

The result should be sorted in ascending order by *full_name*.

Note:

- A relation is "approved" when *relations.is_approved* is true
- A relation is "pending" when *relations.is_approved* is false

▼ Schema

profiles			
name	type	constraint	description
id	INT	PRIMARY KEY	Profile ID
first_name	VARCHAR(255)		Profile first name
last_name	VARCHAR(255)		Profile last name
email	VARCHAR(255)		Profile email address

relations			
name	type	constraint	description
profile_id	INT	FOREIGN KEY (profile_id => profiles.id)	Profile ID
related_to	VARCHAR(255)		Relation to
is_approved	BOOLEAN		Relation approval flag

▼ Sample Data Tables

profiles			
id	first_name	last_name	email
1	Shayne	Shilito	sshilito0@ftc.gov
2	Shell	Shade	ssshade1@paginegialle.it
3	Nobie	Splain	nsplain2@npr.org

relations		
profile_id	related_to	is_approved
1	cbasinigazzii	1

1	ldevered	1
1	edeniskeb	1
1	cstirland4	1
1	ngooddiea	1
1	alockney7	1
1	jsorrillj	0
1	bnodin3	0
1	dwall2	0
1	folivas1	0
2	ksharland6	0
2	pbarosch8	0
2	smacieja9	0
2	bbrasonf	0
2	dabrahartg	0
3	gaymer5	1
3	rwoolcockse	1
3	egilyott0	1
3	agillionc	0
3	fgrubinh	0

▼ Expected Output

full_name	email	total_relations	approved_relations	pending_relations
Shade Shell	sshade1@paginegialle.it	5	0	5
Shilito Shayne	sshilito0@ftc.gov	10	6	4
Splain Nobie	nsplain2@npr.org	5	3	2

Question - 56

SQL: Online Banking Transactions

A client requested a query for a dashboard in their online banking web service. It should return the count and the sum of transaction amounts for each customer in the current month.

The result should have the following columns: *iban* / *transactions* / *total*.

- *iban* - account IBAN
- *transactions* - total number of transactions
- *total* - the sum of the transaction amounts

The result should be sorted in descending order by *total*.

Note:

- Only transactions in the current month should be included in the result.
- The current month is September.

▼ Schema

accounts			
name	type	constraint	description
id	INT	PRIMARY KEY	Account ID
iban	VARCHAR(255)		Account IBAN

transactions			
name	type	constraint	description
account_id	INT	FOREIGN KEY (account_id => accounts.id)	Account ID
dt	DATETIME		Transaction date and time
amount	DECIMAL(5,2)		Transaction amount

▼ Sample Data Tables

accounts	
id	iban
1	SE48 2961 2087 8112 2835 6438
2	BE89 2286 5514 4847
3	MU84 HRGV 2047 2584 5774 3195 856J PZ

transactions		
account_id	dt	amount
2	2022-09-25 19:24:50	75.06
2	2022-09-24 03:09:17	41.10

1	2022-09-19 04:13:17	65.85
3	2022-09-30 07:18:29	44.57
1	2022-09-26 01:51:44	98.93
1	2022-08-28 02:51:04	60.42
1	2022-08-25 23:25:54	45.34
2	2022-09-09 11:00:48	11.05
3	2022-08-25 19:37:02	53.61
2	2022-09-23 09:44:05	89.18
1	2022-08-28 19:48:40	47.60
3	2022-09-12 10:28:10	96.40
3	2022-10-03 16:49:51	45.41
2	2022-09-05 16:20:41	46.78
3	2022-10-03 04:51:29	50.81
1	2022-09-10 17:31:44	78.72
2	2022-08-31 21:59:56	61.09
2	2022-09-14 12:52:13	20.36
1	2022-09-28 11:05:21	70.52
3	2022-09-30 09:21:12	48.00

▼ Expected Output

iban	transactions	total
SE48 2961 2087 8112 2835 6438	4	314.02
BE89 2286 5514 4847	6	283.53
MU84 HRGV 2047 2584 5774 3195 856J PZ	3	188.97

Question - 57

SQL: Internet Service Provider Monthly Report

An ISP is updating its billing back end. They need a query that returns a list of all client MAC addresses, with traffic for the current billing period and its total cost. The current billing period includes all records where *dt* is in May 2022.

Since the business model of the ISP is significantly different from its competitors, each client has its own tariff per megabyte of data.

The result should have the following columns: *mac / traffic / cost*.

- *mac* - MAC address of the client
- *traffic* - total traffic for a specific client in the current billing period
- *cost* - total cost of traffic for a specific client in the current billing period:
 - Record is the sum of all traffic for a specific client in the current billing period, multiplied by the client's tariff and rounded to two decimal places

The result should be sorted in descending order by *cost*.

Note:

- Only traffic whose reporting date is in May 2022 should be included in the report

▼ Schema

clients		
name	type	description
id	SMALLINT	ID
mac	VARCHAR(17)	MAC address
tariff	DECIMAL(6,5)	Tariff per MB

traffic		
name	type	description
client_id	SMALLINT	Client ID
dt	VARCHAR(19)	Reporting date
amount	INT	Amount

▼ Sample Data Tables

clients		
id	mac	tariff
1	A2-53-FC-0C-3E-B4	0.00007
2	DC-80-42-E9-AE-FC	0.00003
3	3F-9B-A9-2A-B1-7B	0.00007
4	D4-6F-E4-AF-47-D5	0.00004
5	B9-65-FC-8E-F0-15	0.00007

traffic		
client_id	dt	amount

1	2022-05-22	9127
1	2022-06-07	62203
1	2022-06-10	88227
2	2022-05-31	99874
2	2022-06-01	78400
2	2022-06-03	61106
2	2022-06-12	20963
2	2022-06-29	98304
2	2022-07-04	6626
3	2022-05-22	8386
3	2022-06-08	22959
3	2022-07-05	52096
3	2022-07-14	70777
4	2022-05-22	93743
5	2022-05-16	84660
5	2022-05-28	63267
5	2022-06-10	80681
5	2022-06-21	55460
5	2022-07-04	91365
5	2022-07-09	23296

▼ Expected Output

mac	traffic	cost
B9-65-FC-8E-F0-15	147927	10.35
D4-6F-E4-AF-47-D5	93743	3.75
DC-80-42-E9-AE-FC	99874	3.00
A2-53-FC-0C-3E-B4	9127	0.64
3F-9B-A9-2A-B1-7B	8386	0.59

Question - 58
SQL: The Yellow Pages Companies Report

A company maintains "The Yellow Pages" website. They need a query that returns a list of companies and their overall rating based on all the categories in which they are listed.

The result should have the following columns: *name / address / phone / overall_review_rating*.

- *name* - company name
- *address* - company address
- *phone* - company phone number
- *overall_review_rating* - overall review rating for a specific company:
 - Record format is ## (## categories) , where the placeholders are ## in the order they appear:
 1. The average review rating across all categories for a specific company, rounded to one decimal place
 2. The total number of categories in which the company is listed

The result should be sorted in descending order of the average rating (before rounding), then in ascending order by *name*.

▼ Schema

companies		
name	type	description
id	SMALLINT	unique id, primary key
name	VARCHAR(255)	
address	VARCHAR(255)	
phone	VARCHAR(255)	

categories		
name	type	description
company_id	SMALLINT	Foreign key, companies.id
name	VARCHAR(255)	
review_rating	SMALLINT	

▼ Sample Data Tables

companies			
id	name	address	phone
1	Casper, Oberbrunner and Williamson	53 Di Loreto Hill	+420 (569) 566-3689
2	Tromp, Kozey and Abbott	84 Mcguire Plaza	+62 (145) 722-2330
3	Gerlach, Hayes and Stamm	80 Service Point	+86 (731) 234-4119
4	Wolff-Fadel	06 Fair Oaks Trail	+7 (894) 233-0976

categories		
company_id	name	review_rating
1	HVAC	2
2	HVAC	2
2	Retaining Wall and Brick Pavers	1
2	Rebar & Wire Mesh Install	2
3	Prefabricated Aluminum Metal Canopies	2
3	Prefabricated Aluminum Metal Canopies	0
3	RF Shielding	2
3	Overhead Doors	0
3	Rebar & Wire Mesh Install	5
3	Termite Control	0
4	Sitework & Site Utilities	0
4	Electrical and Fire Alarm	2
4	Masonry	2
4	Temp Fencing, Decorative Fencing and Gates	0
4	Elevator	1
4	Drywall & Acoustical (FED)	5
5	Asphalt Paving	0
5	Glass & Glazing	1
5	Framing (Steel)	3
5	Structural & Misc Steel Erection	1

▼ Expected Output

name	address	phone	overall_review_rating
Casper, Oberbrunner and Williamson	53 Di Loreto Hill	+420 (569) 566-3689	2.0 (1 categories)
Tromp, Kozey and Abbott	84 Mcguire Plaza	+62 (145) 722-2330	1.7 (3 categories)
Wolff-Fadel	06 Fair Oaks Trail	+7 (894) 233-0976	1.7 (6 categories)

Gerlach, Hayes and Stamm	80 Service Point	+86 (731) 234-4119	1.5 (6 categories)
Kihn-Cronin	483 Nobel Road	+1 (396) 693-1661	1.3 (4 categories)

Question - 59

SQL: Domain Name Registrar Accounts Report

A domain name registration service needs new reporting functionality. Create a query that returns a list of all active accounts, the number of unexpired domain names they have, and the nearest expiration date later than July 15, 2022.

The result should have the following columns: *username / domains / nearest_expiration*.

- *username* - account username
- *domains* - total number of domains for a specific account
- *nearest_expiration* - nearest expiration date for domains for a specific account

The result should be sorted in ascending order by *username*.

Note:

- Only active accounts should be included in the report.
- Only domain names that have not expired (with an expiration date greater than today) should be included in the report.
- Today is July 15, 2022.

▼ Schema

accounts		
name	type	description
id	SMALLINT	unique id, primary key
username	VARCHAR(255)	
is_active	SMALLINT	Account status: 1 = Active, 0 = Not active

domains		
name	type	description
account_id	SMALLINT	foreign key, accounts.id
name	VARCHAR(255)	
expiration_date	VARCHAR(19)	

▼ Sample Data Tables

accounts		
id	username	is_active

1	obeedie0	0
2	stopham1	1
3	ndolder2	1
4	jyanshinov3	1
5	ewilflinger4	0

domains		
account_id	name	expiration_date
1	imgur.com	2022-05-14
1	domainmarket.com	2022-07-02
1	comsenz.com	2022-07-28
1	gizmodo.com	2022-08-09
1	toplist.cz	2022-08-15
1	scientificamerican.com	2022-09-03
1	examiner.com	2022-12-18
1	photobucket.com	2023-01-22
2	merriam-webster.com	2022-02-20
2	tripod.com	2022-08-08
3	ca.gov	2022-04-24
3	ehow.com	2022-06-28
3	purevolume.com	2022-07-01
3	squidoo.com	2022-10-27
3	eepurl.com	2022-12-21
4	digg.com	2022-05-14
4	jugem.jp	2022-08-05
4	artisteer.com	2022-10-21
5	behance.net	2022-03-24
5	cnn.com	2022-05-11

▼ Expected Output

username	domains	nearest_expiration
jyanshinov3	2	2022-08-05
ndolder2	2	2022-10-27
stopham1	1	2022-08-08

Question - 60

SQL: Advertising Network Events Report

As part of a marketing audit of your company's advertising campaigns published in the an online advertising network, they need a list of all advertising campaigns whose average event values as of July 15, 2022 are equal to or greater than 0.7.

The result should have the following columns: *campaign / events / average_value*.

- *campaign* - campaign name
- *events* - total number of events for a specific campaign
- *average_value* - average event value for a specific campaign, rounded to five decimal places (format 0##### where each # is a digit)

The result should be sorted in descending order by *average_value*.

Note:

- Only events for July 15, 2022, should be included in the report
- Only campaigns with average event values equal to or greater than 0.7 should be included in the report

▼ Schema

campaigns		
name	type	description
id	SMALLINT	unique id, primary key
name	VARCHAR(255)	

events		
name	type	description
campaign_id	SMALLINT	foreign key campaign.id
dt	VARCHAR(19)	Event datetime
value	DECIMAL(6,5)	Event value

▼ Sample Data Tables

campaigns

id	name
1	11-080 - Registration Equipment
2	12-700 - Systems Furniture
3	9-900 - Paints and Coatings

events		
campaign_id	dt	value
1	2022-07-14 13:11:38	0.59275
1	2022-07-14 14:55:43	0.12928
1	2022-07-14 18:16:11	0.82350
1	2022-07-15 01:19:44	0.97144
1	2022-07-15 22:52:02	0.60728
1	2022-07-16 08:55:38	0.71158
1	2022-07-16 10:22:44	0.29627
2	2022-07-14 02:36:31	0.42323
2	2022-07-14 04:45:32	0.91077
2	2022-07-14 07:24:11	0.35956
2	2022-07-15 06:43:08	0.16662
2	2022-07-16 08:21:27	0.02559
2	2022-07-16 11:59:41	0.34606
2	2022-07-16 23:26:12	0.62697
3	2022-07-14 00:21:56	0.97297
3	2022-07-14 10:22:11	0.93894
3	2022-07-14 12:29:59	0.44633
3	2022-07-15 01:17:41	0.37531
3	2022-07-15 14:20:48	0.24872
3	2022-07-16 23:02:51	0.80594

▼ Expected Output

campaign	events	average_value
-----------------	---------------	----------------------

Question - 61**SQL: Ecommerce Deal Report**

An online shop needs a new deal history feature. Create a query that returns a list of the top three seller profiles with the highest total deals in June, 2022.

The result should have the following columns: *first_name / last_name / email / total*.

- *total* - the total amount of all deals for a specific profile

The result should be sorted in descending order by *total*.

The result should be limited to the first three records.

Note:

- Only deals for June 2022 should be included in the report

▼ Schema

profiles		
name	type	description
id	SMALLINT	unique id, primary key
first_name	VARCHAR(255)	
last_name	VARCHAR(255)	
email	VARCHAR(255)	

deals		
name	type	description
profile_id	SMALLINT	foreign key into profile.id
dt	VARCHAR(19)	Deal datetime
amount	DECIMAL(5,2)	Deal amount

▼ Sample Data Tables

profiles			
id	first_name	last_name	email
1	Wallis	Treadway	wttreadway0@senate.gov

2	Franklin	Blackston	fblackston1@parallels.com
3	Honorina	Constant	hconstant2@umich.edu
4	Bertine	Hillaby	bhillaby3@artisteer.com
5	Constance	Knutsen	cknutschen4@google.ca

deals		
profile_id	dt	amount
5	2022-05-21 02:44:24	49.10
2	2022-05-22 23:26:59	46.21
1	2022-05-23 09:56:25	58.57
5	2022-05-28 02:38:08	27.81
4	2022-06-04 07:16:27	22.31
4	2022-06-04 14:15:03	36.33
5	2022-06-04 15:03:10	21.41
1	2022-06-07 02:58:06	92.84
4	2022-06-08 05:09:52	24.41
3	2022-06-13 03:28:52	61.55
4	2022-06-16 15:09:39	77.70
5	2022-06-18 16:51:32	58.79
4	2022-06-20 02:55:20	43.61
3	2022-06-22 06:52:10	10.41
1	2022-06-23 04:59:05	6.59
1	2022-06-30 16:11:02	43.07
4	2022-07-05 06:05:28	36.45
5	2022-07-12 07:49:51	14.76
4	2022-07-12 18:58:11	91.61
5	2022-07-14 00:50:45	69.61

▼ Expected Output

first_name	last_name	email	total
------------	-----------	-------	-------

Bertine	Hillaby	bhillaby3@artisteer.com	204.36
Wallis	Treadway	wtreadway0@senate.gov	142.50
Constance	Knutsen	cknutesen4@google.ca	80.20

Question - 62

SQL: Freelance Platform Candidate Review

A company's HR department needs an SQL query against data from a freelance hiring platform.

The result should have the following columns: *first_name / last_name / email / job_success_score*.

It should be sorted in descending order by *job_success_score*, then in ascending order by *first_name* and *last_name*.

Limit the result to the first ten records.

Note:

- Only profiles that have a "JSS" (Job Success Score) greater than 90 should be included.
- Only those profiles that have passed the verification process (*is_verified* is 1) should be included.

▼ Schema

profiles		
name	type	description
id	SMALLINT	unique id, primary key
first_name	VARCHAR(255)	
last_name	VARCHAR(255)	
email	VARCHAR(255)	
is_verified	SMALLINT	1 = True, 0 = False

stats		
name	type	description
profile_id	SMALLINT	Foreign key into profiles.id
job_success_score	SMALLINT	

▼ Sample Data Tables

profiles				
id	first_name	last_name	email	is_verified

1	Junia	Sehorsch	jsehorsch0@oracle.com	0
2	Dave	Halliburton	dhalliburton1@pbs.org	1
3	Agneta	Dutch	adutch2@thetimes.co.uk	1
4	Kendell	Sylvester	ksylvester3@canalblog.com	1
5	Koralle	Ragsdale	kragsdale4@buzzfeed.com	1
6	Roma	Kenelin	rkenelin5@prnewswire.com	1
7	Harold	Molloy	hmolloy6@ycombinator.com	1
8	Berri	Hartzogs	bhartzogs7@ox.ac.uk	0
9	Garrik	Preddle	gpreddle8@topsy.com	1
10	Sophie	Messenger	smessenger9@myspace.com	1
11	Ashby	Philipsson	aPhilipsona@typepad.com	1
12	Kayle	Jesteco	kjestecob@ocn.ne.jp	1
13	Munroe	Chevolleau	mchevolleauc@yandex.ru	1
14	Etheline	Choake	echoaked@hao123.com	1
15	Marten	Zamboniari	mzamboniarie@cbc.ca	1
16	Hersch	Blasdale	hblasdalef@wunderground.com	1
17	Jori	MacFaell	jmacfaellg@va.gov	1
18	Margo	Finnemore	mfinnemoreh@discovery.com	1
19	Felicle	Ramsdale	framsdalei@devhub.com	1
20	Demetris	Arnet	darnetj@livejournal.com	1

stats	
profile_id	job_success_score
4	100
1	100
18	95
5	95
8	95
10	95
12	95
13	95

14	95
20	90
17	90
7	90
15	90
11	85
16	85
9	85
6	85
2	85
3	75
19	75

▼ Expected Output

first_name	last_name	email	job_success_score
Kendell	Sylvester	ksylvester3@canalblog.com	100
Etheline	Choake	echoaked@hao123.com	95
Kayle	Jesteco	kjestecob@ocn.ne.jp	95
Koralle	Ragsdale	kragsdale4@buzzfeed.com	95
Margo	Finnemore	mfinnemoreh@discovery.com	95
Munroe	Chevolleau	mchevolleauc@yandex.ru	95
Sophie	Messenger	smessenger9@myspace.com	95
Demetris	Arnet	darnetj@livejournal.com	90
Harold	Molloy	hmolloy6@ycombinator.com	90
Jori	MacFaell	jmacfaellg@va.gov	90

Question - 63

SQL: Virtual Machine Deployment Report

While working on a new dashboard for a large cloud hosting company, an engineer needs a query that returns a list of CVM (Cloud Virtual Machine) configurations and their total number of deployments in 2021.

The result should have the following columns: *configuration / deployments*.

- *configuration* - configuration name
- *deployments* - the total number of all deployments for a specific configuration

Note:

- Only deployments in 2021 should be included in the report
- The result should be sorted in descending order by *deployments*.

▼ Schema

configurations		
name	type	description
id	VARCHAR(64)	unique id, primary key
name	VARCHAR(255)	

deployments		
name	type	description
configuration_id	VARCHAR(64)	foreign key configurations.id
dt	VARCHAR(19)	deployment date and time

▼ Sample Data Tables

configurations	
id	name
1vcpu_512mb_10gb_500gb	1 CPU / 512 MB RAM / 10 GB SSD Disk / 500 GB transfer
1vcpu_1gb_25gb_1tb	1 CPU / 1 GB RAM / 25 GB SSD Disk / 1000 GB transfer
1vcpu_2gb_50gb_2tb	1 CPU / 2 GB RAM / 50 GB SSD Disk / 2 TB transfer
2vcpu_2gb_60gb_3tb	2 CPUs / 2 GB RAM / 60 GB SSD Disk / 3 TB transfer
2vcpu_4gb_80gb_4tb	2 CPUs / 4 GB RAM / 80 GB SSD Disk / 4 TB transfer
4vcpu_8gb_160gb_5tb	4 CPUs / 8 GB RAM / 160 GB SSD Disk / 5 TB transfer
8vcpu_16gb_320gb_6tb	8 CPUs / 16 GB RAM / 320 GB SSD Disk / 6 TB transfer

deployments	
configuration_id	dt
1vcpu_512mb_10gb_500gb	2020-10-22 05:59:47
1vcpu_1gb_25gb_1tb	2020-11-09 06:07:57

1vcpu_2gb_50gb_2tb	2020-12-02 14:47:24
2vcpu_2gb_60gb_3tb	2021-01-24 15:41:42
4vcpu_8gb_160gb_5tb	2021-01-25 09:31:37
2vcpu_2gb_60gb_3tb	2021-02-08 02:43:14
1vcpu_512mb_10gb_500gb	2021-03-25 06:13:36
8vcpu_16gb_320gb_6tb	2021-03-26 22:23:42
2vcpu_2gb_60gb_3tb	2021-05-24 09:48:16
1vcpu_512mb_10gb_500gb	2021-05-31 05:03:28
1vcpu_2gb_50gb_2tb	2021-08-25 22:24:10
2vcpu_4gb_80gb_4tb	2021-09-05 22:12:17
2vcpu_4gb_80gb_4tb	2021-09-23 08:31:50
2vcpu_2gb_60gb_3tb	2021-10-14 08:26:20
1vcpu_2gb_50gb_2tb	2021-11-01 19:00:30
1vcpu_512mb_10gb_500gb	2021-11-26 10:53:05
1vcpu_1gb_25gb_1tb	2021-12-27 07:07:23
2vcpu_2gb_60gb_3tb	2022-02-13 06:00:29
2vcpu_2gb_60gb_3tb	2022-03-03 09:10:30

▼ Expected Output

configuration	deployments
2 CPUs / 2 GB RAM / 60 GB SSD Disk / 3 TB transfer	4
2 CPUs / 4 GB RAM / 80 GB SSD Disk / 4 TB transfer	3
1 CPU / 512 MB RAM / 10 GB SSD Disk / 500 GB transfer	3
1 CPU / 2 GB RAM / 50 GB SSD Disk / 2 TB transfer	2
8 CPUs / 16 GB RAM / 320 GB SSD Disk / 6 TB transfer	1
4 CPUs / 8 GB RAM / 160 GB SSD Disk / 5 TB transfer	1
1 CPU / 1 GB RAM / 25 GB SSD Disk / 1000 GB transfer	1

Question - 64

SQL: Visitors Behavior Report

You are working on developing visitor tracking software and need to get the total number of purchase events in May 2022. A purchase event has "buy" in the type field.

The result should have the following columns: *purchases*.

- *purchases* - the total number of events of the "buy" type

Note:

- Only events of the type "buy" should be included in the report.
- Only events that occurred in May 2022 should be included in the report.

▼ Schema

events		
name	type	description
dt	VARCHAR(19)	Event timestamp
type	VARCHAR(64)	Event type

▼ Sample Data Tables

events	
dt	type
2022-04-04 03:36:00	buy
2022-04-21 07:05:09	buy
2022-04-02 11:34:24	sell
2022-05-27 16:12:50	buy
2022-05-20 09:09:07	buy
2022-05-22 09:06:37	buy
2022-05-31 07:49:36	buy
2022-05-14 22:29:10	buy
2022-05-13 15:00:54	sell
2022-05-24 15:40:54	sell
2022-05-13 01:20:05	sell
2022-05-16 07:07:44	sell
2022-05-01 16:57:00	sell
2022-06-02 09:42:02	buy

2022-06-01 06:34:59	buy
2022-06-06 17:14:47	buy
2022-06-05 13:37:23	buy
2022-06-17 19:10:13	buy
2022-06-15 21:40:13	sell
2022-06-11 12:26:43	sell

▼ Expected Output

purchases
5

Question - 65

SQL: Advertising Campaigns Report

As part of the summary report on advertising campaigns, get a list of the top three companies whose campaigns have a positive return on their advertising budgets.

The result should have the following columns: *company_name / profit*.

- *profit* is the campaign revenues less expenses

The result should be sorted in descending order by *profit*.

The result should be limited to the three first records.

Note:

- Only companies whose campaigns have a positive return on their advertising budget should be included in the report

▼ Schema

companies		
name	type	description
id	SMALLINT	ID
name	VARCHAR(255)	Company name

campaigns		
name	type	description
company_id	SMALLINT	Company ID

expenses	DECIMAL(7,2)	Campaign expenses
revenue	DECIMAL(7,2)	Campaign revenue

▼ Sample Data Tables

companies	
id	name
1	Lion Biotechnologies, Inc.
2	Boston Private Financial Holdings, Inc.
3	Universal Corporation
4	Arbutus Biopharma Corporation
5	Royal Bank Of Canada
6	Penn West Petroleum Ltd
7	Public Storage
8	Halcon Resources Corporation
9	TTM Technologies, Inc.
10	Atwood Oceanics, Inc.
11	ACADIA Pharmaceuticals Inc.
12	Central European Media Enterprises Ltd.
13	Oxbridge Re Holdings Limited
14	Western Refining Logistics, LP
15	Vaalco Energy Inc
16	Xilinx, Inc.
17	Liberty Global plc
18	Honda Motor Company, Ltd.
19	Great Plains Energy Inc
20	Assurant, Inc.

campaigns		
company_id	expenses	revenue
1	7390.24	8652.18

2	5774.65	7955.47
3	2154.71	5920.23
4	9366.49	3397.85
5	2765.18	9158.63
6	7908.41	5018.85
7	2251.44	6654.52
8	3383.14	9354.79
9	8287.96	9522.53
10	4356.62	4658.52
11	9272.86	9161.77
12	4996.18	5903.57
13	8354.75	2259.26
14	6402.90	8146.16
15	1692.05	686.71
16	5988.48	9089.41
17	6192.33	7580.19
18	3016.37	7761.25
19	9838.05	1293.09
20	4386.52	9513.73

▼ Expected Output

company_name	profit
Royal Bank Of Canada	6393.45
Halcon Resources Corporation	5971.65
Assurant, Inc.	5127.21

Question - 66

SQL: Traffic Audit Report

As part of the traffic audit report for one of the ISPs, you need to get a list of customers whose hardware has a higher downstream rate than upstream rate, but no more than one downtime.

The result should have the following columns: *mac / upstream_rate / downstream_rate / downtime_rate*.

The result should be sorted in ascending order by *mac*.

Note:

- Only clients whose downstream rate is higher than the upstream rate should be included in the report.
- Only clients whose downtime ratio is "never" or "once" should be included in the report.

▼ Schema

clients		
name	type	description
mac	VARCHAR(64)	MAC address
upstream_rate	INT	Upstream rate
downstream_rate	INT	Downstream rate
downtime_rate	VARCHAR(64)	Downtime rate

▼ Sample Data Tables

clients			
mac	upstream_rate	downstream_rate	downtime_rate
78-C1-E5-20-D5-61	925526	5195	never
78-E2-20-71-9C-30	582152	375829	never
0D-09-F7-77-03-E5	359529	710743	never
56-18-67-55-58-EA	78626	562544	once
C9-73-EC-1C-4C-B7	574927	669655	yearly
02-35-3F-7B-CC-76	430072	296196	seldom
3D-95-33-8A-65-F9	894176	489401	monthly
8D-33-3F-0E-04-D5	897666	297063	weekly
0E-0A-63-B9-79-3E	69133	984354	seldom
15-D3-2A-DD-02-A4	19203	995983	seldom
8D-42-B1-97-AB-87	476648	177677	monthly
6B-64-60-47-16-D3	700056	374321	monthly
FD-61-81-00-BF-EC	216401	498229	yearly
95-46-C6-C7-6F-E0	236331	341013	monthly

11-0E-62-32-62-5E	694746	451525	daily
D6-5B-72-D5-FF-4F	2931	992852	monthly
66-BF-AD-F2-E5-45	861075	44216	monthly
E5-C9-C6-74-2E-A8	639487	968494	daily
18-56-3A-93-8E-9F	494945	259910	weekly
51-EB-D7-22-45-99	219419	326479	often

▼ Expected Output

mac	upstream_rate	downstream_rate	downtime_rate
0D-09-F7-77-03-E5	359529	710743	never
56-18-67-55-58-EA	78626	562544	once

Question - 67

SQL: Calendar Application Events Report

You are working on a calendar application and need to get a list of the five earliest events for all event owners who are not on vacation.

The result should have the following columns: *dt / title / full_name / email_address*.

It should be sorted in ascending order by *dt*, and limited to the first five records.

Note:

- Only events whose owner is not on vacation should be included in the report.

▼ Schema

owners		
name	type	description
id	SMALLINT	ID
full_name	VARCHAR(255)	Full name
email_address	VARCHAR(255)	Email address
on_vacation	SMALLINT	Vacation status, 1 = on vacation

events		
name	type	description
owner_id	SMALLINT	Owner ID

dt	VARCHAR(19)	Timestamp
title	VARCHAR(255)	Title

▼ Sample Data Tables

owners			
id	full_name	email_address	on_vacation
2	Aleksandr Fellows	afellows1@instagram.com	0
3	Collete Pack	cpack2@mit.edu	0
4	Lorelle Squibb	lsquibb3@huffingtonpost.com	0
1	Benjamin Sevier	bsevier0@discuz.net	1
5	Cointon Welberry	cwelberry4@theguardian.com	1

events		
owner_id	dt	title
1	2021-08-13 09:17:41	ac consequat metus sapien ut
1	2021-10-26 16:14:28	a libero nam
1	2022-05-18 16:09:36	vestibulum ante ipsum primis in faucibus
1	2022-06-04 20:37:36	nisi volutpat eleifend donec
2	2021-01-17 23:34:59	eu mi nulla ac
2	2021-01-26 15:59:04	nisl venenatis lacinia
2	2021-04-25 05:04:29	arcu adipiscing molestie hendrerit
2	2022-02-22 10:24:50	lectus vestibulum quam sapien varius
2	2022-05-22 03:04:33	pellentesque quisque porta volutpat erat
3	2021-03-11 08:25:21	felis sed interdum venenatis
3	2022-02-15 07:29:45	nulla sed accumsan felis ut
4	2021-02-28 16:10:55	cras mi pede malesuada in imperdiet et
4	2021-05-20 03:35:11	odio in hac habitasse platea dictumst
4	2021-09-01 19:25:41	sagittis nam congue risus semper
4	2022-02-04 03:11:53	luctus cum sociis natoque penatibus et magnis
4	2022-06-18 08:39:41	nec molestie sed justo pellentesque

5	2021-04-04 13:39:38	id massa id nisl venenatis lacinia aenean
5	2021-10-18 19:04:14	adipiscing lorem vitae mattis
5	2021-10-19 00:56:11	vestibulum eget vulputate ut ultrices vel augue
5	2022-04-26 18:55:04	at dolor quis

▼ Expected Output

dt	title	full_name	email_address
2021-01-17 23:34:59	eu mi nulla ac	Aleksandr Fellows	afellows1@instagram.com
2021-01-26 15:59:04	nisl venenatis lacinia	Aleksandr Fellows	afellows1@instagram.com
2021-02-28 16:10:55	cras mi pede malesuada in imperdiet et	Lorelle Squibb	lsquibb3@huffingtonpost.com
2021-03-11 08:25:21	felis sed interdum venenatis	Collete Pack	cpack2@mit.edu
2021-04-25 05:04:29	arcu adipiscing molestie hendrerit	Aleksandr Fellows	afellows1@instagram.com

Question - 68

Firewall Active Clients Tracking

As part of a firewall log analysis process, a team needs a list of all the client's MAC addresses that have had traffic.

The result should be in the following format: *mac*.

- *mac* is a client MAC address, which was active and has traffic (its *clients.id* exists in the *traffic* table).

▼ Schema

There are 2 tables:

clients		
name	type	description
id	SMALLINT	Client ID
mac	VARCHAR(64)	MAC address

traffic		
name	type	description
client_id	SMALLINT	Client ID
amount	INT	Traffic amount

▼ Sample Data Tables

For the sample data in tables:

clients	
id	mac
1	E5-A3-AC-8A-20-F9
2	3B-2F-83-25-A8-81
3	3A-4E-A6-43-1D-B1
4	B7-03-14-91-8F-58
5	63-1A-FD-9A-AF-6F
6	D6-B8-1F-1D-34-04
7	38-83-E5-F8-C8-DC
8	E9-F6-89-7F-8D-34
9	82-2E-B3-67-04-41
10	D1-9D-DE-37-A0-49
11	8A-46-F4-83-29-13
12	4C-1F-7B-C7-08-7E
13	72-57-E6-CA-2C-91
14	57-F7-E7-E7-45-36
15	05-8A-05-1D-2D-20
16	46-06-F1-B9-65-7C
17	0A-E0-26-9D-2A-27
18	F0-86-99-18-36-9B
19	DD-81-BF-53-BD-9B
20	50-53-64-8E-42-BE

traffic	
client_id	amount
3	6385047
8	6490817
14	9109219

16	5558512
17	1870152
17	8228920
18	326127
18	5429741
18	4063477
19	7411789
20	5832337
20	1426585
23	1097368
23	6769594
23	802387
24	5959513
24	1300408
24	4631624
28	1629306
29	2814818

the expected output is:

mac
0A-E0-26-9D-2A-27
3A-4E-A6-43-1D-B1
46-06-F1-B9-65-7C
50-53-64-8E-42-BE
57-F7-E7-E7-45-36
DD-81-BF-53-BD-9B
E9-F6-89-7F-8D-34
F0-86-99-18-36-9B

Explanation

Each of the addresses listed has at least one record in the traffic table.

Question - 69

Active Wallets

As part of an Etherium transactions monitoring process, a team needs a list of all the wallets that have one or more transactions.

The result should be in the following format: *address*

- *address* is a wallet address which has one or more transactions (its *wallet.id* exists in the *transactions* table).

▼ Schema

There are 2 tables:

wallets		
name	type	description
id	SMALLINT	Wallet ID
address	VARCHAR(64)	Wallet address

transactions		
name	type	description
wallet_id	SMALLINT	Wallet ID
credit	DECIMAL(4,2)	Transaction amount

▼ Sample Data Tables

For the sample data in tables:

wallets	
id	address
1	0x1ecc4cefde6dfb773352a2dcd8b5f518ccd24ff4
2	0x0f4487168610dcae7f16b6c000a7ba284bb6703c
3	0x54c5c516b5c601a3e6fdea18db966186274879e6
4	0x50c404a6790d44849c3500d95f147d2102db5f77
5	0x41cc32037fdcc6cbf5cb04ada4e38032a84361ac
6	0xb38ca2076b2c3a0a1c0796ccab5dc3fcbb645336
7	0x923c6ee4debce6c37c512eb90e9b26fc9bbfc3e1
8	0x76cfbf67e1765117a98fd2286bdbcb731b0d29ec

9	0xf4dff43df89365453440dbeb5e53445d3400d218
10	0x010a4b23f985eda6b397864878cb70bf0b233aa6
11	0x106c4c852394d5a6580c11ca0060687a798ec78c
12	0x541da74b9f4b2283524a65156263bcce7429ba61
13	0x86df8e62ad23a35ee2e64b3b4b128d8a3660116a
14	0x6b021b12b779aec27009f18bd6ad227685df5474
15	0x7073e9f48b7211d9940db4c8bad8e31d5d9d0577
16	0x4672e132b9627a7db76e5af431bb5febc93b1b2f
17	0x5cbc82a01d3b9df8f50ee09bf5a1c48afe0cb966
18	0x1c7d30f90081b45004e0d6549534ef8658795672
19	0xf830618ca9f027ce5f23bf270629a6a418ad355c
20	0x1a31ac1b923fcc17a42a340091424ea18192a3e7

transactions	
wallet_id	credit
3	5.21
6	24.04
8	29.66
10	1.67
10	3.32
10	4.27
14	3.32
14	3.35
15	14.34
16	11.94
20	13.86
20	16.54
21	11.64
25	12.86
26	8.77
26	12.36

27	29.65
28	3.00
28	13.26
28	2.85

the expected output is:

address
0x010a4b23f985eda6b397864878cb70bf0b233aa6
0x1a31ac1b923fcc17a42a340091424ea18192a3e7
0x4672e132b9627a7db76e5af431bb5febc93b1b2f
0x54c5c516b5c601a3e6fdea18db966186274879e6
0x6b021b12b779aec27009f18bd6ad227685df5474
0x7073e9f48b7211d9940db4c8bad8e31d5d9d0577
0x76cfbf67e1765117a98fd2286bdbcb731b0d29ec
0xb38ca2076b2c3a0a1c0796ccab5dc3fcbb645336

Explanation

Each of the wallets listed has at least one record in the transactions table.

Question - 70 Animal Tracking

As part of an animal tracking program, a team needs a list of all the animals, tracked at least once.

The result should be in the following format: *name*.

- *name* is the name of an animal, which has at least one record of being tracked (its *animal.id* exists in the *tracklog* table).

▼ Schema

There are 2 tables:

animals		
name	type	description
id	SMALLINT	Animal ID
name	VARCHAR(64)	Animal name

tracklog		
name	type	description
animal_id	SMALLINT	Animal ID
tracked_at	VARCHAR(19)	Track timestamp

▼ Sample Data Tables

For the sample data in tables:

animals	
id	name
1	Blue and yellow macaw
2	Jungle kangaroo
3	Stork, woolly-necked
4	North American river otter
5	Square-lipped rhinoceros
6	Black-fronted bulbul
7	American beaver
8	Capybara
9	Black-backed jackal
10	Dragon, ornate rock
11	Wombat, southern hairy-nosed
12	Snake, carpet
13	Egyptian cobra
14	Green heron
15	Indian star tortoise
16	Roan antelope
17	Rhea, common
18	Fairy penguin
19	Black-eyed bulbul
20	Starling, cape

tracklog	
animal_id	tracked_at
2	2021-07-08 12:30:34
5	2021-09-15 03:00:04
8	2021-12-14 11:20:50
9	2021-05-15 14:54:15
11	2021-07-03 17:04:14
11	2021-02-23 09:02:11
15	2021-11-02 22:49:37
15	2021-07-31 04:51:25
15	2021-06-06 14:57:23
19	2021-03-25 02:00:31
19	2021-04-21 22:02:51
19	2021-07-17 17:12:35
21	2021-05-06 03:35:59
22	2021-09-21 22:49:36
22	2021-12-04 00:26:29
24	2021-12-27 08:19:52
26	2021-05-14 07:30:11
28	2021-06-08 00:47:00
28	2021-04-24 00:42:31
30	2021-05-06 21:03:03

the expected output is:

name
Black-backed jackal
Black-eyed bulbul
Capybara
Indian star tortoise
Jungle kangaroo
Square-lipped rhinoceros

Explanation

Each of the animals listed has at least one record in the tracklog table.

Question - 71**SQL: Exchange Rates**

In exchange for a transaction fee, a stock exchange lets people trade stocks online. For every buy order they charge 0.1% of the order amount, and for every sell order they charge 0.15%. Calculate the fees paid by customers for their orders.

Fetch the name of the customer and their total fees paid. Order the result by customer names and round the fees to 2 decimal places.

To round the data:

In MySQL, Oracle, or DB2, use ROUND(val, 2).

In MS SQL use FORMAT(val, 'N2').

▼ Schema

You are provided 2 tables: *customers*, *orders*.

CUSTOMERS		
Name	Type	Description
id	int	Unique id of the customer.
customer_name	varchar(30)	Name of the customer.

ORDERS		
Name	Type	Description
order_id	int	Unique id of the order.
customer_id	int	Id of the customer referring to the customers table.
order_type	varchar(5)	Type of order placed (Buy or Sell).
order_amount	decimal(18,2)	Amount of the order.

▼ Sample Data Tables

CUSTOMERS	
id	customer_name
401	Hubert Keesler
402	Devin Vert
403	Lashawna Bowerman

404	Brigid Wellborn
405	Josefine Perl

ORDERS			
order_id	customer_id	order_type	order_amount
4361	401	Sell	912.77
3478	405	Sell	741.69
7292	405	Sell	436.05
5833	405	Sell	231.30
3472	402	Buy	950.92
4472	401	Sell	367.70
2624	404	Buy	218.15
7198	405	Buy	797.29
7660	403	Buy	131.18
5192	401	Buy	362.44
5260	402	Buy	636.26
2726	403	Sell	138.15
6594	401	Buy	234.51
4657	404	Buy	427.30
9744	402	Sell	623.36

OUTPUT	
customer_name	total_fees
Brigid Wellborn	0.65
Devin Vert	2.52
Hubert Keesler	2.52
Josefine Perl	2.91
Lashawna Bowerman	0.34

Explanation

For Devin Vert, total amount of buy orders was 1587.18 and total amount of buy orders was 623.36. Fee applied on buy and sell orders will be 1.58718 and 0.93504 with rate of 0.1% and 0.15% respectively. Total fee would be 2.52222 after rounding final fee would be 2.52.

Question - 72
SQL: Credit Dues

A credit card issuer sets its customers' interest rates based on their credit scores. Given details of customer transactions, calculate the amount due for every customer based on their individual interest rates.

Write a query to fetch the full name of the customer along with the amount due (i.e. total transaction amount + interest on that amount) for the customers whose interest rate is more than 12%. Round the amounts to 2 places after the decimal and order them from high to low.

▼ Schema

There are 2 tables: `credit_holders`, `transactions`.

CREDIT_HOLDERS		
Name	Type	Description
id	int	Unique id of the credit holder (customer).
first_name	varchar(15)	First name of the customer.
last_name	varchar(15)	Last name of the customer.
interest_rate	int	Interest based on their credit score (in percentage).

TRANSACTIONS		
Name	Type	Description
transaction_id	int	Unique id of the transaction.
credit_holder_id	int	Id of the customer referring to the credit_holders table.
amount	decimal(18,2)	Amount of the transaction.

▼ Sample Data Tables

CREDIT_HOLDERS			
id	first_name	last_name	interest_rate
101	Clemencia	Hutsell	12
102	Susannah	Ismail	18
103	Sixta	Hagy	18
104	Otto	Izquierdo	18
105	Anita	Degroot	15

TRANSACTIONS		
transaction_id	credit_holder_id	amount
4361	101	65.22
3478	104	51.85
7292	104	64.60

5833	105	72.15
3472	102	96.28
4472	101	80.06
2624	101	85.27
7198	104	23.73
7660	103	81.86
5192	101	69.64
5260	101	71.72
2726	102	57.66
6594	103	23.23
4657	101	81.68
9744	104	99.57
2054	103	51.13
7156	105	12.78
3273	105	36.15
9756	101	45.41
9702	105	69.75

OUTPUT	
full_name	dues
Otto Izquierdo	282.91
Anita Degroot	219.45
Sixta Hagy	184.34
Susannah Ismail	181.65

Explanation

For Otto Izquierdo, there are 4 transactions of 51.85, 64.60, 23.73, 99.57 which sums to 239.75. After applying 18% interest, the due amount is 282.905 rounded to 282.91.

Question - 73

SQL: Interest Earned

A bank offers 5% interest on deposits. Given account holder names and amounts deposited, calculate the interest earned in 1 year. Report the depositor names and the interest they earn ordered ascending by name.

Notes:

The *amount* field is a string where the first character is a currency symbol, and the remainder is the deposit amount.

▼ Schema

There is 1 table: accounts.

ACCOUNTS		
Name	Type	Description
id	int	Unique id of the account.
account_holder	varchar(30)	Name of the account holder.
amount	varchar(10)	Amount deposited in the fixed deposit.

▼ Sample Data Tables

ACCOUNTS		
id	account_holder	amount
1	Ellis Beane	\$5582.03
2	Drew Nolf	\$2470.3
3	Jordan Chatmon	\$6211.52
4	Robin Hansard	\$8133.31
5	Spencer Days	\$5273.81
6	Morgan Criss	\$3741.85
7	Wesley Waugh	\$7056.14
8	Alex Cantly	\$2590.45
9	Blake Hawbaker	\$3987.27
10	Taylor Blackston	\$8351.98

OUTPUT	
account_holder	interest
Taylor Blackston	\$417.60
Wesley Waugh	\$352.81
Jordan Chatmon	\$310.58
Robin Hansard	\$406.67
Alex Cantly	\$129.52
Drew Nolf	\$123.52
Blake Hawbaker	\$199.36
Spencer Days	\$263.69

Ellis Beane	\$279.10
Morgan Criss	\$187.09

Explanation

For Taylor Blackston, after 1 year the interest is $\$8351.98 * 0.05 = \417.599 which rounds to \$417.60.

Question - 74

SQL: Monthly Revenue

A company has been selling beauty products for the last three years. They need a report of transactions by month over the period. The transaction id contains the year and month of the transaction, e.g. 19JAN1092 was executed in January of 2019.

Write a query to fetch year, month, and total transactions in each month. Order the results ascending by year and month.

▼ Schema

There is 1 table: `transactions`.

TRANSACTIONS		
Name	Type	Description
transaction_id	varchar(10)	Unique id of the transaction YYMMMD[the rest]
amount	decimal(18,2)	Amount of the transaction.

▼ Sample Data Tables

TRANSACTIONS	
transaction_id	amount
19SEP2187	785.72
19OCT4361	752.64
19APR3478	197.92
19SEP7292	910.26
21MAR5833	344.70
20MAY3472	939.61
20DEC4472	154.98
20DEC2624	935.44
21JUN7198	309.81
19APR7660	528.10
20MAR5192	995.22
19OCT5260	861.11

21JUN2726	611.94
19OCT6594	478.54
19APR4657	183.20

OUTPUT		
year	month	total_transactions
19	APR	909.22
19	OCT	2092.29
19	SEP	1695.98
20	DEC	1090.42
20	MAR	995.22
20	MAY	939.61
21	JUN	921.75
21	MAR	344.70

Explanation

For the month of April, 19 there are three transactions of 197.92, 528.10, and 183.20 which sum up to 909.22.

Question - 75

SQL: Final Result

A college offers a 4-year B.Tech degree. Each year students receive grades out of 10 (CGPA).

Write a query to fetch the full name of each student with the average GPA over the 4 years. Round the average to one decimal place. Order the result from highest to lowest average GPA.

▼ Schema

There is 1 table: `results`.

RESULTS		
Name	Type	Description
id	int	Unique id of the student.
first_name	varchar(20)	First name of the student.
last_name	varchar(20)	Last name of the student.
cgpa_first_year	float	CGPA year 1
cgpa_second_year	float	CGPA year 2
cgpa_third_year	float	CGPA year 3

▼ Sample Data Tables

RESULTS						
id	first_name	last_name	cgpa_first_year	cgpa_second_year	cgpa_third_year	cgpa_fourth_year
1	Pearlene	Beane	7	5.1	8.4	8.9
2	Franklin	Nolf	7.7	7.2	5.2	8.3
3	Bell	Chatmon	7.3	8.4	8.9	10
4	Belva	Hansard	6.2	9.2	5.8	6.7
5	Missy	Days	8.3	10	7.3	6.7
6	Vicenta	Criss	5.4	9.5	6.1	9
7	Annelle	Waugh	6.5	7.9	9.6	9.3
8	Darby	Canty	5.5	9	8.6	5.9
9	Ka	Hawbaker	5.7	6.4	5.2	6.8
10	Alease	Blackston	5.3	7.5	9.3	6

OUTPUT	
full_name	average_gpa
Bell Chatmon	8.6
Annelle Waugh	8.3
Missy Days	8.1
Vicenta Criss	7.5
Pearlene Beane	7.3
Darby Canty	7.3
Franklin Nolf	7.1
Belva Hansard	7.0
Alease Blackston	7.0
Ka Hawbaker	6.0

Explanation

For Bell Chatmon, the cgpa sum is $7.3 + 8.4 + 8.9 + 10 = 34.6$. The average is $34.6/4 = 8.65$. After rounding to 1 decimal place the final value is 8.6. All values should show 1 place after the decimal, even in a case like Belva Hansard where the final value should be 7.0.

Question - 76

SQL: Mutual Funds

A bank offers several mutual funds to investors. They receive investments each month and would like a summary.

Write a query to fetch the month, name of the fund, and the total amount invested that month. Order the report by month and mutual fund name.

▼ Schema

There is 1 table: `funds`.

FUNDS		
Name	Type	Description
id	int	Unique id of the order
order_date	date	transaction date
fund_name	varchar(50)	Name of the fund
order_amount	int	Amount invested

▼ Sample Data Tables

FUNDS			
id	order_date	fund_name	order_amount
1	2021-09-05	Mid-Cap	151
2	2020-11-27	Small-Cap	784
3	2020-11-22	Multi-Cap	761
4	2020-02-26	Large-Cap	778
5	2020-01-04	Mid-Cap	949
6	2020-02-01	Large-Cap	392
7	2020-02-07	Mid-Cap	629
8	2020-06-01	Mid-Cap	529
9	2020-08-05	Large-Cap	258
10	2021-09-23	Mid-Cap	739
11	2020-02-14	Large-Cap	563
12	2021-09-29	Small-Cap	817
13	2020-05-11	Large-Cap	121
14	2021-09-18	Mid-Cap	341
15	2021-09-07	Large-Cap	260
16	2021-10-11	Small-Cap	102
17	2020-06-04	Mid-Cap	496
18	2021-07-26	Multi-Cap	321

19	2020-04-11	Small-Cap	323
20	2020-09-24	Multi-Cap	938
21	2020-06-16	Small-Cap	668
22	2021-03-08	Small-Cap	733
23	2021-04-27	Small-Cap	597
24	2020-01-11	Small-Cap	767
25	2021-11-14	Small-Cap	489

OUTPUT		
month	fund_name	total_investments
1	Mid-Cap	949
1	Small-Cap	767
2	Large-Cap	1733
2	Mid-Cap	629
3	Small-Cap	733
4	Small-Cap	920
5	Large-Cap	121
6	Mid-Cap	1025
6	Small-Cap	668
7	Multi-Cap	321
8	Large-Cap	258
9	Large-Cap	260
9	Mid-Cap	1231
9	Multi-Cap	938
9	Small-Cap	817
10	Small-Cap	102
11	Multi-Cap	761
11	Small-Cap	1273

Explanation

For the 1st month the total investment for the Mid-Cap fund is 949 and for Small-Cap is 767. Since there is no Large-Cap or Multi-Cap fund investments that month, they are not shown.

Question - 77
DPI Software Protocols Report

As part of HackerSniff's DPI (Deep Packet Inspection) software analytics, a team needs a list of all the protocols for which incoming traffic is higher than outgoing.

The result should be in the following format: *protocol*, *traffic_in*, *traffic_out*.

- Results should be sorted ascending by *protocol*.

▼ Schema

There is 1 table:

traffic		
name	type	description
client	VARCHAR(17)	Client MAC address
protocol	VARCHAR(64)	Protocol name
traffic_in	INT	Traffic in
traffic_out	INT	Traffic out

▼ Sample Data Tables

For the sample data in table:

traffic			
client	protocol	traffic_in	traffic_out
02-E1-80-76-EC-4B	BGP	0	234737
43-15-AA-26-0F-A4	BGP	402860	606565
90-E7-B0-14-7E-8C	BGP	840772	988197
FB-60-23-C1-5E-D6	DNS	341155	356569
4D-6D-7F-62-F4-00	FTP	8346	413322
09-89-26-46-C4-21	FTP	210656	470568
B1-6A-35-2F-1A-C2	FTP	897097	161083
0C-CA-68-2D-4B-F5	HTTP	918793	550403
A4-C6-52-10-2E-9C	HTTPS	520856	185387
95-B8-7D-78-06-42	POP	150880	423073
B9-C1-1B-32-55-95	POP	862946	979544
14-FD-21-F6-5E-67	SMTP	139389	280646
70-E1-2D-B1-B2-9B	SMTP	163986	450401

C6-F1-59-FF-5D-BE	SMTP	271295	878246
62-01-CF-AD-32-A7	SMTP	388933	81625
41-80-FB-86-D1-93	SMTP	752842	253981
93-3F-01-57-5F-4A	SSH	496717	599280
52-F2-BF-45-84-74	SSH	632534	128765
87-66-B5-A5-2F-7B	SSH	835441	354950
CE-FC-80-F3-95-58	UDP	903443	120298

the expected output is:

protocol ▲	traffic_in	traffic_out
FTP	1116099	1044973
HTTP	918793	550403
HTTPS	520856	185387
SSH	1964692	1082995
UDP	903443	120298

Question - 78

Advertising System Failures Report

As part of HackerAd's advertising system analytics, a team needs a list of customers who have a maximum number of failure events (*status* = "failure") in their campaigns.

For all customers with more than 3 events with *status* = 'failure', report the customer name and their number of failures.

The result should be in the following format: *customer, failures*.

- *customer* is a candidate's full name, the *first_name* and *last_name* separated by a single space.
- The order of the output is not important.

▼ Schema

There are 3 tables:

customers		
name	type	description
id	SMALLINT	Customer ID
first_name	VARCHAR(64)	Customer first name
last_name	VARCHAR(64)	Customer last name

campaigns		
name	type	description
id	SMALLINT	Campaign ID
customer_id	SMALLINT	Customer ID
name	VARCHAR(64)	Campaign name

events		
name	type	description
dt	VARCHAR(19)	Event timestamp
campaign_id	SMALLINT	Campaign ID
status	VARCHAR(64)	Event status

▼ Sample Data Tables

For the sample data in tables:

customers		
id	first_name	last_name
1	Whitney	Ferrero
2	Dickie	Romera

campaigns		
id	customer_id	name
1	1	Upton Group
2	1	Roob, Hudson and Rippin
3	1	McCullough, Rempel and Larson
4	1	Lang and Sons
5	2	Ruecker, Hand and Haley

events		
dt	campaign_id	status
2021-12-02 13:52:00	1	Failure

2021-12-02 08:17:48	2	failure
2021-12-02 08:18:17	2	failure
2021-12-01 11:55:32	3	failure
2021-12-01 06:53:16	4	failure
2021-12-02 04:51:09	4	failure
2021-12-01 06:34:04	5	failure
2021-12-02 03:21:18	5	failure
2021-12-01 03:18:24	5	failure
2021-12-02 15:32:37	1	success
2021-12-01 04:23:20	1	success
2021-12-02 06:53:24	1	success
2021-12-02 08:01:02	2	success
2021-12-01 15:57:19	2	success
2021-12-02 16:14:34	3	success
2021-12-02 21:56:38	3	success
2021-12-01 05:54:43	4	success
2021-12-02 17:56:45	4	success
2021-12-02 11:56:50	4	success
2021-12-02 06:08:20	5	success

the expected output is:

customer	failures ▼
Whitney Ferrero	6

Question - 79 Election Exit Poll Report

As part of HackerPoll's election exit poll analytics, a team needs a list of candidates with votes they received.

The result should be in the following format: *candidate, votes*.

- *candidate* is a candidate's full name, the *first_name* and *last_name* separated by a single space.
- *votes* is a total count of all *vote_at* of the candidate.
- Votes that cannot be matched to a candidate should be ignored.
- Results should be sorted descending by *votes*, then ascending by *candidate*.

▼ Schema

There are 2 tables:

candidates		
name	type	description
id	SMALLINT	Candidate ID
first_name	VARCHAR(64)	Candidate first name
last_name	VARCHAR(64)	Candidate last name

results		
name	type	description
candidate_id	SMALLINT	Candidate ID
vote_at	VARCHAR(19)	Vote timestamp

▼ Sample Data Tables

For the sample data in tables:

candidates		
id	first_name	last_name
1	Xavier	Ping
2	Westley	Drewell
3	Dominick	Scoble

results	
candidate_id	vote_at
0	2021-12-01 14:15:52
1	2021-12-01 03:55:23
1	2021-12-01 21:53:26
1	2021-12-02 07:57:40
1	2021-12-02 13:56:06
2	2021-12-01 11:46:40
2	2021-12-01 14:56:05

2	2021-12-01 21:54:50
2	2021-12-02 00:43:18
2	2021-12-02 06:59:33
2	2021-12-02 08:36:35
2	2021-12-02 10:20:33
2	2021-12-02 14:02:38
3	2021-12-01 05:18:34
3	2021-12-02 03:55:37
3	2021-12-02 05:30:24
3	2021-12-02 08:32:06
4	2021-12-02 05:05:55
5	2021-12-02 15:50:50
5	2021-12-02 20:45:08

the expected output is:

candidate	votes ▼
Westley Drewell	8
Xavier Ping	4
Dominick Scoble	4

Question - 80 Billing Analytics Customer Report

As part of HackerPay's billing analytics, a team needs a list of MVCs (Most Valued Customers) and their total transactions in December. An MVC is a customer with 3 or more transactions in a month.

List all customers who are MVC in December. For each MVC, include the customer's name (*customer*), their transactions count, and the sum of their transactions *amount* for the month.

The result should be in the following format: *customer, transactions, total*.

Sort the result ascending by *customer*.

▼ Schema

There is 1 table:

events		
name	type	description
dt	VARCHAR(19)	Transaction timestamp
customer	VARCHAR(64)	Customer name
amount	DECIMAL(5,2)	Transaction amount

▼ Sample Data Tables

For the sample data in table:

events		
dt	customer	amount
2021-11-22 06:41:01	Donaugh Furneaux	0.89
2021-12-22 20:07:04	Donaugh Furneaux	10.51
2021-12-31 05:22:11	Donaugh Furneaux	55.92
2021-12-12 21:26:42	Harley Lyddiard	37.68
2021-11-22 21:24:30	Kippy Jelly	85.87
2021-11-25 07:00:29	Kippy Jelly	7.25
2021-12-16 16:48:32	Kippy Jelly	65.49
2021-11-22 23:30:55	Latrina Jackman	93.49
2021-11-24 19:38:52	Latrina Jackman	82.28
2021-11-30 22:59:33	Latrina Jackman	96.87
2021-12-30 13:05:34	Latrina Jackman	88.19
2021-11-22 02:08:02	Maribel Braim	20.19
2021-12-13 00:14:58	Maribel Braim	97.99
2021-12-26 13:22:20	Maribel Braim	57.06
2021-12-29 00:20:27	Maribel Braim	24.35
2021-11-25 14:29:29	Orrin Curley	6.69
2021-12-08 06:22:16	Orrin Curley	36.85
2021-12-09 15:32:16	Orrin Curley	11.04
2021-11-28 00:15:20	Rasla Venny	14.59
2021-12-25 09:58:23	Rasla Venny	6.41

the expected output is:

customer ▲	transactions	total
Maribel Braim	3	179.40

Question - 81

Aggregate Marks

There is a database containing the marks of some students in various subjects. The data may contain any number of subjects for a student.

Retrieve the records of students who have a sum of marks greater than or equal to 500. The result should be in the following format: *STUDENT_ID SUM_OF_MARKS* sorted descending by *STUDENT_ID*.

▼ Schema

There is 1 table: `marks` .

marks		
Name	Type	Description
STUDENT_ID	INTEGER	This is the student's unique ID.
MARKS	INTEGER	These are the marks obtained.

▼ Sample Data Table

marks	
STUDENT_ID	MARKS
1	450
2	200
3	260
2	300
3	250

Sample Output

```
3 510
2 500
```

Explanation

- 3 has a sum of **510**, so it is printed in the query results.
- 1 has a sum of **450**, so it does not get printed in the query results.
- 2 has a sum of **500**, so it is printed in the query results.

Question - 82

Trip Query

A travel and tour company has 2 tables that relate to customers: *FAMILIES* and *COUNTRIES*. Each tour offers a discount if a minimum number of people book at the same time.

Write a query to print the maximum number of discounted tours any 1 family in the *FAMILIES* table can choose from.

▼ Schema

There are 2 tables: *FAMILIES* , *COUNTRIES*.

FAMILIES		
Name	Type	Description
ID	STRING	Unique ID of the family.
NAME	STRING	Name of the primary contact.
FAMILY_SIZE	INTEGER	Size of the family.

COUNTRIES		
Name	Type	Description
ID	STRING	Unique ID of the country.
NAME	STRING	Name of the country.
MIN_SIZE	INTEGER	Minimum size group to get a discount.

▼ Sample Data Tables

FAMILIES		
ID	NAME	FAMILY_SIZE
c00dac11bde74750b4d207b9c182a85f	Alex Thomas	9
eb6f2d3426694667ae3e79d6274114a4	Chris Gray	2

COUNTRIES		
ID	NAME	MIN_SIZE
023fd23615bd4ff4b2ae0a13ed7efec9	Bolivia	2
be247f73de0f4b2d810367cb26941fb9	Cook Islands	4
3e85ab80a6f84ef3b9068b21dbcc54b3	Brazil	4

Sample Output

3

Explanation

Question - 83

Activity Query

A parent keeps track of the activities of a child and their friends in two tables: *FRIENDS* and *ACTIVITIES*. Write a query to print the names of all the activities with neither the maximum nor minimum number of participants.

▼ Schema

There are 2 tables: *FRIENDS*, *ACTIVITIES*.

FRIENDS		
Name	Type	Description
ID	INTEGER	The ID of a friend. This is the primary key.
NAME	STRING	The name of the friend.
ACTIVITY	STRING	Name of the activity which the friend takes part in.

ACTIVITIES		
Name	Type	Description
ID	INTEGER	ID of the activity.
NAME	STRING	Name of the activity.

▼ Sample Data Tables

Sample Input

FRIENDS		
ID	NAME	ACTIVITY
1	James Smith	Horse Riding
2	Eric Jenkins	Eating
3	Sean Cox	Eating
4	Eric Schmidt	Horse Riding
5	Chris Evans	Eating
6	Jessica Breeds	Playing

Activities	
ID	NAME
1	Horse Riding
2	Eating
3	Playing

Sample Output

Horse Riding

Question - 84 Restaurant's Growth

A restaurant is visited by various customers during a day. At the same time, the restaurant is advertising to increase customer revenue. Write an SQL query to compute the moving average of how much customers spent over a window of 7 days (current day and 6 days before) to analyze their business growth. The output should contain three columns {visited_on, amount, avg_amount(avg over 7 days)}

▼ Schema

There is 1 table: customers

CUSTOMERS		
Name	Type	Description
id	INTEGER	The customer's id, Primary Key
name	VARCHAR	Name of the customer
phone	VARCHAR	Phone of the customer
visited_on	DATE	Date that the customer visited
amount	INTEGER	Amount he customer spent

▼ Sample Data Tables

CUSTOMERS				
id	name	phone	visited_on	amount
1	Julia	1234567890	2015-05-01	100
2	Samantha	1234567890	2015-05-02	200
3	Julia-Samantha	1234567890	2015-05-03	300

Sample Output

2015-05-01 100 100
2015-05-02 200 150
2015-05-03 300 200

Explanation

Row 1 : On May 1, Julia visited and spent 100. $\text{avg_amount} = 100/1$
Row 2 : On May 2, Samantha visited and spent 200. $\text{avg_amount} = (100+200)/2 = 150$
Row 3 : On May 3, they both visited and spent 300. $\text{avg_amount} = (100+200+300)/3 = 200$

Question - 85

Examination Data Management

A college maintains the data of its students and their respective appeared examinations in two tables: STUDENT and EXAMINATION. A student might appear for an exam in a particular subject more than once. Write a query to print the student ID, the subject name, and the total number of times the student appeared in a particular subject's examination. The order of output does not matter. The result should be in the following format: STUDENT.ID EXAMINATION.SUBJECT NUMBER_OF_TIMES

▼ Schema

STUDENT		
Name	Type	Description
ID	Integer	The student's ID number. This is the primary key.
NAME	String	The student's name.

EXAMINATION		
Name	Type	Description
STUDENT_ID	Integer	The student's ID number.
SUBJECT	String	The subject's name.

▼ Sample Data Tables

STUDENT	
ID	NAME
1	Taylor
2	Wesley
3	Jordan
4	Robin
5	Alex

EXAMINATION	
STUDENT_ID	SUBJECT
1	Biology
1	Physics
3	History
4	Geography
4	Geography

Sample Output

```
1 Biology 1
1 Physics 1
3 History 1
4 Geography 2
```

Explanation

- Taylor appeared in *Biology* examination one time.
- Taylor appeared in the *Physics* examination one time.
- Jordan appeared in the *History* examination one time.
- Robin appeared in the *Geography* examination two times.

Question - 86

The Perfect Arrangement

Write a query to print the *id*, *first_name* and *last_name*. To filter the names, concatenate the first and last names to create a *combined name*. Return the names of customers whose combined names are less than 12 letters long. Sort the results by their combined name lengths, then alphabetically, case insensitive, by combined name, then by id. All sorts are ascending.

Input Format

CUSTOMER

Name	Type	Description
ID	Integer	unique id, primary key.
FIRST_NAME	String	
LAST_NAME	String	
COUNTRY	String	
CREDIT_LIMIT	Float	

Output Format

```
CUSTOMER.ID CUSTOMER.FIRST_NAME CUSTOMER.LAST_NAME
```

Sample Input

CUSTOMER

ID	FIRST_NAME	LAST_NAME	COUNTRY	CREDIT_LIMIT
1	Alex	White	USA	200350.54
2	Tyler	Hanson	UK	15354.23
3	Jordan	Fernandez	France	359200.67
4	Drew	Bradley	Albania	1060.57
5	Blake	Fuller	USA	14789.00
6	Spencer	Johnston	China	100243.35
7	Ellis	Gutierrez	USA	998999.20
8	Morgan	Thomas	Canada	500500.23
9	Riley	Garza	UK	18782.44

10	Peyton	Harris	USA	158367.00
----	--------	--------	-----	-----------

Sample Output

1	Alex White
9	Riley Garza
5	Blake Fuller
4	Drew Bradley
2	Tyler Hanson

Explanation

AlexWhite is 9 letters, so it is included in the results. JordanFernandez is 15 letters, so it is omitted. The last 3 names are the same length, so they are sorted alphabetically ascending.

Names that are excluded and their lengths

MorganThomas	12
PeytonHarris	12
EllisGutierrez	14
JordanFernandez	15
SpencerJohnston	15

Question - 87

Students Score

The Math scores of students have been stored in the *STUDENT* table. Write a query to print the ids and the names of the students who have secured higher than the average score. The result should be sorted in ascending order of student ID.

Input Format

STUDENT		
Name	Type	Description
ID	Integer	A student ID in the inclusive range [1, 1000]. This field is the primary key.
NAME	String	A student name. This field contains between 1 and 100 characters (inclusive).
SCORE	Float	The Math score of the student.

Output Format

The result should contain the ids and the names of the students who secured higher than the average score. The result should be sorted in ascending order of student ID.

STUDENT.ID	STUDENT.NAME
------------	--------------

Sample Input

STUDENT

ID	NAME	SCORE
1	Bob	50
2	John	65.5
3	Harry	45
4	Dick	85
5	Dev	25
6	Sid	98
7	Tom	90
8	Julia	70.5
9	Erica	81
10	Jerry	85

Sample Output

```
4 Dick
6 Sid
7 Tom
8 Julia
9 Erica
10 Jerry
```

Explanation

The average score is $(50 + 65.5 + 45 + 85 + 25 + 98 + 90 + 70.5 + 81 + 85) / 10 = 69.5$.

- *Bob* scored 50, which is less than 69.5.
- *John* scored 65.5, which is less than 69.5.
- *Harry* scored 45, which is less than 69.5.
- *Dick* scored 85, which is more than 69.5.
- *Dev* scored 25, which is less than 69.5.
- *Sid* scored 98, which is more than 69.5.
- *Tom* scored 90, which is more than 69.5.
- *Julia* scored 70.5, which is more than 69.5.
- *Erica* scored 81, which is more than 69.5.
- *Jerry* scored 85, which is more than 69.5.

Sid, Dick, Tom, Julia, Erica, and Jerry scored more than the average score.

Question - 88**The First Orders**

A company maintains information about its orders in the *ORDERS* table. Write a query to print details of the earliest *five* orders (sorted by *ORDER_DATE*, ascending) that have not been delivered (i.e., *STATUS* is not *DELIVERED*). If there are more than five orders to choose from, select the ones with the lowest order ID. Sort the output in the increasing order of order *ID*. The output should contain *ID*, *ORDER_DATE*, *STATUS*, *CUSTOMER_ID*.

▼ Schema

Table: Orders

column name	column type
id	int
order_date	date
status	varchar(50)
customer_id	int

▼ Sample Data Tables

Sample Input

ORDERS			
ID	ORDER_DATE	STATUS	CUSTOMER_ID
10100	2003-01-06	PLACED	363
10101	2003-01-06	PLACED	128
10102	2003-01-06	IN TRANSIT	181
10103	2003-01-06	DELIVERED	121
10104	2003-01-07	DELIVERED	114
10106	2003-01-07	IN TRANSIT	278
10120	2003-01-07	PLACED	114
10122	2003-05-05	IN TRANSIT	350
10123	2003-05-05	DELIVERED	103

Sample Output

```
10100 2003-01-06 PLACED 363
10101 2003-01-06 PLACED 128
10102 2003-01-06 IN TRANSIT 181
10106 2003-01-07 IN TRANSIT 278
10120 2003-01-07 PLACED 114
```

Explanation

The orders with order numbers *10100*, *10101*, *10102*, *10106*, and *10120* are the earliest placed orders and also have order status not equal to *DELIVERED*, so all their information in the increasing order of their order ID is printed.

Question - 89

Customers Credit Limit

A company maintains the data of its customers in the *CUSTOMER* table. Write a query to print the *IDs* and the *NAMES* of the customers who are from the USA and whose credit limit is greater than *100000*, ordered by increasing *ID* number.

Input Format

CUSTOMER

Name	Type	Description
ID	Integer	A customer ID in the inclusive range [1, 1000]. This is the primary key.
NAME	String	A customer name. This field contains between 1 and 100 characters (inclusive).
COUNTRY	String	The country of the customer.
CREDITS	Integer	The credit limit of the customer.

Output Format

The result should print the *IDs* and the *NAMEs* of those customers who are from the USA and whose credit limit is greater than 100000, in ascending *ID* order and in the following format:

```
CUSTOMER.ID CUSTOMER.NAME
```

Sample Input

CUSTOMER			
ID	NAME	COUNTRY	CREDITS
1	Frances White	USA	200350
2	Carolyn Bradley	UK	15354
3	Annie Fernandez	France	359200
4	Ruth Hanson	Albania	1060
5	Paula Fuller	USA	14789
6	Bonnie Johnston	China	100243
7	Ruth Gutierrez	USA	998999
8	Ernest Thomas	Canada	500500
9	Joe Garza	UK	18782
10	Anne Harris	USA	158367

Sample Output

```
1 Frances White
7 Ruth Gutierrez
10 Anne Harris
```

Explanation

Description of some of the customers is given below:

- *Frances White* is from the *USA*, and his credit limit is 200350, which is greater than 100000.
- *Carolyn Bradley* is from the *UK*, and her credit limit is 15354, which is less than 100000.
- *Paula Fuller* is from the *USA*, and her credit limit is 14789, which is less than 100000.
- Remaining records are analyzed similarly.

So, *Frances White*, *Ruth Gutierrez*, and *Anne Harris* are from *USA* and credit is greater than 100000.

Question - 90

The Beautiful Collection

A shopkeeper maintains the count of the different colored balls (Red, green, and blue) in the *COLLECTION* table. Each row of the table represents one of the following types:

- **GOOD** : If the count of the red, green, and blue balls are equal.
- **BAD** : If the count of any two colored balls are equal, i.e., only one of the following conditions is true:
 - Red balls count is equal to green balls.
 - Red balls count is equal to blue balls.
 - Green balls count is equal to blue balls.
- **WORSE** : If all the colored balls have different counts.

Write a query to print the type which is represented by each row of the table. Note that the output is *case-sensitive*, so make sure to output only **GOOD**, **BAD**, or **WORSE**.

Input Format

COLLECTION		
Name	Type	Description
RED	Integer	This describes the count of red balls. The count can be between 30 units and 100 units (inclusive).
GREEN	Integer	This describes the count of green balls. The count can be between 30 units and 100 units (inclusive).
BLUE	Integer	This describes the count of blue balls. The count can be between 30 units and 100 units (inclusive).

Output Format

Each row of the result should contain one of the types which are described above, in the following format. Note that the output is *case-sensitive*, so make sure to output only **GOOD**, **BAD**, or **WORSE**.

TYPE_OF_COLLECTION

Sample Input

COLLECTION		
RED	GREEN	BLUE
65	65	87
50	50	50
30	50	100
40	50	90
92	50	50

Sample Output

BAD
GOOD
WORSE
WORSE
BAD

Explanation

The type of collection represented by each row is explained below:

- The count of the red balls is equal to the count of the green balls. But the count of blue balls is not equal to the count of red or green balls. So, the type is **BAD**.
- The count of the red, green, and blue balls are equal. So, the type is **GOOD**.
- The count of the red, green, and blue balls are different from each other. So, the type is **WORSE**.
- The count of the red, green, and blue balls are different from each other. So, the type is **WORSE**.
- The count of the green balls is equal to the count of the blue balls. But the count of red balls is not equal to the count of green or blue balls. So, the type is **BAD**.

Question - 91

Big Companies

An organization maintains employment data in three tables: *EMPLOYEE*, *COMPANY*, and *SALARY*. Write a query to print the names of every company where the average salary is greater than *40000*. Each distinct row of results in the output must contain the name of a company whose average employee salary is > *40,000* in the *COMPANY.NAME* format.

▼ Schema

EMPLOYEE		
Name	Type	Description
ID	Integer	An employee ID in the inclusive range [1, 1000]. This is the primary key.
NAME	String	An employee name. This field contains between 1 and 100 characters (inclusive).

COMPANY		
Name	Type	Description
ID	Integer	A company ID in the inclusive range [1, 1000]. This is the primary key.
NAME	String	A company name. This field contains between 1 and 100 characters (inclusive).

SALARY		
Name	Type	Description
EMPLOYEE_ID	Integer	An employee ID in the inclusive range [1, 1000].
COMPANY_ID	Integer	A company ID in the inclusive range [1, 1000].
SALARY	Integer	The salary of the employee in the inclusive range [10000, 100000].

▼ Sample Data Tables

EMPLOYEE	
ID	NAME
1	Frances White
2	Carolyn Bradley

3	Annie Fernandez
4	Ruth Hanson
5	Paula Fuller
6	Bonnie Johnston
7	Ruth Gutierrez
8	Ernest Thomas
9	Joe Garza
10	Anne Harris

COMPANY	
ID	NAME
1	PeopleSoft Inc
2	Baker Hughes Incorporated
3	MDU Resources Group Inc.
4	DST Systems, Inc.
5	Williams Companies Inc
6	Fisher Scientific International Inc.
7	Emcor Group Inc.
8	Genuine Parts Company
9	MPS Group Inc.
10	Novellus Systems Inc

SALARY		
EMPLOYEE_ID	COMPANY_ID	SALARY
2	4	27779
2	9	36330
3	9	71466
3	10	22804
5	5	49892
6	4	31493
6	10	26888
7	3	87118
7	7	70767
7	9	39929

Sample Output

MDU Resources Group Inc.
Williams Companies Inc
Emcor Group Inc.
MPS Group Inc.

Explanation

The following companies have average employee salaries > 40,000.

- *MDU Resources Group Inc.* has one employee (*ID 7*) whose *SALARY* is 87118. The company's average employee salary is $87118 \div 1 = 87118$ and $87118 > 40000$.
- *Williams Companies Inc* has one employee (*ID 5*) whose *SALARY* is 49892. The company's average employee salary is $49892 \div 1 = 49892$.
- *Emcor Group Inc.* has one employee (*ID 7*) whose *SALARY* is 70767. The company's average employee salary is $70767 \div 1 = 70767$.
- *MPS Group Inc.* has three employees (*IDs 2, 3, and 7*) whose *SALARY* values are 36330, 71466, and 39929. The company's average employee salary $(36330 + 71466 + 39929) \div 3 = 49241.67$.

Question - 92

Scheduling Errors

Write a query to return a list of professor names and their associated courses for all courses outside of their departments. There should be no duplicate rows, but they can be in any order.

The output should contain two columns: *professor.name, course.name*.

▼ Schema

PROFESSOR		
Name	Type	Description
ID	Integer	unique id, primary key
NAME	String	
DEPARTMENT_ID	Integer	foreign key, <i>department.id</i>
SALARY	Integer	

DEPARTMENT		
Name	Type	Description
ID	Integer	unique id, primary key
NAME	String	

COURSE		
Name	Type	Description
ID	Integer	unique id, primary key
NAME	String	
DEPARTMENT_ID	Integer	foreign key, <i>department.id</i>
CREDITS	Integer	

SCHEDULE		
Name	Type	Description
PROFESSOR_ID	Integer	foreign key, <i>professor.id</i>
COURSE_ID	Integer	foreign key, <i>course.id</i>
SEMESTER	Integer	
YEAR	Integer	

▼ Sample Data Tables

PROFESSOR			
ID	NAME	DEPARTMENT_ID	SALARY
1	Alex Daniels	4	7169
2	Drew Knight	1	9793
3	Jordan Myers	4	25194
4	Tyler Rodriguez	3	9686
5	Blake Gome	2	30860
6	Spencer George	5	10487
7	Ellis Vasquez	4	6353
8	Morgan Flores	1	25796
9	Riley Gilbert	5	35678
10	Peyton Stevens	2	26648

DEPARTMENT	
ID	NAME
3	Biological Sciences
5	Technology
6	Humanities & Social Sciences
2	Clinical Medicine
4	Arts and Humanities
1	Physical Sciences

COURSE			
ID	NAME	DEPARTMENT_ID	CREDITS
9	Clinical Biochemistry	2	3
4	Astronomy	1	6
10	Clinical Neuroscience	2	5
1	Pure Mathematics and Mathematical Statistics	1	3
6	Geography	1	7
8	Chemistry	1	1
5	Physics	1	8
3	Earth Science	1	7
7	Materials Science and Metallurgy	1	5
2	Applied Mathematics and Theoretical Physics	1	5

SCHEDULE			
PROFESSOR_ID	COURSE_ID	SEMESTER	YEAR
4	4	3	2003
3	3	1	2011
1	7	5	2011
7	7	1	2010
4	6	1	2001
9	3	1	2012
10	2	4	2009
1	1	3	2014
1	2	3	2008
1	7	5	2007

Sample Output

Tyler Rodriguez Astronomy
 Jordan Myers Earth Sciences
 Alex Daniels Materials Science and Metallurgy
 Ellis Vasquez Materials Science and Metallurgy
 Tyler Rodriguez Geography
 Riley Gilbert Earth Sciences
 Peyton Stevens Applied Mathematics and Theoretical Physics
 Alex Daniels Pure Mathematics and Mathematical Statistics
 Alex Daniels Applied Mathematics and Theoretical Physics
 Alex Daniels Materials Science and Metallurgy

Explanation

Example logic

1. Professor *Tyler Rodriguez's department_id* is 3, but the *Astronomy* course's *department_id* is 1.
2. Professor *Jordan Myers's department_id* is 4, but the *Earth Sciences* course's *department_id* is 1

Question - 93

List the Course Names

Write a query to return a list of professor names and their associated courses. The results can be in any order but must not contain duplicate rows.

▼ Schema

PROFESSOR		
Name	Type	Description
ID	Integer	unique id, primary key
NAME	String	
DEPARTMENT_ID	Integer	This is a foreign key to DEPARTMENT.ID.
SALARY	Integer	

DEPARTMENT		
Name	Type	Description
ID	Integer	unique id, primary key
NAME	String	

COURSE		
Name	Type	Description
ID	Integer	unique id, primary key
NAME	String	
DEPARTMENT_ID	Integer	foreign key to DEPARTMENT.ID
CREDITS	Integer	

SCHEDULE		
Name	Type	Description
PROFESSOR_ID	Integer	foreign key to PROFESSOR.ID
COURSE_ID	Integer	foreign key to COURSE.ID
SEMESTER	Integer	
YEAR	Integer	

▼ Sample Data Tables

PROFESSOR			
ID	NAME	DEPARTMENT_ID	SALARY
1	Alex Burton	5	7340
8	Jordan Diaz	1	17221
9	Drew Hicks	5	16613
2	Tyler Matthews	2	14521
10	Blake Foster	4	28526
3	Spencer Peters	1	10487
4	Ellis Marshall	3	6353
7	Morgan Lee	2	25796
5	Riley Peterson	1	35678
6	Peyton Fields	5	26648

DEPARTMENT	
ID	NAME
3	Biological Sciences
5	Technology
6	Humanities & Social Sciences
2	Clinical Medicine
4	Arts and Humanities
1	Physical Sciences

COURSE			
ID	NAME	DEPARTMENT_ID	CREDITS
9	Clinical Biochemistry	2	3
4	Astronomy	1	6
10	Clinical Neuroscience	2	5
1	Pure Mathematics and Mathematical Statistics	1	3
6	Geography	1	7
8	Chemistry	1	1
5	Physics	1	8
3	Earth Science	1	7
7	Materials Science and Metallurgy	1	5
2	Applied Mathematics and Theoretical Physics	1	5

SCHEDULE			
PROFESSOR_ID	COURSE_ID	SEMESTER	YEAR
5	3	6	2012
7	3	1	2013
5	7	6	2010
2	10	2	2004
5	1	1	2011
2	9	4	2005
7	10	6	2009
5	6	4	2007
7	9	1	2014
9	9	5	2011

Sample Output

```

Tyler Matthews Clinical Biochemistry
Tyler Matthews Clinical Neuroscience
Drew Hicks Clinical Biochemistry
Morgan Lee Clinical Biochemistry
Morgan Lee Clinical Neuroscience
Morgan Lee Earth Science
Riley Peterson Earth Science
Riley Peterson Geography
Riley Peterson Materials Science and Metallurgy
Riley Peterson Pure Mathematics and Mathematical Statistics

```

Question - 94

Professor Names and Salaries

A university maintains data on professors and departments in two tables: *PROFESSOR* and *DEPARTMENT*. Write a query to print the NAME and SALARY for each professor who satisfies the following two requirements:

- The professor does not work in the *Arts and Humanities* department.
- The professor's salary is *greater than the smallest* salary of any professor in the *Arts and Humanities* department.

The name must be printed before the salary, but row order does not matter.

▼ Schema

DEPARTMENT		
Name	Type	Description
ID	Integer	A department ID in the inclusive range [1, 1000]. This is a primary key.
NAME	String	A department name. This field contains between 1 and 100 characters.

PROFESSOR

Name	Type	Description

ID	Integer	A professor's ID in the inclusive range [1, 1000]. This is a primary key.
NAME	String	A professor's name. This field contains between 1 and 100 characters.
DEPARTMENT_ID	Integer	A professor's department ID. This is a foreign key to DEPARTMENT.ID.
SALARY	Integer	A professor's salary in the inclusive range [5000, 40000].

▼ Sample Data Tables

DEPARTMENT	
ID	NAME
3	Biological Sciences
5	Technology
6	Humanities & Social Sciences
2	Clinical Medicine
4	Arts and Humanities
1	Physical Sciences

PROFESSOR			
ID	NAME	DEPARTMENT_ID	SALARY
1	Shauna Rivera	1	22606
8	Ruth Price	3	9287
9	Julie Gonzalez	4	18870
2	Craig Elliott	5	27524
10	Scott Butler	1	26200
3	Nancy Russell	2	7076
4	Clarence Johnson	1	7249
7	Louis Schmidt	1	13437
5	Terri Thompson	3	28432
6	Keith Gilbert	5	12610

Sample Output

```
Shauna Rivera 22606
Craig Elliott 27524
Terri Thompson 28432
Scott Butler 26200
```

Explanation

Julie Gonzalez has a salary of 18870 , which is smaller than the salary of any other professor in the Arts and Humanities department. The following employees of other departments have salaries higher than Julie's:

- Shauna Rivera 's salary of 22606 is higher than Julie Gonzalez 's.

- Craig Elliott's salary of 27524 is higher than Julie Gonzalez's.
- Terri Thompson's salary of 28432 is higher than Julie Gonzalez's.
- Scott Butler's salary of 26200 is higher than Julie Gonzalez's.

Question - 95 Student's Major

A university maintains data on students and their majors in three tables: STUDENTS, MAJORS, and REGISTER. The university needs a list of STUDENT_NAME and MAJOR_NAME. Sort the list by STUDENT_ID and return the first 20 records.

▼ Table Schema

STUDENTS			MAJORS		
Name	Type	Description	Name	Type	Description
STUDENT_ID	Integer	The ID of a student. This is the <i>primary key</i> .	MAJOR_ID	Integer	The ID of a major. This is the <i>primary key</i> .
STUDENT_NAME	String	The name of the student.	MAJOR_NAME	String	The name of a major.
STUDENT_AGE	Integer	The age of the student.			

REGISTER		
Name	Type	Description
STUDENT_ID	Integer	The ID of a student. This is a <i>foreign key</i> .
MAJOR_ID	Integer	The ID of a major. This is a <i>foreign key</i> .

▼ Sample Case 0

Sample Input 0

STUDENTS			MAJORS		REGISTER	
STUDENT_ID	STUDENT_NAME	STUDENT_AGE	MAJOR_ID	MAJOR_NAME	STUDENT_ID	MAJOR_ID
1	John	20	1000	Computer Science	2	1000
2	Masie	21	2000	Biology	3	3000
3	Harry	21	3000	Physics	1	2000

Sample Output

```
John Biology
Masie Computer Science
Harry Physics
```

Question - 96 Student Rank

A university stores students' standardized test scores in a table named *STUDENT*. Student X placed 213th on the test. Write a query to find Student X's test score (i.e., the 213th highest *STUDENT.SCORE* in *STUDENT*).

▼ Schema

A table is provided: STUDENT .

STUDENT		
Name	Type	Description
ID	Integer	The student's unique ID. This is a <i>primary key</i> .
AGE	Integer	The student's age.
SCORE	Integer	The student's standardized test score.

▼ Sample Data Tables

STUDENT		
ID	AGE	SCORE
1	19	91
2	20	90
3	20	87
4	21	72
5	19	98
6	20	50

The table's scores (from highest to lowest) are {98, 91, 90, 87, 72, 50}. As an example, the fourth-highest score is 87. For the real query, find the 213th highest score.

Question - 97

Clumsy Administrator

A company maintains the records of all employees. The company pays the database administrator too little so the work has been quite clumsy. The administrator carelessly inserted the records of many employees into the employee records table multiple times. An employee's record is considered duplicate only if all columns (fields) of the employee's record are duplicated.

Write a query to find the names of employees whose records occur more than once in the table. Each name should appear only once in the results. The order of results does not matter.

▼ Schema

You are provided 1 table: EMPLOYEE .

EMPLOYEE		
Name	Type	Description
NAME	STRING	The name of the employee.
PHONE	STRING	The telephone number of the employee.
AGE	INTEGER	The age of the employee.

▼ Sample Data Tables

EMPLOYEE		
NAME	PHONE	AGE
Sam	1000040000	30
Alex	1000020000	60
Alex	1000020012	65
Sam	1000040000	30
Chris	1000012000	34
Chris	1000012000	34

Here, the exact records of Sam and Chris occur more than once in the table. Hence, Sam and Chris are the employees in the resultant output.

Question - 98

Accounting Software Balance Report

HackerFinance is developing an accounting software and needs to generate a list of customers and the change in their account balances during December.

The output should display each customer and their corresponding balance, formatted as "customer balance". The balance should show two places after the decimal, e.g., 0.00.

The change is calculated as the sum of debit transactions minus the sum of credit transactions during December. The results should be sorted in ascending order by customer name.

▼ Schema

There is 1 table:

transactions		
name	type	description
dt	VARCHAR(19)	Transaction timestamp
customer	VARCHAR(64)	Customer name
debit	DECIMAL(5,2)	Transaction debit
credit	DECIMAL(5,2)	Transaction credit

▼ Sample Data Tables

For the sample data in table:

transactions

dt	customer	debit	credit
2021-11-30 12:48:22	Arney Cuff	6.43	16.12
2021-12-25 19:00:46	Arney Cuff	97.78	12.53
2021-11-27 21:34:24	Donaugh Furneaux	89.71	85.04
2021-11-25 07:31:37	Ferrell Brunn	63.58	28.58
2021-11-25 15:30:56	Gibbie Jurisic	25.81	13.75
2021-11-21 00:09:50	Harley Lyddiard	57.49	7.11
2021-12-01 07:37:42	Harley Lyddiard	48.33	82.35
2021-12-02 13:08:52	Harley Lyddiard	12.13	63.81
2021-11-24 03:51:13	Kippy Jelly	50.34	12.91
2021-12-04 10:11:40	Latrina Jackman	10.73	39.51
2021-12-12 13:02:50	Latrina Jackman	5.35	96.74
2021-12-20 17:31:44	Latrina Jackman	54.99	92.73
2021-12-30 13:40:43	Maribel Braim	57.06	21.37
2021-12-02 09:57:35	Orrin Curley	65.44	51.31
2021-12-14 19:57:25	Orrin Curley	40.04	96.44
2021-11-20 09:28:11	Rasla Venny	80.33	20.69
2021-12-08 03:31:31	Rasla Venny	55.43	99.04
2021-12-09 21:31:29	Rasla Venny	87.96	5.87
2021-12-11 13:02:54	Rasla Venny	45.42	55.81
2021-12-28 18:04:52	Rasla Venny	68.17	85.30

the expected output is:

customer ▲	balance
Arney Cuff	85.25
Harley Lyddiard	-85.70
Latrina Jackman	-157.91
Maribel Braim	35.69
Orrin Curley	-42.27
Rasla Venny	10.96